

The **OpenCL C++** Specification

Version: 1.0

Document Revision: 08

Khronos OpenCL Working Group

Editor: Aaftab Munshi

Provisional Specification

Copyright (c) 2008-2015 The Khronos Group Inc. All Rights Reserved.

This specification is protected by copyright laws and contains material proprietary to the Khronos Group, Inc. It or any components may not be reproduced, republished, distributed, transmitted, displayed, broadcast or otherwise exploited in any manner without the express prior written permission of Khronos Group. You may use this specification for implementing the functionality therein, without altering or removing any trademark, copyright or other notice from the specification, but the receipt or possession of this specification does not convey any rights to reproduce, disclose, or distribute its contents, or to manufacture, use, or sell anything that it may describe, in whole or in part.

Khronos Group grants express permission to any current Promoter, Contributor or Adopter member of Khronos to copy and redistribute UNMODIFIED versions of this specification in any fashion, provided that NO CHARGE is made for the specification and the latest available update of the specification for any version of the API is used whenever possible. Such distributed specification may be re-formatted AS LONG AS the contents of the specification are not changed in any way. The specification may be incorporated into a product that is sold as long as such product includes significant independent work developed by the seller. A link to the current version of this specification on the Khronos Group web-site should be included whenever possible with specification distributions.

Khronos Group makes no, and expressly disclaims any, representations or warranties, express or implied, regarding this specification, including, without limitation, any implied warranties of merchantability or fitness for a particular purpose or non-infringement of any intellectual property. Khronos Group makes no, and expressly disclaims any, warranties, express or implied, regarding the correctness, accuracy, completeness, timeliness, and reliability of the specification. Under no circumstances will the Khronos Group, or any of its Promoters, Contributors or Members or their respective partners, officers, directors, employees, agents or representatives be liable for any damages, whether direct, indirect, special or consequential damages for lost revenues, lost profits, or otherwise, arising from or in connection with these materials.

Khronos, StreamInput, WebGL, COLLADA, OpenKODE, OpenVG, OpenWF, OpenSL ES, OpenMAX, OpenMAX AL, OpenMAX IL and OpenMAX DL are trademarks and WebCL is a certification mark of the Khronos Group Inc. OpenCL is a trademark of Apple Inc. and OpenGL and OpenML are registered trademarks and the OpenGL ES and OpenGL SC logos are trademarks of Silicon Graphics International used under license by Khronos. All other product names, trademarks, and/or company names are used solely for identification and belong to their respective owners.

1	The OpenCL C++ Programming Language	5
1.1	Supported Data Types	5
1.1.1	Built-in Scalar Data Types	5
1.1.2	Built-in Vector Data Types	8
1.1.3	Other Built-in Data Types	9
1.1.4	Reserved Data Types	11
1.1.5	Alignment of Types	12
1.1.6	Vectors	12
1.1.7	Keywords	18
1.2	Conversions and Type Casting	19
1.2.1	Implicit Conversions	19
1.2.2	Explicit Casts	19
1.2.3	Explicit Conversions	20
1.2.4	Reinterpreting Data As Another Type	23
1.3	Operators	26
1.4	Vector Operations	31
1.5	Address Spaces	32
1.6	Function Qualifiers	36
1.7	Storage-Class Specifiers	36
1.8	Restrictions	37
1.9	Pre-processor Directives and Macros	38
1.10	Attribute Qualifiers	39
1.10.1	Optional Attribute Qualifiers	39
1.10.2	Specifying Attributes for Unrolling Loops	40
1.10.3	Extending Attribute Qualifiers	43
2	OpenCL C++ Standard Library	44
2.1	Work-Item Functions	44
2.2	Math Functions	48
2.2.1	Floating-point macros and pragmas	55
2.3	Integer Functions	60
2.4	Common Functions	64
2.5	Geometric Functions	67
2.6	Relational Functions	68
2.7	Vector Data Load and Store Functions	72
2.8	Synchronization Functions	76
2.9	Address Space Qualifier Functions	80
2.10	Atomics	81
2.11	printf	96
2.11.1	printf output synchronization	96
2.11.2	printf format string	97
2.11.3	Differences between OpenCL C++ and C++14 printf	103
2.12	Images	104
2.12.1	Image Types	104
2.12.2	Samplers	104
2.12.3	Image Sample Functions	108
2.12.4	Image Read Functions	109
2.12.5	Image Write Functions	111
2.12.6	Image Query Functions	112
2.12.7	Reading and writing to the same image in a kernel	115

2.12.8 Mapping image channels to color values returned by image_sample, image_read and color values passed to image_write to image channels	116
2.13 Work-group Functions	118
2.14 Pipe Functions.....	121
2.15 Enqueuing kernels.....	129
2.16 Metaprogramming and type traits.....	142
2.17 Diagnostics.....	143
3 OpenCL Numerical Compliance.....	144
3.1 Rounding Modes	144
3.2 INF, NaN and Denormalized Numbers	144
3.3 Floating-Point Exceptions.....	145
3.4 Relative Error as ULPs	145
3.5 Edge Case Behavior	154
3.5.1 Additional Requirements Beyond ISO/IEC 9899:TC2	154
3.5.2 Changes to ISO/IEC 9899: TC2 Behavior	157
3.5.3 Edge Case Behavior in Flush To Zero Mode	158
4 Image Addressing and Filtering.....	159
4.1 Image Coordinates	159
4.2 Addressing and Filter Modes.....	159
4.3 Conversion Rules	166
4.3.1 Conversion rules for normalized integer channel data types.....	166
4.3.2 Conversion rules for half precision floating-point channel data type.....	170
4.3.3 Conversion rules for floating-point channel data type	170
4.3.4 Conversion rules for signed and unsigned 8-bit, 16-bit and 32-bit integer channel data types	171
4.3.5 Conversion rules for sRGBA and sBGRA images	172
4.4 Selecting an Image from an Image Array.....	173

1 The OpenCL C++ Programming Language

This section describes the OpenCL C++ programming language used to create kernels that are executed on OpenCL device(s). The OpenCL C++ programming language is based on the ISO/IEC JTC1 SC22 WG21 N 3690 language specification (a.k.a. C++14 specification) with specific extensions and restrictions. Please refer to this specification for a detailed description of the language grammar. This section describes modifications and restrictions to the C++14 specification supported in OpenCL C++.

1.1 Supported Data Types

The following data types are supported.

1.1.1 Built-in Scalar Data Types

Table 1.1 describes the list of built-in scalar data types.

Type	Description
<code>bool</code> ¹	A conditional data type which is either <i>true</i> or <i>false</i> . The value <i>true</i> expands to the integer constant 1 and the value <i>false</i> expands to the integer constant 0.
<code>char</code>	A signed two's complement 8-bit integer.
<code>unsigned char</code> , <code>uchar</code>	An unsigned 8-bit integer.
<code>short</code>	A signed two's complement 16-bit integer.
<code>unsigned short</code> , <code>ushort</code>	An unsigned 16-bit integer.
<code>int</code>	A signed two's complement 32-bit integer.
<code>unsigned int</code> , <code>uint</code>	An unsigned 32-bit integer.
<code>long</code>	A signed two's complement 64-bit integer.
<code>unsigned long</code> , <code>ulong</code>	An unsigned 64-bit integer.
<code>float</code>	A 32-bit floating-point. The float data type must conform to the IEEE 754 single precision storage format.
<code>double</code> ²	A 64-bit floating-point. The double data type must conform

¹ When any scalar value is converted to **bool**, the result is 0 if the value compares equal to 0; otherwise, the result is 1.

² The `double` data type is an optional type that is supported if `CL_DEVICE_DOUBLE_FP_CONFIG` in table 4.3 for a device is not zero.

	to the IEEE 754 double precision storage format.
<code>half</code>	A 16-bit floating-point. The half data type must conform to the IEEE 754-2008 half precision storage format.
<code>size_t</code>	The unsigned integer type of the result of the sizeof operator. This is a 32-bit unsigned integer if <code>CL_DEVICE_ADDRESS_BITS</code> defined in <i>table 4.3</i> is 32-bits and is a 64-bit unsigned integer if <code>CL_DEVICE_ADDRESS_BITS</code> is 64-bits.
<code>ptrdiff_t</code>	A signed integer type that is the result of subtracting two pointers. This is a 32-bit signed integer if <code>CL_DEVICE_ADDRESS_BITS</code> defined in <i>table 4.3</i> is 32-bits and is a 64-bit signed integer if <code>CL_DEVICE_ADDRESS_BITS</code> is 64-bits.
<code>intptr_t</code>	A signed integer type with the property that any valid pointer to void can be converted to this type, then converted back to pointer to void , and the result will compare equal to the original pointer. This is a 32-bit signed integer if <code>CL_DEVICE_ADDRESS_BITS</code> defined in <i>table 4.3</i> is 32-bits and is a 64-bit signed integer if <code>CL_DEVICE_ADDRESS_BITS</code> is 64-bits.
<code>uintptr_t</code>	An unsigned integer type with the property that any valid pointer to void can be converted to this type, then converted back to pointer to void , and the result will compare equal to the original pointer. This is a 32-bit signed integer if <code>CL_DEVICE_ADDRESS_BITS</code> defined in <i>table 4.3</i> is 32-bits and is a 64-bit signed integer if <code>CL_DEVICE_ADDRESS_BITS</code> is 64-bits.
<code>void</code>	The void type comprises an empty set of values; it is an incomplete type that cannot be completed.

Table 1.1 *Built-in Scalar Data Types*

Most built-in scalar data types are also declared as appropriate types in the OpenCL API (and header files) that can be used by an application. The following table describes the built-in scalar data type in the OpenCL C++ programming language and the corresponding data type available to the application:

Type in OpenCL Language	API type for application
<code>bool</code>	n/a
<code>char</code>	<code>cl_char</code>
<code>unsigned char,</code> <code>uchar</code>	<code>cl_uchar</code>
<code>short</code>	<code>cl_short</code>
<code>unsigned short,</code> <code>ushort</code>	<code>cl_ushort</code>

Provisional Specification

<code>int</code>	<code>cl_int</code>
<code>unsigned int,</code> <code>uint</code>	<code>cl_uint</code>
<code>long</code>	<code>cl_long</code>
<code>unsigned long,</code> <code>ulong</code>	<code>cl_ulong</code>
<code>float</code>	<code>cl_float</code>
<code>double</code>	<code>cl_double</code>
<code>half</code>	<code>cl_half</code>
<code>size_t</code>	<code>n/a</code>
<code>ptrdiff_t</code>	<code>n/a</code>
<code>intptr_t</code>	<code>n/a</code>
<code>uintptr_t</code>	<code>n/a</code>
<code>void</code>	<code>void</code>

1.1.1.1 The half data type

The `half` data type must be IEEE 754-2008 compliant. `half` numbers have 1 sign bit, 5 exponent bits, and 10 mantissa bits. The interpretation of the sign, exponent and mantissa is analogous to IEEE 754 floating-point numbers. The exponent bias is 15. The `half` data type must represent finite and normal numbers, denormalized numbers, infinities and NaN. Denormalized numbers for the `half` data type which may be generated when converting a `float` to a `half` using `vstore_half` and converting a `half` to a `float` using `vload_half` cannot be flushed to zero. Conversions from `float` to `half` correctly round the mantissa to 11 bits of precision. Conversions from `half` to `float` are lossless; all `half` numbers are exactly representable as `float` values.

The `half` data type can only be used to declare a pointer to a buffer that contains half values. A few valid examples are given below:

Example 1:

```
void
bar(half *p)
{
    ...
}

kernel void
foo(half *pg)
{
    half *ptr;
    int  offset;
```

```

    ptr = pg + offset;
    bar(ptr);
}

```

Below are some examples that are not valid usage of the `half` type.

```

half a;
half b[100];

half *p = b;
a = *p;    ← not allowed.  must use vload_half function

```

The `half` scalar data type is required to be supported as a data storage format. Vector data load and store functions (described in *section 2.7*) must be supported. Support for basic arithmetic operations and library functions described in *section 2* (except functions in *section 2.7*) that operate on `half` types is optional.

1.1.2 Built-in Vector Data Types

The `bool`, `char`, `unsigned char`, `short`, `unsigned short`, `int`, `unsigned int`, `long`, `unsigned long`, `half`, `float` and `double` vector data types are supported. The vector data type is defined with the type name i.e. `bool`, `char`, `uchar`, `short`, `ushort`, `int`, `uint`, `long`, `ulong`, `half`, `float` or `double` followed by a literal value n that defines the number of elements in the vector. Supported values of n are 2, 3, 4, 8, and 16 for all vector data types.

Table 1.2 describes the list of built-in vector data types.

Type	Description
<code>booln</code>	A vector of n boolean values.
<code>charn</code>	A vector of n 8-bit signed two's complement integer values.
<code>ucharn</code>	A vector of n 8-bit unsigned integer values.
<code>shortn</code>	A vector of n 16-bit signed two's complement integer values.
<code>ushortn</code>	A vector of n 16-bit unsigned integer values.
<code>intn</code>	A vector of n 32-bit signed two's complement integer values.
<code>uintn</code>	A vector of n 32-bit unsigned integer values.
<code>longn</code>	A vector of n 64-bit signed two's complement integer values.
<code>ulongn</code>	A vector of n 64-bit unsigned integer values.
<code>halfn</code>	A vector of n 16-bit floating-point values.
<code>floatn</code>	A vector of n 32-bit floating-point values.
<code>doublen</code>	A vector of n 64-bit floating-point values.

Table 1.2 Built-in Vector Data Types

The built-in vector data types are also declared as appropriate types in the OpenCL API (and header files) that can be used by an application. The following table describes the built-in vector data type in the OpenCL C++ programming language and the corresponding data type available to the application:

Type in OpenCL Language	API type for application
<code>booln</code>	<code>cl_booln</code>
<code>charn</code>	<code>cl_charn</code>
<code>ucharn</code>	<code>cl_ucharn</code>
<code>shortn</code>	<code>cl_shortn</code>
<code>ushortn</code>	<code>cl_ushortn</code>
<code>intn</code>	<code>cl_intn</code>
<code>uintn</code>	<code>cl_uintn</code>
<code>longn</code>	<code>cl_longn</code>
<code>ulongn</code>	<code>cl_ulongn</code>
<code>halfn</code>	<code>cl_halfn</code>
<code>floatn</code>	<code>cl_floatn</code>
<code>doublen</code>	<code>cl_doublen</code>

The `halfn` vector data type is required to be supported as a data storage format. Vector data load and store functions (described in *section 2.7*) must be supported. Support for basic arithmetic operations and library functions described in *section 2* (except functions in *section 2.7*) that operate on half vector types is optional.

Support for the `doublen` vector data type is optional.

1.1.3 Other Built-in Data Types

Table 1.3 describes the list of additional data types supported by OpenCL C++. These are defined in the `opencl` namespace.

Type	Description
<code>image2d</code>	A 2D image. Refer to <i>section 2.12</i> for a detailed description of the built-in functions that use this type.
<code>image3d</code>	A 3D image. Refer to <i>section 2.12</i> for a detailed description of the built-in functions that use this type.
<code>image2d_array</code>	A 2D image array. Refer to <i>section 2.12</i> for a detailed description of the built-in functions that use this type.
<code>image1d</code>	A 1D image. Refer to <i>section 2.12</i> for a detailed description of the built-in functions that use this type.

<code>image1d_buffer</code>	A 1D image created from a buffer object. Refer to <i>section 2.12</i> for a detailed description of the built-in functions that use this type.
<code>image1d_array</code>	A 1D image array. Refer to <i>section 2.12</i> for a detailed description of the built-in functions that use this type.
<code>image2d_depth</code>	A 2D depth image. Refer to <i>section 2.12</i> for a detailed description of the built-in functions that use this type.
<code>image2d_array_depth</code>	A 2D depth image array. Refer to <i>section 2.12</i> for a detailed description of the built-in functions that use this type.
<code>sampler</code>	A sampler type. Refer to <i>section 2.12</i> for a detailed description the built-in functions that use of this type.
<code>queue</code>	A device command queue. This queue can only be used to enqueue commands from kernels executing on the device.
<code>ndrange</code>	The N-dimensional range over which a kernel executes.
<code>event</code>	A device side event that identifies a command enqueue to a device command queue.
<code>pipe</code>	A device pipe.
<code>program_pipe</code>	A device pipe object declared in program scope.
<code>reserve_id</code>	A reservation ID. This opaque type is used to identify the reservation for reading and writing a pipe. Refer to <i>section 2.14</i> .
<code>mem_type</code>	<p>This is a bitfield and can be 0 or a combination of the following values ORed together:</p> <pre>mem_type_local mem_type_global mem_type_image</pre> <p>These flags are described in detail in <i>section 2.8</i>.</p>

Table 1.3 Other Built-in Data Types

NOTE: The `image2d`, `image3d`, `image2d_array`, `image1d`, `image1d_buffer`, `image1d_array`, `image2d_depth`, `image2d_array_depth` and `sampler` types are only defined if the device supports images.

Derived types (arrays, structs, unions, functions, and pointers), constructed from the built-in data types described in *sections 1.1.1*, *1.1.2* and *1.1.3* are supported, with restrictions described in *section 1.8*.

The following tables describe the other built-in data types in OpenCL described in table 1.3 and the corresponding data type available to the application:

Type in OpenCL C	API type for application
queue	cl_command_queue
event	cl_event

1.1.4 Reserved Data Types

The data type names described in *table 1.4* are reserved and cannot be used by applications as type names. The vector data type names that are defined in *table 1.2*, where n is any value other than 2, 3, 4, 8 and 16, are also reserved.

Type	Description
quad, quadn	A 128-bit floating-point scalar and vector.
complex half, complex halfn	A complex 16-bit floating-point scalar and vector.
imaginary half, imaginary halfn	An imaginary 16-bit floating-point scalar and vector.
complex float, complex floatn	A complex 32-bit floating-point scalar and vector.
imaginary float, imaginary floatn	An imaginary 32-bit floating-point scalar and vector.
complex double, complex doublen	A complex 64-bit floating-point scalar and vector.
imaginary double, imaginary doublen	An imaginary 64-bit floating-point scalar and vector.
complex quad, complex quadn	A complex 128-bit floating-point scalar and vector.
imaginary quad, imaginary quadn	An imaginary 128-bit floating-point scalar and vector.
floatnxm	An $n \times m$ matrix of single precision floating-point values stored in column-major order.
doublenxm	An $n \times m$ matrix of double precision floating-point values stored in column-major order.
long double, long doublen	A floating-point scalar and vector type with at least as much precision and range as a double and no more precision and range than a quad.
long long, long longn	A 128-bit signed integer scalar and vector.
unsigned long long, unsigned long longn	A 128-bit unsigned integer scalar and vector.

Table 1.4 *Reserved Data Types*

1.1.5 Alignment of Types

A data item declared to be a data type in memory is always aligned to the size of the data type in bytes. For example, a `float4` variable will be aligned to a 16-byte boundary, a `char2` variable will be aligned to a 2-byte boundary.

For 3-component vector data types, the size of the data type is $4 * \text{sizeof}(\text{component})$. This means that a 3-component vector data type will be aligned to a $4 * \text{sizeof}(\text{component})$ boundary. The `vload3` and `vstore3` built-in functions can be used to read and write, respectively, 3-component vector data types from an array of packed scalar data type.

A built-in data type that is not a power of two bytes in size must be aligned to the next larger power of two. This rule applies to built-in types only, not structs or unions.

The OpenCL C++ compiler is responsible for aligning data items to the appropriate alignment as required by the data type. For arguments to a `kernel` function declared to be a pointer to a data type, the OpenCL compiler can assume that the pointee is always appropriately aligned as required by the data type. The behavior of an unaligned load or store is undefined, except for the `vloadn`, `vload_halfn`, `vstoren`, and `vstore_halfn` functions defined in *section 2.7*. The vector load functions can read a vector from an address aligned to the element type of the vector. The vector store functions can write a vector to an address aligned to the element type of the vector.

1.1.6 Vectors

In this section we describe supported operations on vector types.

1.1.6.1 Vector Component Access

The components of vector data types with 1 ... 4 components can be addressed as `<vector_data_type>.xyzw` or `<vector_data_type>.rgba`. Vector data types of type `char2`, `uchar2`, `short2`, `ushort2`, `int2`, `uint2`, `long2`, `ulong2`, `half2`, `float2` and `double2` can access `.xy` or `.rg` elements. Vector data types of type `char3`, `uchar3`, `short3`, `ushort3`, `int3`, `uint3`, `long3`, `ulong3`, `half3`, `float3` and `double3` can access `.xyz` or `.rgb` elements. Vector data types of type `char4`, `uchar4`, `short4`, `ushort4`, `int4`, `uint4`, `long4`, `ulong4`, `half4`, `float4` and `double4` can access `.xyzw` or `.rgba` elements.

Accessing components beyond those declared for the vector type is an error so, for example:

```
float2 pos;  
pos.x = 1.0f; // is legal  
pos.z = 1.0f; // is illegal  
  
float3 pos;  
pos.z = 1.0f; // is legal  
pos.w = 1.0f; // is illegal
```

The component selection syntax allows multiple components to be selected by appending their names after the period (.).

```
float4 c;  
  
c.xyzw = float4(1.0f, 2.0f, 3.0f, 4.0f);  
c.z = 1.0f;  
c.xy = float2(3.0f, 4.0f);  
c.xyz = float3(3.0f, 4.0f, 5.0f);
```

The component selection syntax also allows components to be permuted or replicated.

```
float4 pos = float4(1.0f, 2.0f, 3.0f, 4.0f);  
  
float4 swiz= pos.wzyx; //swiz=(4.0f, 3.0f, 2.0f, 1.0f)  
  
float4 dup = pos.xxyy; //dup=(1.0f, 1.0f, 2.0f, 2.0f)
```

The component group notation can occur on the left hand side of an expression. To form an l-value, swizzling must be applied to an l-value of vector type, contain no duplicate components, and it results in an l-value of scalar or vector type, depending on number of components specified. Each component must be a supported scalar or vector type.

```
float4 pos = float4(1.0f, 2.0f, 3.0f, 4.0f);  
  
// pos = (5.0f, 2.0f, 3.0f, 6.0f)  
pos.xw = float2(5.0f, 6.0f);  
  
// pos = (8.0f, 2.0f, 3.0f, 7.0f)  
pos.wx = float2(7.0f, 8.0f);  
  
// pos = (3.0f, 5.0f, 9.0f, 4.0f)
```

Provisional Specification

```
pos.xyz = float3(3.0f, 5.0f, 9.0f);

// illegal - 'x' used twice
pos.xx = float2(3.0f, 4.0f);

// illegal - mismatch between float2 and float4
pos.xy = float4(1.0f, 2.0f, 3.0f, 4.0f);

float4  a, b, c, d;
float16 x;
x = float16(a, b, c, d);
x = float16(a.xxxx, b.xyz, c.xyz, d.xyz, a.yzw);

// illegal - component a.xxxxxxx
//           is not a valid vector type
x = float16(a.xxxxxxx, b.xyz, c.xyz, d.xyz);
```

Elements of vector data types can also be accessed using a numeric index to refer to the appropriate element in the vector. The numeric indices that can be used are given in the table below:

Vector Components	Numeric indices that can be used
2-component	0, 1
3-component	0, 1, 2
4-component	0, 1, 2, 3
8-component	0, 1, 2, 3, 4, 5, 6, 7
16-component	0, 1, 2, 3, 4, 5, 6, 7, 8, 9, a, A, b, B, c, C, d, D, e, E, f, F

Table 1.5 *Numeric indices for built-in vector data types*

The numeric indices must be preceded by the letter *s* or *S*.

In the following example

```
float8      f;
```

`f.s0` refers to the 1st element of the `float8` variable `f` and `f.s7` refers to the 8th element of the `float8` variable `f`.

In the following example

```
float16     x;
```

Provisional Specification

`x.sa` (or `x.sA`) refers to the 11th element of the `float16` variable `x` and `x.sf` (or `x.sF`) refers to the 16th element of the `float16` variable `x`.

The numeric indices used to refer to an appropriate element in the vector cannot be intermixed with `.xyzw` notation used to access elements of a 1 .. 4 component vector.

For example

```
float4      f, a;

a = f.x12w; // illegal use of numeric indices with
            // .xyzw

a.xyzw = f.s0123;    // valid
```

Vector data types can use the `.lo` (or `.even`) and `.hi` (or `.odd`) suffixes to get smaller vector types or to combine smaller vector types to a larger vector type. Multiple levels of `.lo` (or `.even`) and `.hi` (or `.odd`) suffixes can be used until they refer to a scalar term.

The `.lo` suffix refers to the lower half of a given vector. The `.hi` suffix refers to the upper half of a given vector.

The `.even` suffix refers to the even elements of a vector. The `.odd` suffix refers to the odd elements of a vector.

Some examples to help illustrate this are given below:

```
float4      vf;

float2      low = vf.lo;    // returns vf.xy
float2      high = vf.hi;   // returns vf.zw

float2      even = vf.even; // returns vf.xz
float2      odd = vf.odd;   // returns vf.yw
```

The suffixes `.lo` (or `.even`) and `.hi` (or `.odd`) for a 3-component vector type operate as if the 3-component vector type is a 4-component vector type with the value in the `w` component undefined.

Some examples are given below:

```
float8      vf;
float4      odd = vf.odd;
```

Provisional Specification

```
float4    even = vf.even;
float2    high = vf.even.hi;
float2    low  = vf.odd.lo;

// interleave L+R stereo stream
float4    left,  right;
float8    interleaved;
interleaved.even = left;
interleaved.odd  = right;

// deinterleave
left = interleaved.even;
right = interleaved.odd;

// transpose a 4x4 matrix
void transpose( float4 m[4] )
{
    // read matrix into a float16 vector
    float16 x = float16( m[0], m[1], m[2], m[3] );

    float16 t;

    //transpose
    t.even = x.lo;
    t.odd  = x.hi;
    x.even = t.lo;
    x.odd  = t.hi;

    //write back
    // { m[0][0], m[1][0], m[2][0], m[3][0] }
    m[0] = x.lo.lo;

    // { m[0][1], m[1][1], m[2][1], m[3][1] }
    m[1] = x.lo.hi;

    // { m[0][2], m[1][2], m[2][2], m[3][2] }
    m[2] = x.hi.lo;

    // { m[0][3], m[1][3], m[2][3], m[3][3] }
    m[3] = x.hi.hi;
}

float3    vf = float3(1.0f, 2.0f, 3.0f);
float2    low = vf.lo;  // (1.0f, 2.0f);
float2    high = vf.hi; // (3.0f, undefined);
```


It is an error to take the address of a vector element and will result in a compilation error. For example:

```
float8      vf;

float       *f = &vf.x;           // is illegal
float2      *f2 = &vf.s07;        // is illegal

float4      *odd = &vf.odd;        // is illegal
float4      *even = &vf.even;      // is illegal
float2      *high = &vf.even.hi;   // is illegal
float2      *low = &vf.odd.lo;     // is illegal
```

1.1.6.2 Vector Constructors

Vector constructors are defined to initialize a vector data type from a list of scalar or vectors. The forms of the constructors that are available is the set of possible argument lists for which all arguments have the same element type as the result vector, and the total number of elements is equal to the number of elements in the result vector. In addition, a form with a single scalar of the same type as the element type of the vector is available. For example, the following forms are available for float4:

```
float4( float, float, float, float )
float4( float2, float, float )
float4( float, float2, float )
float4( float, float, float2 )
float4( float2, float2 )
float4( float3, float )
float4( float, float3 )

float4( float )
```

Operands are evaluated by standard rules for function evaluation, except that implicit scalar widening shall not occur. The order in which the operands are evaluated is undefined. The operands are assigned to their respective positions in the result vector as they appear in memory order. That is, the first element of the first operand is assigned to result.x, the second element of the first operand (or the first element of the second operand if the first operand was a scalar) is assigned to result.y, etc. In the case of the form that has a single scalar operand, the operand is replicated across all lanes of the vector.

Examples:

```
float4  f = float4(1.0f, 2.0f, 3.0f, 4.0f);
```

```

uint4    u = uint4(1); ← u will be (1, 1, 1, 1).

float4    f = float4(float2(1.0f, 2.0f),
                     float2(3.0f, 4.0f));

float4    f = float4(1.0f, float2(2.0f, 3.0f), 4.0f);

float4    f = float4(1.0f, 2.0f); ← error

```

1.1.6.3 Vector Types and Usual Arithmetic Conversions


Many operators that expect operands of arithmetic type cause conversions and yield result types in a similar way. The purpose is to determine a common real type for the operands and result. For the specified operands, each operand is converted, without change of type domain, to a type whose corresponding real type is the common real type. For this purpose, all vector types shall be considered to have higher conversion ranks than scalars. Unless explicitly stated otherwise, the common real type is also the corresponding real type of the result, whose type domain is the type domain of the operands if they are the same, and complex otherwise. This pattern is called the usual arithmetic conversions. If the operands are of more than one vector type, then an error shall occur. Implicit conversions between vector types are not permitted, per *section 1.2.1*.

Otherwise, if there is only a single vector type, and all other operands are scalar types, the scalar types are converted to the type of the vector element, then widened into a new vector containing the same number of elements as the vector, by duplication of the scalar value across the width of the new vector. An error shall occur if any scalar operand has greater rank than the type of the vector element

1.1.7 Keywords

The following names are reserved for use as keywords in OpenCL C++ and shall not be used otherwise.

- ✚ Names reserved as keywords by C++14.
- ✚ OpenCL C++ data types defined in *tables 1.1, 1.2, 1.3 and 1.4*.
- ✚ Function qualifiers: `__kernel` and `kernel`.
- ✚ Access qualifiers: `__read_only`, `read_only`, `__write_only`, `write_only`, `__read_write` and `read_write`.

 uniform.

1.2 Conversions and Type Casting

1.2.1 Implicit Conversions

Implicit conversions between scalar built-in types defined in *table 1.1* (except `void`) are supported. When an implicit conversion is done, it is not just a re-interpretation of the expression's value but a conversion of that value to an equivalent value in the new type. For example, the integer value 5 will be converted to the floating-point value 5.0.

Implicit conversions from a scalar type to a vector type are allowed. In this case, the scalar may be subject to the usual arithmetic conversion to the element type used by the vector. The scalar type is then widened to the vector.

Implicit conversions between built-in vector data types are disallowed.

Implicit conversions for pointer types follow the rules described in the C++ 14 specification.

1.2.2 Explicit Casts

The C++ explicit typecasts for built-in scalar data types defined in *table 1.1* will perform appropriate conversion (except `void`). In the example below:

```
float    f = 1.0f;
int      i = (int)f;
int      j = static_cast<int>(f);
```

`f` stores 0x3F800000 and `i` and `j` store 0x1 which is the floating-point value 1.0f in `f` converted to an integer value.

Explicit casts between vector types are not legal. The examples below will generate a compilation error.

```
int4      i;
uint4     u = (uint4) i;    ← not allowed

float4    f;
int4      i = static_cast<int4> f;    ← not allowed

float4    f;
```

Provisional Specification

```
int8      i = (int8) f;  ← not allowed
```

Scalar to vector conversions may be performed by casting the scalar to the desired vector data type. Type casting will also perform appropriate arithmetic conversion. The round to zero rounding mode will be used for conversions to built-in integer vector types. The default rounding mode will be used for conversions to floating-point vector types. When casting a `bool` to a vector integer data type, the vector components will be set to -1 (i.e. all bits set) if the `bool` value is *true* and 0 otherwise.

Below are some correct examples of explicit casts.

```
float      f = 1.0f;
// va is a float4 vector with elements (f, f, f, f).
float4 va = (float4)f;

uchar      u = 0xFF;

// vb is a float4 vector with elements
//      ((float)u, (float)u, (float)u, (float)u).
float4 vb = static_cast<float4>u;

float      f = 2.0f;

// vc is an int2 vector with elements
//      ((int)f, (int)f).
int2 vc = (int2)f;

// vtrue is a uchar4 vector with elements
//      (0xff, 0xff, 0xff, 0xff).
uchar4 vtrue = static_cast<uchar4>true;
```

1.2.3 Explicit Conversions

Explicit type conversions for scalar and vector types may be performed using the

```
convert_cast<T>
```

type conversion operator. This operator provide a full set of type conversions between supported scalar and vector data types (see *sections 1.1.1* and *1.1.2*) except for the following types: `size_t`, `ptrdiff_t`, `intptr_t`, `uintptr_t`, and `void`.

The number of elements in the source and destination vectors must match.

In the example below:

```
uchar4  u;
int4    c = convert_cast<int4>(u);
```

`convert_cast<int4>` converts a `uchar4` vector `u` to an `int4` vector `c`.

```
float    f;
int      i = convert_cast<int>(f);
```

`convert_cast<int>` converts a `float` scalar `f` to an `int` scalar `i`.

The behavior of the conversion may be modified by one or two optional modifiers that specify saturation for out-of-range inputs and rounding behavior.

The `convert_cast<T>` type conversion operator that specifies a rounding mode and saturation is also provided. This operator is defined as:

```
enum class roundingmode { rte, rtz, rtp, rtn };
enum class saturate { off, on };
convert_cast<T, roundingmode>
convert_cast<T, saturate>
convert_cast<T, roundingmode, saturate>
```

1.2.3.1 Data Types

Conversions are available for the following scalar types: `bool`, `char`, `uchar`, `short`, `ushort`, `int`, `uint`, `long`, `ulong`, `half`, `float`, `double`, and built-in vector types derived therefrom. The operand and result type must have the same number of elements. The operand and result type may be the same type in which case the conversion has no effect on the type or value of an expression.

Conversions between integer types follow the conversion rules specified in the C++14 specification except for out-of-range behavior and saturated conversions which are described in *section 1.2.3.3* below.

1.2.3.2 Rounding Modes

Conversions to and from floating-point type shall conform to IEEE-754 rounding rules. Conversions may have an optional rounding mode specified as described in *table 1.6*.

Rounding Mode	Description
<code>rte</code>	Round to nearest even
<code>rtz</code>	Round toward zero
<code>rtp</code>	Round toward positive infinity
<code>rtn</code>	Round toward negative infinity

Table 1.6 *Rounding Modes*

If a rounding mode is not specified, conversions to integer type use the `rtz` (round toward zero) rounding mode and conversions to floating-point type³ uses the `rte` rounding mode.

1.2.3.3 Out-of-Range Behavior and Saturated Conversions

When the conversion operand is either greater than the greatest representable destination value or less than the least representable destination value, it is said to be out-of-range. The result of out-of-range conversion is determined by the conversion rules specified by the C++14 specification in *section 4.9*. When converting from a floating-point type to integer type, the behavior is implementation-defined.

Conversions to integer type may opt to convert using the optional saturation mode. When in saturated mode, values that are outside the representable range shall clamp to the nearest representable value in the destination format. (NaN should be converted to 0).

Conversions to floating-point type shall conform to IEEE-754 rounding rules. The `convert_cast` operator with a `saturate` argument may not be used for conversions to floating-point formats.

1.2.3.4 Explicit Conversion Examples

Example 1:

```
short4 s;

// negative values clamped to 0
ushort4 u = convert_cast<ushort4,saturate::on>( s );
```

³ For conversions to floating-point format, when a finite source value exceeds the maximum representable finite floating-point destination value, the rounding mode will affect whether the result is the maximum finite floating-point value or infinity of same sign as the source value, per IEEE-754 rules for rounding.

```
// values > CHAR_MAX converted to CHAR_MAX
// values < CHAR_MIN converted to CHAR_MIN
char4 c = convert_cast<char4, saturate::on>( s );
```

Example 2:

```
float4 f;

// values implementation defined for
// f > INT_MAX, f < INT_MIN or NaN
int4 i = convert_cast<int4>( f );

// values > INT_MAX clamp to INT_MAX, values < INT_MIN clamp
// to INT_MIN. NaN should produce 0.
// The rtz rounding mode is used to produce the integer
// values.
int4 i2 = convert_cast<int4,saturate::on>( f );

// similar to convert_cast<int4>, except that floating-point
// values are rounded to the nearest integer instead of
// truncated
int4 i3 = convert_cast<int4, roundingmode::rte>( f );

// similar to convert_cast<int4, saturate::on>, except that
// floating-point values are rounded to the nearest integer
// instead of truncated
int4 i4 = convert_cast<int4, saturate::on,
                      roundingmode::rte>( f );
```

Example 3:

```
int4 i;

// convert ints to floats using the default rounding mode.
float4 f = convert_cast<float4>( i );

// convert ints to floats. integer values that cannot
// be exactly represented as floats should round up to the
// next representable float.
float4 f = convert_cast<float4, roundingmode::rtp>( i );
```

1.2.4 Reinterpreting Data As Another Type

It is frequently necessary to reinterpret bits in a data type as another data type in OpenCL C++. This is typically required when direct access to the bits in a floating-point type is needed, for example to mask off the sign bit or make use of the result of a vector relational operator (see *section 1.3, item d*) on floating-point data.

1.2.4.1 Reinterpreting Types Using Unions

The OpenCL C++ language extends the union to allow the program to access a member of a union object using a member of a different type. The relevant bytes of the representation of the object are treated as an object of the type used for the access. If the type used for access is larger than the representation of the object, then the value of the additional bytes is undefined.

Examples:

```
union{ float f; uint u; double d4; } u;

u.u = 1;      // u.f contains 2**-149.  u.d is undefined -
              // depending on endianness the low or high half
              // of d is unknown

u.f = 1.0f;   // u.u contains 0x3f800000, u.d contains an
              // undefined value -- depending on endianness
              // the low or high half of d is unknown

u.d = 1.0;    // u.u contains 0x3ff00000 (big endian) or 0
              // (little endian). u.f contains either
              // 0x1.ep0f(big endian) or 0.0f(little endian)
```

1.2.4.2 Reinterpreting Types Using `as_type<T>`

All data types described in tables 1.1 and 1.2 (except `bool` and `void`) may be also reinterpreted as another data type of the same size using the **`as_type()`**⁵ operator for scalar and vector data types. When the operand and result type contain the same number of elements, the bits in the operand shall be returned directly without modification as the new type. The usual type promotion for function arguments shall not be performed.

⁴ Only if double precision is supported.

⁵ While the union is intended to reflect the organization of data in memory, the **`as_type<T>`** operator is intended to reflect the organization of data in register. The **`as_type<T>`** construct is intended to compile to no instructions on devices that use a shared register file designed to operate on both the operand and result types. Note that while differences in memory organization are expected to largely be limited to those arising from endianness, the register based representation may also differ due to size of the element in register. (For example, an architecture may load a char into a 32-bit register, or a char vector into a SIMD vector register with fixed 32-bit element size.) If the element count does not match, then the implementation should pick a data representation that most closely matches what would happen if an appropriate result type operator was applied to a register containing data of the source type. If the number of elements matches, then the `as_type()` should faithfully reproduce the behavior expected from a similar data type reinterpretation using memory/unions. So, for example if an implementation stores all single precision data as double in register, it should implement `as_int(float)` by first downconverting the double to single precision and then (if necessary) moving the single precision bits to a register suitable for operating on integer data. If data stored in different address spaces do not have the same endianness, then the "dominant endianness" of the device should prevail.

For example, `as_type<float>(0x3f800000)` returns `1.0f`, which is the value that the bit pattern `0x3f800000` has if viewed as an IEEE-754 single precision value.

When the operand and result type contain a different number of elements, the result shall be implementation-defined except if the operand is a 4-component vector and the result is a 3-component vector. In this case, the bits in the operand shall be returned directly without modification as the new type. That is, a conforming implementation shall explicitly define a behavior, but two conforming implementations need not have the same behavior when the number of elements in the result and operand types does not match. The implementation may define the result to contain all, some or none of the original bits in whatever order it chooses. It is an error to use the `as_type<T>` operator to reinterpret data to a type of a different number of bytes.

Examples:

```
float f = 1.0f;
uint u = as_type<uint>(f); // Legal. Contains: 0x3f800000

float4 f = (float4)(1.0f, 2.0f, 3.0f, 4.0f);
// Legal. Contains:
// (int4)(0x3f800000, 0x40000000, 0x40400000, 0x40800000)
int4 i = as_type<int4>(f);

float4 f, g;
int4 is_less = f < g;

// Legal. f[i] = f[i] < g[i] ? f[i] : 0.0f
f = as_type<float4>(as_type<int4>(f) & is_less);

int i;
// Legal. Result is implementation-defined.
short2 j = as_type<short2>(i);

int4 i;
// Legal. Result is implementation-defined.
short8 j = as_type<short8>(i);




float4 f;
// Error. Result and operand have different sizes
double4 g = as_type<double4>6(f);

float4 f;
// Legal. g.xyz will have same values as f.xyz. g.w is
// undefined
float3 g = as_type<float3>(f);
```

⁶ Only if double precision is supported.

1.3 Operators

- a. The arithmetic operators add (+), subtract (-), multiply (*) and divide (/) operate on built-in integer and floating-point scalar, and vector data types. The remainder (%) operates on built-in integer scalar and integer vector data types. All arithmetic operators return result of the same built-in type (integer or floating-point) as the type of the operands, after operand type conversion. After conversion, the following cases are valid:

-  The two operands are scalars. In this case, the operation is applied, resulting in a scalar.
-  One operand is a scalar, and the other is a vector. In this case, the scalar may be subject to the usual arithmetic conversion to the element type used by the vector operand. The scalar type is then widened to a vector that has the same number of components as the vector operand. The operation is done component-wise resulting in the same size vector.
-  The two operands are vectors of the same type. In this case, the operation is done component-wise resulting in the same size vector.

All other cases of implicit conversions are illegal. Division on integer types which results in a value that lies outside of the range bounded by the maximum and minimum representable values of the integer type will not cause an exception but will result in an unspecified value. A divide by zero with integer types does not cause an exception but will result in an unspecified value. Division by zero for floating-point types will result in \pm infinity or NaN as prescribed by the IEEE-754 standard. Use the built-in functions **dot** and **cross** to get, respectively, the vector dot product and the vector cross product.

- b. The arithmetic unary operators (+ and -) operate on built-in scalar and vector types.
- c. The arithmetic post- and pre-increment and decrement operators (-- and ++) operate on built-in scalar and vector types except the built-in scalar and vector float types⁷. All unary operators work component-wise on their operands. These result with the same type they operated on. For post- and pre-increment and decrement, the expression must be one that could be assigned to (an l-

⁷ The pre- and post- increment operators may have unexpected behavior on floating-point values and are therefore not supported for floating-point scalar and vector built-in types. For example, if variable *a* has type float and holds the value 0x1.0p25f, then *a*++ returns 0x1.0p25f. Also, (*a*++)-- is not guaranteed to return *a*, if *a* has fractional value. In non-default rounding modes, (*a*++)-- may produce the same result as *a*++ or *a*-- for large *a*.

value). Pre-increment and pre-decrement add or subtract 1 to the contents of the expression they operate on, and the value of the pre-increment or pre-decrement expression is the resulting value of that modification. Post-increment and post-decrement expressions add or subtract 1 to the contents of the expression they operate on, but the resulting expression has the expression's value before the post-increment or post-decrement was executed.

- d. The relational operators⁸ greater than ($>$), less than ($<$), greater than or equal ($>=$), and less than or equal ($<=$) operate on scalar and vector types. All relational operators result in a boolean (scalar or vector) type. After operand type conversion, the following cases are valid:

- ✚ The two operands are scalars. In this case, the operation is applied, resulting in a `bool` scalar.
- ✚ One operand is a scalar, and the other is a vector. In this case, the scalar may be subject to the usual arithmetic conversion to the element type used by the vector operand. The scalar type is then widened to a vector that has the same number of components as the vector operand. The operation is done component-wise resulting in the same size vector.
- ✚ The two operands are vectors of the same type. In this case, the operation is done component-wise resulting in the same size vector.

All other cases of implicit conversions are illegal.

The result is a scalar boolean of type `bool` if the source operands are scalar and a vector boolean type of the same size as the source operands if one or both of the source operands are vector types. The relational operators shall return `false` if the specified relation is *false* and `true` if the specified relation is *true*. The relational operations always return `false` if either operand is not a number (NaN).

- e. The equality operators⁹ equal (`==`), and not equal (`!=`) operate on built-in scalar and vector types. All equality operators result in a boolean (scalar or vector) type. After operand type conversion, the following cases are valid:

- ✚ The two operands are scalars. In this case, the operation is applied, resulting in a `bool` scalar.

⁸ To test whether any or all elements in the result of a vector relational operator test true, for example to use in the context in an `if ()` statement, please see the **any** and **all** builtins in *section 6.11.6*.

⁹ To test whether any or all elements in the result of a vector equality operator test true, for example to use in the context in an `if ()` statement, please see the **any** and **all** builtins in *section 6.11.6*.

- ✚ One operand is a scalar, and the other is a vector. In this case, the scalar may be subject to the usual arithmetic conversion to the element type used by the vector operand. The scalar type is then widened to a vector that has the same number of components as the vector operand. The operation is done component-wise resulting in the same size vector.
- ✚ The two operands are vectors of the same type. In this case, the operation is done component-wise resulting in the same size vector.

All other cases of implicit conversions are illegal.

The result is a scalar boolean of type `bool` if the source operands are scalar and a vector boolean type of the same size as the source operands if one or both of the source operands are vector types. The equality operator `equal (==)` returns `false` if one or both arguments are not a number (NaN). The equality operator `not equal (!=)` returns `true` if one or both operands are not a number (NaN).

- f. The bitwise operators `and (&)`, `or (|)`, `exclusive or (^)`, `not (~)` operate on all scalar and vector built-in types except the built-in scalar and vector float types. For vector built-in types, the operators are applied component-wise. If one operand is a scalar and the other is a vector, the scalar may be subject to the usual arithmetic conversion to the element type used by the vector operand. The scalar type is then widened to a vector that has the same number of components as the vector operand. The operation is done component-wise resulting in the same size vector.
- g. The logical operators `and (&&)`, `or (||)` operate on all scalar and vector built-in types. For scalar built-in types only, `and (&&)` will only evaluate the right hand operand if the left hand operand compares unequal to 0. For scalar built-in types only, `or (||)` will only evaluate the right hand operand if the left hand operand compares equal to 0. For built-in vector types, both operands are evaluated and the operators are applied component-wise. If one operand is a scalar and the other is a vector, the scalar may be subject to the usual arithmetic conversion to the element type used by the vector operand. The scalar type is then widened to a vector that has the same number of components as the vector operand. The operation is done component-wise resulting in the same size vector.

The logical operator `exclusive or (^)` is reserved.

The result is a scalar or vector boolean.

- h. The logical unary operator `not (!)` operates on all scalar and vector built-in types. For built-in vector types, the operators are applied component-wise.

The result is a scalar or vector boolean.

- i. The ternary selection operator (`? :`) operates on three expressions (`exp1 ? exp2 : exp3`). This operator evaluates the first expression `exp1`, which must result in a scalar boolean. If the result is true it selects to evaluate the second expression, otherwise it selects to evaluate the third expression. The second and third expressions can be any type, as long their types match, or there is a conversion in *section 1.2.1 Implicit Conversions* that can be applied to one of the expressions to make their types match, or one is a vector and the other is a scalar and the scalar may be subject to the usual arithmetic conversion to the element type used by the vector operand and widened to the same type as the vector type. This resulting matching type is the type of the entire expression.
- j. The operators right-shift (`>>`), left-shift (`<<`) operate on all scalar and vector built-in types except the built-in scalar and vector float types. For built-in vector types, the operators are applied component-wise. For the right-shift (`>>`), left-shift (`<<`) operators, the rightmost operand must be a scalar if the first operand is a scalar, and the rightmost operand can be a vector or scalar if the first operand is a vector.

The result of `E1 << E2` is `E1` left-shifted by $\log_2(N)$ least significant bits in `E2` viewed as an unsigned integer value, where `N` is the number of bits used to represent the data type of `E1` after integer promotion¹⁰, if `E1` is a scalar, or the number of bits used to represent the type of `E1` elements, if `E1` is a vector. The vacated bits are filled with zeros.

The result of `E1 >> E2` is `E1` right-shifted by $\log_2(N)$ least significant bits in `E2` viewed as an unsigned integer value, where `N` is the number of bits used to represent the data type of `E1` after integer promotion, if `E1` is a scalar, or the number of bits used to represent the type of `E1` elements, if `E1` is a vector. If `E1` has an unsigned type or if `E1` has a signed type and a nonnegative value, the vacated bits are filled with zeros. If `E1` has a signed type and a negative value, the vacated bits are filled with ones.

- k. The `sizeof` operator yields the size (in bytes) of its operand, including any padding bytes (refer to *section 1.1.5*) needed for alignment, which may be an expression or the parenthesized name of a type. The size is determined from the type of the operand. The result is of type `size_t`. If the type of the operand is a variable length array¹¹ type, the operand is evaluated; otherwise, the operand is not evaluated and the result is an integer constant.

¹⁰ Integer promotion is described in ISO/IEC 9899:1999 in *section 6.3.1.1*.

¹¹ Variable length arrays are not supported in OpenCL 2.1. Refer to *section 1.9.d*.

When applied to an operand that has type `char`, `uchar`, the result is 1. When applied to an operand that has type `short`, `ushort`, or `half` the result is 2. When applied to an operand that has type `int`, `uint` or `float`, the result is 4. When applied to an operand that has type `long`, `ulong` or `double`, the result is 8. When applied to an operand that is a vector type, the result¹² is number of components * size of each scalar component. When applied to an operand that has array type, the result is the total number of bytes in the array. When applied to an operand that has structure or union type, the result is the total number of bytes in such an object, including internal and trailing padding. The `sizeof` operator shall not be applied to an expression that has function type or an incomplete type, to the parenthesized name of such a type, or to an expression that designates a bit-field struct member¹³.

The behavior of applying the `sizeof` operator to the `bool` or data types defined in table 1.3 is implementation-defined.

- l. The comma (,) operator operates on expressions by returning the type and value of the right-most expression in a comma separated list of expressions. All expressions are evaluated, in order, from left to right.
- m. The unary (*) operator denotes indirection. If the operand points to an object, the result is an lvalue designating the object. If the operand has type “pointer to *type*”, the result has type “*type*”. If an invalid value has been assigned to the pointer, the behavior of the unary * operator is undefined¹⁴.
- n. The unary (&) operator returns the address of its operand. If the operand has type “*type*”, the result has type “pointer to *type*”. If the operand is the result of a unary * operator, neither that operator nor the & operator is evaluated and the result is as if both were omitted, except that the constraints on the operators still apply and the result is not an lvalue. Similarly, if the operand is the result of a [] operator, neither the & operator nor the unary * that is implied by the [] is evaluated and the result is as if the & operator were removed and the [] operator were changed to a + operator. Otherwise, the result is a pointer to the object designated by its operand¹⁵.
- o. Assignments of values to variable names are done with the assignment operator

¹² Except for 3-component vectors whose size is defined as 4 * size of each scalar component.

¹³ Bit-field struct members are not supported in OpenCL 2.1. Refer to *section 1.9.c*.

¹⁴ Among the invalid values for dereferencing a pointer by the unary * operator are a null pointer, an address inappropriately aligned for the type of object pointed to, and the address of an object after the end of its lifetime. If ***P** is an lvalue and **T** is the name of an object pointer type, ***(T)P** is an lvalue that has a type compatible with that to which **T** points.

¹⁵ Thus, **&*E** is equivalent to **E** (even if **E** is a null pointer), and **&(E1[E2])** to **((E1)+(E2))**. It is always true that if **E** is an lvalue that is a valid operand of the unary & operator, ***&E** is an lvalue equal to **E**.

(=), like

```
lvalue = expression
```

The assignment operator stores the value of *expression* into *lvalue*. The *expression* and *lvalue* must have the same type, or the expression must have a type in *table 6.1*, in which case an implicit conversion will be done on the expression before the assignment is done.

If *expression* is a scalar type and *lvalue* is a vector type, the scalar is converted to the element type used by the vector operand. The scalar type is then widened to a vector that has the same number of components as the vector operand. The operation is done component-wise resulting in the same size vector.

Any other desired type-conversions must be specified explicitly. L-values must be writable. Variables that are built-in types, entire structures or arrays, structure fields, l-values with the field selector (.) applied to select components or swizzles without repeated fields, l-values within parentheses, and l-values dereferenced with the array subscript operator ([]) are all l-values. Other binary or unary expressions, function names, swizzles with repeated fields, and constants cannot be l-values. The ternary operator (?:) is also not allowed as an l-value.

The order of evaluation of the operands is unspecified. If an attempt is made to modify the result of an assignment operator or to access it after the next sequence point, the behavior is undefined. Other assignment operators are the assignments add into (+=), subtract from (-=), multiply into (*=), divide into (/=), modulus into (%=), left shift by (<<=), right shift by (>>=), and into (&=), inclusive or into (|=), and exclusive or into (^=).

The expression

```
lvalue op= expression
```

is equivalent to

```
lvalue = lvalue op expression
```

and the l-value and expression must satisfy the requirements for both operator *op* and assignment (=).

1.4 Vector Operations

Vector operations are component-wise. Usually, when an operator operates on a vector, it is operating independently on each component of the vector, in a

component-wise fashion.

For example,

```
float4    v, u;
float     f;

v = u + f;
```

will be equivalent to

```
v.x = u.x + f;
v.y = u.y + f;
v.z = u.z + f;
v.w = u.w + f;
```

And

```
float4    v, u, w;

w = v + u;
```

will be equivalent to

```
w.x = v.x + u.x;
w.y = v.y + u.y;
w.z = v.z + u.z;
w.w = v.w + u.w;
```

and likewise for most operators and all integer and floating-point vector types.

1.5 Address Spaces

Unlike OpenCL C, OpenCL C++ does not require an address space to be specified for objects allocated in global, constant, local or private memory. The generic address space is used instead.

Arguments to a kernel function that are allocated in local memory must be declared as one of the following templated type:

```
local_ptr<T>
local_array<T, size_t N>
```

Variables allocated in local memory inside a kernel function must be declared as one of the following templated type:


```
local_array<T, size_t N>
local<T>
```

***** Question *** - Should we allow `local_ptr<T>` type for variables declared inside a kernel function?**

The `local_ptr<T>` should be used to declare variables of type `T` which are allocated in local memory at runtime. The `local_array<T, size_t N>` describes a fixed size array of type `T` allocated in local memory. The `local<T>` describes a single variable of type `T` allocated in local memory.

The following member functions are available to `local_ptr<T>` and `local_array<T, size_t N>` templated types.

Function	Description
reference operator [] (size_type pos);	Returns a reference to the element at specified location <code>pos</code> . No bounds checking is performed.
const_reference operator [] (size_type pos) const;	
T *data() ; const T *data() const;	Returns a pointer to the underlying storage. The pointer refers to the allocation of the underlying storage in the generic address space.
size_type size() ; size_type size() const;	Returns the size of the underlying storage in bytes

***** Question *** - Should we support the iterators `begin`, `cbegin`, `end`, `cend`, `rbegin`, `crbegin`, `rend`, `crend`?**

The following operators and member functions are available to `local<T>` templated type.

Function	Description
operator=	The assignment operator.
constant_ptr<T> data() const;	Returns a <code>constant_ptr<T></code> .

Arguments to a kernel function that are allocated in constant memory must be declared as one of the following templated types:

```
constant_ptr<T>
```

```
constant_array<T, size_t N>
```

Variables allocated in constant memory inside a kernel function or in program scope must be declared as one of the following templated type:

```
constant_array<T, size_t N>
constant<T>
```

***** Question *** - Should we allow `constant_ptr<T>` type for variables declared inside a kernel function?**

The `constant_ptr<T>` should be used to declare variables of type `T` which are allocated in constant memory at runtime. The `constant_array<T, size_t N>` describes a fixed size array of type `T` allocated in constant memory. The `constant<T>` describes a single variable of type `T` allocated in constant memory.

The following member functions are available to `constant_ptr<T>` templated type:

Function	Description
const_reference operator[] (size_type pos) const;	Returns a reference to the element at specified location pos. No bounds checking is performed.
size_type size (); size_type size () const;	Returns the size of the storage in bytes

The following member functions are available to `constant_array<T, size_t N>` templated type:

Function	Description
const_reference operator[] (size_type pos) const;	Returns a reference to the element at specified location pos. No bounds checking is performed.
constant_ptr<T> data () const;	Returns a <code>constant_ptr<T></code> that can be used to reference the underlying storage.
size_type size (); size_type size () const;	Returns the size of the storage in bytes

The following operators and member functions are available to `constant<T>` templated type.

Function	Description
----------	-------------

<code>operator=</code>	The assignment operator.
<code>constant_ptr<T> data() const;</code>	Returns a <code>constant_ptr<T></code> that can be used to reference the underlying storage

Restrictions

1. Variables allocated in local memory inside a kernel function must occur at kernel function scope; otherwise such variables must be declared at program scope or with the `static` qualifier. Some examples of variables allocated in the `local` address space inside a kernel function are:

```
extern local_ptr<float> elf; // OK
local_array<int, 23> sli; // OK
local<int> libad = 42; // Error - no initializer
                        // allowed

kernel void my_func(...)
{
    // A single float allocated in
    // local address space
    local<float> a;

    // An array of 10 floats allocated in
    // local address space.
    local_array<float,10> b;

    if (...)
    {
        // example of variable in local address space
        // but not declared at __kernel function scope.
        local_array<float> c; ← not allowed.
    }
}
```

2. Variables allocated in the `local` address space inside a kernel function or in program scope cannot be initialized.

```
kernel void my_func(...)
{
    local<float> a = 1; ← not allowed

    local<float> b;
    b = 1; ← allowed
}
```

3. It is not legal to take the address of an element stored within a `constant_array`, `constant`, or through a `constant_ptr` object. Constant locations are not within the generic address space and cannot be assigned to pointer variables.

NOTE: Variables allocated in local memory inside a kernel function are allocated for each work-group executing the kernel and exist only for the lifetime of the work-group executing the kernel.

1.6 Function Qualifiers

The **kernel** (or **__kernel**) qualifier declares a function to be a kernel that can be executed by an application on an OpenCL device(s). The following rules apply to functions that are declared with this qualifier:

- ✚ It can be executed on the device only
- ✚ It can be enqueued by the host or on the device

The `kernel` and `__kernel` names are reserved for use as functions qualifiers and shall not be used otherwise.

Restrictions

- ✚ A `kernel` function **cannot** be called by another kernel function.
- ✚ Templated functions cannot be declared with the `kernel` qualifier.
- ✚ Kernel functions are by default declared as extern C.
- ✚ Arguments of a kernel function cannot be specified with default values. [***
Question – should we support default values? ***]

1.7 Storage-Class Specifiers

The `static`, `extern` and `mutable` storage class specifiers (as defined in *section 7.1.1* of the C++14 specification) are supported.

The `extern` storage-class specifier can only be used for functions (kernel and non-kernel functions) and global variables declared in program scope or variables declared inside a function (kernel and non-kernel functions). The `static` storage-class specifier can only be used for non-kernel functions, global variables declared in program scope and variables inside a function.

Examples:

```

#include <opencl_stdlib>
using namespace cl;

extern float4 noise_table[256];
static float4 color_table[256];

extern kernel void my_foo(image2d<float4> img);
extern void my_bar(float *a);

kernel void my_func(image2d<float4> img, float *a)
{
    extern float4 a;
    // "a" and "b" will be allocated in
    // the global address space
    static float4 b = float4(1.0f); // OK.
    static float c; // OK
    ...
    my_foo(img);
    ...
    my_bar(a);
    ...
    while (1)
    {
        static int inside; // OK.
        ...
    }
    ...
}

```

1.8 Restrictions

The following C++14 features are not supported by OpenCL C++

- The `dynamic_cast` operator (*section 5.2.7*)
- Type identification (*section 5.2.8*)
- Recursive function calls (*section 5.2.2, item 9*)
- `new` and `delete` operators (*sections 5.3.4, 5.3.5*)
- `noexcept` operator (*section 5.3.7*)
- `goto` statement (*section 6.6*)
- `register`, `thread_local` storage qualifiers (*section 7.1.1*)
- virtual function qualifier (*section 7.1.2*)
- function pointers (*sections 8.3.5, 8.5.3*)
- virtual functions and abstract classes (*sections 10.3, 10.4*)
- exception handling (*section 15*)
- the C++ standard library (*sections 17 ... 30*)

1.9 Pre-processor Directives and Macros

The preprocessing directives defined by the C++14 specification (*section 16*) are supported.

The **# pragma** directive is described as:

pragma *pp-tokens_{opt} new-line*

A **# pragma** directive where the preprocessing token **OPENCL** (used instead of **STDC**) does not immediately follow **pragma** in the directive (prior to any macro replacement) causes the implementation to behave in an implementation-defined manner. The behavior might cause translation to fail or cause the translator or the resulting program to behave in a non-conforming manner. Any such **pragma** that is not recognized by the implementation is ignored. If the preprocessing token **OPENCL** does immediately follow **pragma** in the directive (prior to any macro replacement), then no macro replacement is performed on the directive, and the directive shall have one of the following forms whose meanings are described elsewhere:

#pragma OPENCL FP_CONTRACT *on-off-switch*
on-off-switch: one of **ON OFF DEFAULT**

#pragma OPENCL EXTENSION *extensionname* : *behavior*
#pragma OPENCL EXTENSION *all* : *behavior*

The following predefined macro names are available.

__FILE__ The presumed name of the current source file (a character string literal).

__LINE__ The presumed line number (within the current source file) of the current source line (an integer constant).

__OPENCL_VERSION__ substitutes an integer reflecting the version number of the OpenCL supported by the OpenCL device. The version of OpenCL described in this document will have **__OPENCL_VERSION__** substitute the integer 210.

CL_VERSION_2_1 substitutes the integer 210 reflecting the OpenCL 2.1 version.

__OPENCL_C++_VERSION__ substitutes an integer reflecting the OpenCL C version specified when compiling the OpenCL C++ program. The version of OpenCL C++ described in this document will have **__OPENCL_C++_VERSION__** substitute

the integer 100.

The macro names defined by the C++14 specification in *section 16* but not currently supported by OpenCL are reserved for future use.

The predefined identifier `__func__` is available.

1.10 Attribute Qualifiers

The following additional attribute qualifiers are supported:

1.10.1 Optional Attribute Qualifiers

The `kernel` qualifier can be used with the `[[]]` attribute syntax to declare additional information about the kernel function. The following attributes are supported:

The optional `[[vec_type_hint(<type>)]]` is a hint to the compiler and is intended to be a representation of the computational *width* of the `kernel`, and should serve as the basis for calculating processor bandwidth utilization when the compiler is looking to autovectorize the code. In the `[[vec_type_hint(<type>)]]` qualifier `<type>` is one of the built-in vector types listed in *table 1.2* or the constituent scalar element types. If `vec_type_hint(<type>)` is not specified, the kernel is assumed to have the `[[vec_type_hint(int)]]` qualifier.

For example, where the developer specified a width of `float4`, the compiler should assume that the computation usually uses up to 4 lanes of a float vector, and would decide to merge work-items or possibly even separate one work-item into many threads to better match the hardware capabilities. A conforming implementation is not required to autovectorize code, but shall support the hint. A compiler may autovectorize, even if no hint is provided. If an implementation merges `N` work-items into one thread, it is responsible for correctly handling cases where the number of global or local work-items in any dimension modulo `N` is not zero.

Examples:

```
// autovectorize assuming float4 as the
// basic computation width
kernel [[vec_type_hint(float4)]]
void foo(float4 *p ) { .... }

// autovectorize assuming double as the
```

```
// basic computation width
kernel [[vec_type_hint(double)]]
void foo(float4 *p ) { .... }

// autovectorize assuming int (default)
// as the basic computation width
kernel
void foo(float4 *p ) { .... }
```

The optional `[[work_group_size_hint(X, Y, Z)]]` is a hint to the compiler and is intended to specify the work-group size that may be used i.e. value most likely to be specified by the *local_work_size* argument to **clEnqueueNDRangeKernel**. For example the `[[work_group_size_hint(1, 1, 1)]]` is a hint to the compiler that the kernel will most likely be executed with a work-group size of 1.

The optional `[[required_work_group_size(X, Y, Z)]]` is the work-group size that must be used as the *local_work_size* argument to **clEnqueueNDRangeKernel**. This allows the compiler to optimize the generated code appropriately for this kernel.

If *Z* is one, the *work_dim* argument to **clEnqueueNDRangeKernel** can be 2 or 3. If *Y* and *Z* are one, the *work_dim* argument to **clEnqueueNDRangeKernel** can be 1, 2 or 3.

The optional `[[nosvm]]` qualifier can be used with a pointer variable to inform the compiler that the pointer does not refer to a shared virtual memory region.

1.10.2 Specifying Attributes for Unrolling Loops

The `[[opencl_unroll_hint]]` and `[[opencl_unroll_hint(n)]]` attribute qualifiers can be used to specify that a loop (*for*, *while* and *do* loops) can be unrolled. This attribute qualifier can be used to specify full unrolling or partial unrolling by a specified amount. This is a compiler hint and the compiler may ignore this directive.

n is the loop unrolling factor and must be a positive integral compile time constant expression. An unroll factor of 1 disables unrolling. If *n* is not specified, the compiler determines the unrolling factor for the loop.

NOTE: The `[[opencl_unroll_hint(n)]]` attribute qualifier must appear immediately before the loop to be affected.

Examples:

Provisional Specification

```
[[opencl_unroll_hint(2)]]
while (*s != 0)
    *p++ = *s++;
```

The tells the compiler to unroll the above while loop by a factor of 2.

```
[[opencl_unroll_hint]]
for (int i=0; i<2; i++)
{
    ...
}
```

In the example above, the compiler will determine how much to unroll the loop.

```
[[opencl_unroll_hint(1)]]
for (int i=0; i<32; i++)
{
    ...
}
```

The above is an example where the loop should not be unrolled.

Below are some examples of invalid usage of `[[opencl_unroll_hint(n)]]`.

```
[[opencl_unroll_hint(-1)]]
while (...)
{
    ...
}
```

The above example is an invalid usage of the loop unroll factor as the loop unroll factor is negative.

```
[[opencl_unroll_hint]]
if (...)
{
    ...
}
```

The above example is invalid because the unroll attribute qualifier is used on a non-loop construct

```
kernel void
my_kernel( ... )
{
    int x;
```

Provisional Specification

```
[[opencl_unroll_hint(x)]]  
for (int i=0; i<x; i++)  
{  
    ...  
}  
}
```

The above example is invalid because the loop unroll factor is not a compile-time constant expression.

The `[[ivdep]]` (ignore vector dependencies) attribute qualifier is a hint to the compiler and may appear in loops to indicate that the compiler may assume there are no memory dependencies across loop iterations in order to autovectorize consecutive iterations of the loop. This attribute qualifier may appear in one of the following forms:

```
[[ivdep]]  
[[ivdep, safelen(len)]]
```

If the optional attribute `safelen` is specified, it is used to specify the maximum number of consecutive iterations without loop-carried dependencies. `safelen` is a lower bound on the distance of any loop-carried dependence, and it applies to arbitrary alignment. For example, any 4 consecutive iterations can be vectorized with `safelen(4)`. The `len` parameter must be a positive integer. `safelen` may only be specified after `ivdep`. The final decision whether to autovectorize the complete loop may be subject to other compiler heuristics as well as flags e.g., `-cl-fast-relaxed-math` to ignore non-associated operations.

Examples:

```
[[ivdep]]  
for (int i=0; i<N; i++)  
{  
    C[i+offset] = A[i+offset] * B[i+offset];  
}
```

In the example above, assuming that A and B are not restricted pointers, it is unknown if C aliases A or B. Placing the `[[ivdep]]` attribute before the loop lets the compiler assume there are no memory dependencies across the loop iterations.

```
[[ivdep, safelen(8)]]  
for (int i=0; i<N; i++)  
{  
    A[i+K] = A[i] * B[i];  
}
```

In the example above, buffer A is read from and written to in the loop iterations. In each iteration, the read and write to A are to different indices. In this case it is not safe to vectorize the loop to a vector length greater than K, so the `safelen` attribute is specified with a parameter value that is known to be not greater than any value that K may take during the execution of loop. In this example we are guaranteed (by `safelen`) that K will always be greater than or equal to 8.

Below are some examples of invalid usage of `[[ivdep]]`.

```
[[ivdep, safelen(-1)]]
for (int i=0; i<N; i++)
{
    C[i+offset] = A[i+offset] * B[i+offset];
}
```

The above example is an invalid usage of the attribute qualifier as `safelen` is negative.

```
[[safelen(8), ivdep]]
for (int i=0; i<N; i++)
{
    C[i+offset] = A[i+offset] * B[i+offset];
}
```

The above example is invalid because the attribute qualifier `safelen` does not appear after `ivdep`.

1.10.3 Extending Attribute Qualifiers

The attribute syntax can be extended for standard language extensions and vendor specific extensions. Any extensions should follow the naming conventions outlined in the introduction to *section 9* in the OpenCL 2.1 Extension Specification.

Attributes are intended as useful hints to the compiler. It is our intention that a particular implementation of OpenCL be free to ignore all attributes and the resulting executable binary will produce the same result. This does not preclude an implementation from making use of the additional information provided by attributes and performing optimizations or other transformations as it sees fit. In this case it is the programmer's responsibility to guarantee that the information provided is in some sense correct.

2 OpenCL C++ Standard Library

OpenCL C++ does not support the C++14 standard library but instead implements its own standard library. All types, keywords and functions defined by OpenCL C++ are in the namespace `cl`. The `opencl_stdlib` header is also defined which includes the headers for various library functions that make the OpenCL C++ standard library.

2.1 Work-Item Functions

Table 2.1 describes the list of built-in work-item functions that can be used to query the number of dimensions, the global and local work size specified to **clEnqueueNDRangeKernel**, and the global and local identifier of each work-item when this kernel is being executed on a device.

Defined in header `<opencl_work_item>`

Function	Description
uint get_work_dim ()	Returns the number of dimensions in use. This is the value given to the <i>work_dim</i> argument specified in clEnqueueNDRangeKernel .
size_t get_global_size (uint <i>dimindx</i>)	Returns the number of global work-items specified for dimension identified by <i>dimindx</i> . This value is given by the <i>global_work_size</i> argument to clEnqueueNDRangeKernel . Valid values of <i>dimindx</i> are 0 to get_work_dim () – 1. For other values of <i>dimindx</i> , get_global_size () returns 1.
size_t get_global_id (uint <i>dimindx</i>)	Returns the unique global work-item ID value for dimension identified by <i>dimindx</i> . The global work-item ID specifies the work-item ID based on the number of global work-items specified to execute the kernel. Valid values of <i>dimindx</i> are 0 to get_work_dim () – 1. For other values of <i>dimindx</i> , get_global_id () returns 0.
size_t get_local_size (uint <i>dimindx</i>)	Returns the number of local work-items specified in dimension identified by <i>dimindx</i> . This value is at most the value given by the <i>local_work_size</i> argument to clEnqueueNDRangeKernel if <i>local_work_size</i> is not NULL; otherwise the OpenCL implementation chooses an appropriate <i>local_work_size</i> value which is returned by this

	<p>function. If the kernel is executed with a non-uniform work-group size¹⁶, calls to this built-in from some work-groups may return different values than calls to this built-in from other work-groups.</p> <p>Valid values of <i>dimindx</i> are 0 to get_work_dim() – 1. For other values of <i>dimindx</i>, get_local_size() returns 1.</p>
size_t get_enqueued_local_size (uint <i>dimindx</i>)	<p>Returns the same value as that returned by get_local_size(dimindx) if the kernel is executed with a uniform work-group size.</p> <p>If the kernel is executed with a non-uniform work-group size, returns the number of local work-items in each of the work-groups that make up the uniform region of the global range in the dimension identified by <i>dimindx</i>. If the <i>local_work_size</i> argument to clEnqueueNDRangeKernel is not NULL, this value will match the value specified in <i>local_work_size[dimindx]</i>. If <i>local_work_size</i> is NULL, this value will match the local size that the implementation determined would be most efficient at implementing the uniform region of the global range.</p> <p>Valid values of <i>dimindx</i> are 0 to get_work_dim() – 1. For other values of <i>dimindx</i>, get_enqueued_local_size() returns 1.</p>
size_t get_local_id (uint <i>dimindx</i>)	<p>Returns the unique local work-item ID i.e. a work-item within a specific work-group for dimension identified by <i>dimindx</i>. Valid values of <i>dimindx</i> are 0 to get_work_dim() – 1. For other values of <i>dimindx</i>, get_local_id() returns 0.</p>
size_t get_num_groups (uint <i>dimindx</i>)	<p>Returns the number of work-groups that will execute a kernel for dimension identified by <i>dimindx</i>. Valid values of <i>dimindx</i> are 0 to get_work_dim() – 1. For other values of <i>dimindx</i>, get_num_groups() returns 1.</p>
size_t get_group_id (uint <i>dimindx</i>)	<p>get_group_id returns the work-group ID which is a number from 0 .. get_num_groups(dimindx) – 1. Valid values of <i>dimindx</i> are 0 to get_work_dim() – 1. For other values, get_group_id() returns 0.</p>
size_t get_global_offset (<p>get_global_offset returns the offset values specified</p>

¹⁶ i.e. the *global_work_size* values specified to **clEnqueueNDRangeKernel** are not evenly divisible by the *local_work_size* values for each dimension.

uint <i>dimindx</i>)	in <i>global_work_offset</i> argument to clEnqueueNDRangeKernel . Valid values of <i>dimindx</i> are 0 to get_work_dim() – 1. For other values, get_global_offset() returns 0.
size_t get_global_linear_id ()	<p>Returns the work-items 1-dimensional global ID. For 1D work-groups, it is computed as get_global_id(0) – get_global_offset(0).</p> <p>For 2D work-groups, it is computed as (get_global_id(1) – get_global_offset(1)) * get_global_size(0) + (get_global_id(0) – get_global_offset(0)).</p> <p>For 3D work-groups, it is computed as ((get_global_id(2) – get_global_offset(2)) * get_global_size(1) * get_global_size(0)) + ((get_global_id(1) – get_global_offset(1)) * get_global_size(0)) + (get_global_id(0) – get_global_offset(0)).</p>
size_t get_local_linear_id ()	<p>Returns the work-items 1-dimensional local ID. For 1D work-groups, it is the same value as get_local_id(0).</p> <p>For 2D work-groups, it is computed as get_local_id(1) * get_local_size(0) + get_local_id(0).</p> <p>For 3D work-groups, it is computed as (get_local_id(2) * get_local_size(1) * get_local_size(0)) + (get_local_id(1) * get_local_size(0)) + get_local_id(0).</p>
size_t get_sub_group_size ()	Returns the number of work-items in the subgroup. This value is no more than the maximum subgroup size and is implementation-defined based on a combination of the compiled kernel and the dispatch dimensions. This will be a constant value for the lifetime of the subgroup.
size_t get_max_sub_group_size ()	Returns the maximum size of a subgroup within the dispatch. This value will be invariant for a given set of dispatch dimensions and a kernel object compiled for a given device.
size_t get_num_sub_groups ()	Returns the number of subgroups that the current workgroup is divided into.



	This number will be constant for the duration of a workgroup's execution. If the kernel is executed with a non-uniform work-group size ¹⁷ values for any dimension, calls to this built-in from some work-groups may return different values than calls to this built-in from other work-groups.
size_t get_enqueued_num_sub_groups ()	Returns the same value as that returned by get_num_sub_groups if the kernel is executed with a uniform work-group size. If the kernel is executed with a non-uniform work-group size, returns the number of sub groups in each of the work groups that make up the uniform region of the global range.
size_t get_sub_group_id ()	get_sub_group_id returns the sub-group ID which is a number from 0 .. get_num_sub_groups() – 1. For clEnqueueTask , this returns 0.
size_t get_sub_group_local_id ()	Returns the unique work-item ID within the current subgroup. The mapping from get_local_id(dimindx) to get_sub_group_local_id will be invariant for the lifetime of the workgroup.

Table 2.1 *Work-Item Functions Table*

¹⁷ i.e. the `global_work_size` values specified to **clEnqueueNDRangeKernel** are not evenly divisible by the `local_work_size` values for each dimension.

2.2 Math Functions

The list of built-in math functions is described in *table 2.2.1*. The built-in math functions are categorized into the following:

-  A list of built-in functions that have scalar or vector argument versions, and,
-  A list of built-in functions that only take scalar float arguments.

The vector versions of the math functions operate component-wise. The description is per-component.

The built-in math functions are not affected by the prevailing rounding mode in the calling environment, and always return the same value as they would if called with the round to nearest even rounding mode.

Table 2.2.1 describes the list of built-in math functions that can take scalar or vector arguments. We use the generic type name `gentype` to indicate that the function can take `half`, `half2`, `half3`, `half4`, `half8`, `half16`, `float`, `float2`, `float3`, `float4`, `float8`, `float16`, `double`, `double2`, `double3`, `double4`, `double8` or `double16` as the type for the arguments. We use the generic type name `gentypeh` to indicate that the function can take `half`, `half2`, `half3`, `half4`, `half8`, or `half16` as the type for the arguments. We use the generic type name `gentypef` to indicate that the function can take `float`, `float2`, `float3`, `float4`, `float8`, or `float16` as the type for the arguments. We use the generic type name `gentyped` to indicate that the function can take `double`, `double2`, `double3`, `double4`, `double8` or `double16` as the type for the arguments. For any specific use of a function, the actual type has to be the same for all arguments and the return type, unless otherwise specified.

Defined in header `<opencl_math>`

Function	Description
<code>gentype acos (gentype)</code>	Arc cosine function.
<code>gentype acosh (gentype)</code>	Inverse hyperbolic cosine.
<code>gentype acospi (gentype x)</code>	Compute acos (x) / π .
<code>gentype asin (gentype)</code>	Arc sine function.
<code>gentype asinh (gentype)</code>	Inverse hyperbolic sine.
<code>gentype asinpi (gentype x)</code>	Compute asin (x) / π .
<code>gentype atan (gentype y_over_x)</code>	Arc tangent function.
<code>gentype atan2 (gentype y, gentype x)</code>	Arc tangent of y / x .
<code>gentype atanh (gentype)</code>	Hyperbolic arc tangent.
<code>gentype atanpi (gentype x)</code>	Compute atan (x) / π .

Provisional Specification

gentype atan2pi (gentype y , gentype x)	Compute atan2 (y, x) / π .
gentype cbrt (gentype)	Compute cube-root.
gentype ceil (gentype)	Round to integral value using the round to positive infinity rounding mode.
gentype copysign (gentype x , gentype y)	Returns x with its sign changed to match the sign of y .
gentype cos (gentype)	Compute cosine.
gentype cosh (gentype)	Compute hyperbolic consine.
gentype cospi (gentype x)	Compute cos (πx).
gentype erfc (gentype)	Complementary error function.
gentype erf (gentype)	Error function encountered in integrating the normal distribution.
gentype exp (gentype x)	Compute the base- e exponential of x .
gentype exp2 (gentype)	Exponential base 2 function.
gentype exp10 (gentype)	Exponential base 10 function.
gentype expm1 (gentype x)	Compute $e^x - 1.0$.
gentype fabs (gentype)	Compute absolute value of a floating-point number.
gentype fdim (gentype x , gentype y)	$x - y$ if $x > y$, +0 if x is less than or equal to y .
gentype floor (gentype)	Round to integral value using the round to negative infinity rounding mode.
gentype fma (gentype a , gentype b , gentype c)	Returns the correctly rounded floating-point representation of the sum of c with the infinitely precise product of a and b . Rounding of intermediate products shall not occur. Edge case behavior is per the IEEE 754-2008 standard.
gentype fmax (gentype x , gentype y) gentypeh fmax (gentypeh x , half y) gentypef fmax (gentypef x , float y) gentyped fmax (gentyped x , double y)	Returns y if $x < y$, otherwise it returns x . If one argument is a NaN, fmax() returns the other argument. If both arguments are NaNs, fmax() returns a NaN.
gentype fmin ¹⁸ (gentype x , gentype y) gentypeh fmin (gentypeh x , half y) gentypef fmin (gentypef x , float y) gentyped fmin (gentyped x , double y)	Returns y if $y < x$, otherwise it returns x . If one argument is a NaN, fmin() returns the other argument. If both arguments are NaNs, fmin() returns a NaN.

¹⁸ **fmin** and **fmax** behave as defined by C++ 14 and may not match the IEEE 754-2008 definition for **minNum** and **maxNum** with regard to signaling NaNs. Specifically, signaling NaNs may behave as quiet NaNs.

Provisional Specification

gentype fmod (gentype <i>x</i> , gentype <i>y</i>)	Modulus. Returns $x - y * \text{trunc}(x/y)$.
gentype fract (gentype <i>x</i> , gentype <i>*iptr</i>) ¹⁹	Returns fmin ($x - \text{floor}(x)$, 0x1.fffffep-1f). floor (<i>x</i>) is returned in <i>iptr</i> .
halfn frexp (halfn <i>x</i> , intn <i>*exp</i>) half frexp (half <i>x</i> , int <i>*exp</i>)	Extract mantissa and exponent from <i>x</i> . For each component the mantissa returned is a half with magnitude in the interval [1/2, 1) or 0. Each component of <i>x</i> equals mantissa returned * 2 ^{<i>exp</i>} .
floatn frexp (floatn <i>x</i> , intn <i>*exp</i>) float frexp (float <i>x</i> , int <i>*exp</i>)	Extract mantissa and exponent from <i>x</i> . For each component the mantissa returned is a float with magnitude in the interval [1/2, 1) or 0. Each component of <i>x</i> equals mantissa returned * 2 ^{<i>exp</i>} .
doublen frexp (doublen <i>x</i> , intn <i>*exp</i>) double frexp (double <i>x</i> , int <i>*exp</i>)	Extract mantissa and exponent from <i>x</i> . For each component the mantissa returned is a double with magnitude in the interval [1/2, 1) or 0. Each component of <i>x</i> equals mantissa returned * 2 ^{<i>exp</i>} .
gentype hypot (gentype <i>x</i> , gentype <i>y</i>)	Compute the value of the square root of $x^2 + y^2$ without undue overflow or underflow.
intn ilogb (floatn <i>x</i>) int ilogb (float <i>x</i>) intn ilogb (doublen <i>x</i>) int ilogb (double <i>x</i>) intn ilogb (halfn <i>x</i>) int ilogb (half <i>x</i>)	Return the exponent as an integer value.
floatn ldexp (floatn <i>x</i> , intn <i>k</i>) floatn ldexp (floatn <i>x</i> , int <i>k</i>) float ldexp (float <i>x</i> , int <i>k</i>) doublen ldexp (doublen <i>x</i> , intn <i>k</i>) doublen ldexp (doublen <i>x</i> , int <i>k</i>) double ldexp (double <i>x</i> , int <i>k</i>) halfn ldexp (halfn <i>x</i> , intn <i>k</i>) halfn ldexp (halfn <i>x</i> , int <i>k</i>) half ldexp (half <i>x</i> , int <i>k</i>)	Multiply <i>x</i> by 2 to the power <i>k</i> .
gentype lgamma (gentype <i>x</i>) floatn lgamma_r (floatn <i>x</i> , intn <i>*signp</i>) float lgamma_r (float <i>x</i> , int <i>*signp</i>)	Log gamma function. Returns the natural logarithm of the absolute value of the gamma function. The sign of the gamma function is returned in the <i>signp</i> argument of lgamma_r .

¹⁹ The min() operator is there to prevent **fract**(-small) from returning 1.0. It returns the largest positive floating-point number less than 1.0.

Provisional Specification

doublen lgamma_r (doublen <i>x</i> , intn * <i>signp</i>) double lgamma_r (double <i>x</i> , int * <i>signp</i>) halfn lgamma_r (halfn <i>x</i> , intn * <i>signp</i>) half lgamma_r (half <i>x</i> , int * <i>signp</i>)	
gentype log (gentype)	Compute natural logarithm.
gentype log2 (gentype)	Compute a base 2 logarithm.
gentype log10 (gentype)	Compute a base 10 logarithm.
gentype log1p (gentype <i>x</i>)	Compute $\log_e(1.0 + x)$.
gentype logb (gentype <i>x</i>)	Compute the exponent of <i>x</i> , which is the integral part of $\log_r x $.
gentype mad (gentype <i>a</i> , gentype <i>b</i> , gentype <i>c</i>)	mad approximates $a * b + c$. Whether or how the product of $a * b$ is rounded and how supernormal or subnormal intermediate products are handled is not defined. mad is intended to be used where speed is preferred over accuracy ²⁰ .
gentype maxmag (gentype <i>x</i> , gentype <i>y</i>)	Returns <i>x</i> if $ x > y $, <i>y</i> if $ y > x $, otherwise fmax (<i>x</i> , <i>y</i>).
gentype minmag (gentype <i>x</i> , gentype <i>y</i>)	Returns <i>x</i> if $ x < y $, <i>y</i> if $ y < x $, otherwise fmin (<i>x</i> , <i>y</i>).
gentype modf (gentype <i>x</i> , gentype * <i>iptr</i>)	Decompose a floating-point number. The modf function breaks the argument <i>x</i> into integral and fractional parts, each of which has the same sign as the argument. It stores the integral part in the object pointed to by <i>iptr</i> .
floatn nan (uintn <i>nancode</i>) float nan (uint <i>nancode</i>) doublen nan (ulongn <i>nancode</i>) double nan (ulong <i>nancode</i>) halfn nan (uintn <i>nancode</i>) half nan (uint <i>nancode</i>)	Returns a quiet NaN. The <i>nancode</i> may be placed in the significand of the resulting NaN.
gentype nextafter (gentype <i>x</i> , gentype <i>y</i>)	Computes the next representable single-precision floating-point value following <i>x</i> in the direction of <i>y</i> . Thus, if <i>y</i> is less than <i>x</i> , nextafter () returns the largest representable

²⁰ The user is cautioned that for some usages, e.g. **mad**(*a*, *b*, -*a***b*), the definition of **mad**() is loose enough that almost any result is allowed from **mad**() for some values of *a* and *b*.


Provisional Specification


	floating-point number less than x .
gentype pow (gentype x , gentype y)	Compute x to the power y .
floatn pown (floatn x , intn y) float pown (float x , int y) doublen pown (doublen x , intn y) double pown (double x , int y) halfn pown (halfn x , intn y) half pown (half x , int y)	Compute x to the power y , where y is an integer.
gentype powr (gentype x , gentype y)	Compute x to the power y , where x is ≥ 0 .
gentype remainder (gentype x , gentype y)	Compute the value r such that $r = x - n*y$, where n is the integer nearest the exact value of x/y . If there are two integers closest to x/y , n shall be the even one. If r is zero, it is given the same sign as x .
floatn remquo (floatn x , floatn y , intn $*quo$) float remquo (float x , float y , int $*quo$) doublen remquo (doublen x , doublen y , intn $*quo$) double remquo (double x , double y , int $*quo$) halfn remquo (halfn x , halfn y , intn $*quo$) half remquo (half x , half y , int $*quo$)	The remquo function computes the value r such that $r = x - k*y$, where k is the integer nearest the exact value of x/y . If there are two integers closest to x/y , k shall be the even one. If r is zero, it is given the same sign as x . This is the same value that is returned by the remainder function. remquo also calculates the lower seven bits of the integral quotient x/y , and gives that value the same sign as x/y . It stores this signed value in the object pointed to by quo .
gentype rint (gentype)	Round to integral value (using round to nearest even rounding mode) in floating-point format. Refer to <i>section 1.2.3.2</i> for description of rounding modes.
floatn rootn (floatn x , intn y) float rootn (float x , int y) doublen rootn (doublen x , intn y) doublen rootn (double x , int y)	Compute x to the power $1/y$.

<code>halfn rootn (halfn x, intn y)</code> <code>half rootn (half x, int y)</code>	
<code>gentype round (gentype x)</code>	Return the integral value nearest to x rounding halfway cases away from zero, regardless of the current rounding direction.
<code>gentype rsqrt (gentype)</code>	Compute inverse square root.
<code>gentype sin (gentype)</code>	Compute sine.
<code>gentype sincos (gentype x, gentype *cosval)</code>	Compute sine and cosine of x . The computed sine is the return value and computed cosine is returned in <i>cosval</i> .
<code>gentype sinh (gentype)</code>	Compute hyperbolic sine.
<code>gentype sinpi (gentype x)</code>	Compute sin (πx).
<code>gentype sqr (gentype)</code>	Compute square root.
<code>gentype tan (gentype)</code>	Compute tangent.
<code>gentype tanh (gentype)</code>	Compute hyperbolic tangent.
<code>gentype tanpi (gentype x)</code>	Compute tan (πx).
<code>gentype tgamma (gentype)</code>	Compute the gamma function.
<code>gentype trunc (gentype)</code>	Round to integral value using the round to zero rounding mode.

Table 2.2.1 *Scalar and Vector Argument Built-in Math Function Table*

Table 2.2.2 describes the following functions:

 A subset of functions from *table 2.2.1* that are defined in the `cl::native` namespace. These functions may map to one or more native device instructions and will typically have better performance compared to the corresponding functions (without the `native__` prefix) described in *table 2.2.1*. The accuracy (and in some cases the input range(s)) of these functions is implementation-defined.

 `cl::native` functions for following basic operations: divide and reciprocal.

We use the generic type name `gentype` to indicate that the functions in *table 6.9* can take `float`, `float2`, `float3`, `float4`, `float8` or `float16` as the type for the arguments.

Function	Description
<code>gentype native::cos (gentype x)</code>	Compute cosine over an implementation-defined range. The maximum error is implementation-defined.
<code>gentype native::divide (gentype x, gentype y)</code>	Compute x / y over an implementation-defined range. The maximum error is implementation-

	defined.
gentype native::exp (gentype x)	Compute the base- e exponential of x over an implementation-defined range. The maximum error is implementation-defined.
gentype native::exp2 (gentype x)	Compute the base- 2 exponential of x over an implementation-defined range. The maximum error is implementation-defined.
gentype native::exp10 (gentype x)	Compute the base- 10 exponential of x over an implementation-defined range. The maximum error is implementation-defined.
gentype native::log (gentype x)	Compute natural logarithm over an implementation-defined range. The maximum error is implementation-defined.
gentype native::log2 (gentype x)	Compute a base 2 logarithm over an implementation-defined range. The maximum error is implementation-defined.
gentype native::log10 (gentype x)	Compute a base 10 logarithm over an implementation-defined range. The maximum error is implementation-defined.
gentype native::powr (gentype x , gentype y)	Compute x to the power y , where x is ≥ 0 . The range of x and y are implementation-defined. The maximum error is implementation-defined.
gentype native::recip (gentype x)	Compute reciprocal over an implementation-defined range. The maximum error is implementation-defined.
gentype native::rsqrt (gentype x)	Compute inverse square root over an implementation-defined range. The maximum error is implementation-defined.
gentype native::sin (gentype x)	Compute sine over an implementation-defined range. The maximum error is implementation-defined.
gentype native::sqrt (gentype x)	Compute square root over an implementation-defined range. The maximum error is implementation-defined.
gentype native::tan (gentype x)	Compute tangent over an implementation-defined range. The maximum error is implementation-defined.

Table 2.2.2 *Scalar and Vector Argument Built-in cl::native:: Math Functions*

Support for denormal values is implementation-defined for **cl::native** functions.

The following symbolic constants are available. Their values are of type `float` and are accurate within the precision of a single precision floating-point number.

Constant Name	Description
MAXFLOAT	Value of maximum non-infinite single-precision floating-point number.
HUGE_VALF	A positive float constant expression. HUGE_VALF evaluates to +infinity. Used as an error value returned by the built-in math functions.
INFINITY	A constant expression of type float representing positive or unsigned infinity.
NAN	A constant expression of type float representing a quiet NaN.

If double precision is supported by the device, the following symbolic constant will also be available:

Constant Name	Description
HUGE_VAL	A positive double constant expression. HUGE_VAL evaluates to +infinity. Used as an error value returned by the built-in math functions.

2.2.1 Floating-point macros and pragmas

The **FP_CONTRACT** pragma can be used to allow (if the state is `on`) or disallow (if the state is `off`) the implementation to contract expressions. Each pragma can occur either outside external declarations or preceding all explicit declarations and statements inside a compound statement. When outside external declarations, the pragma takes effect from its occurrence until another **FP_CONTRACT** pragma is encountered, or until the end of the translation unit. When inside a compound statement, the pragma takes effect from its occurrence until another **FP_CONTRACT** pragma is encountered (including within a nested compound statement), or until the end of the compound statement; at the end of a compound statement the state for the pragma is restored to its condition just before the compound statement. If this pragma is used in any other context, the behavior is undefined.

The pragma definition to set **FP_CONTRACT** is:

```
#pragma OPENCL FP_CONTRACT on-off-switch
```

`on-off-switch` is one of:

ON, **OFF** or **DEFAULT**.

The **DEFAULT** value is **ON**.

The macro names given in the following list must use the values specified. These constant expressions are suitable for use in `#if` preprocessing directives.

Provisional Specification

```
#define FLT_DIG 6
#define FLT_MANT_DIG 24
#define FLT_MAX_10_EXP +38
#define FLT_MAX_EXP +128
#define FLT_MIN_10_EXP -37
#define FLT_MIN_EXP -125
#define FLT_RADIX 2
#define FLT_MAX 0x1.fffffep127f
#define FLT_MIN 0x1.0p-126f
#define FLT_EPSILON 0x1.0p-23f
```

The following table describes the built-in macro names given above in the OpenCL C++ programming language and the corresponding macro names available to the application.

Macro in OpenCL Language	Macro for application
FLT_DIG	CL_FLT_DIG
FLT_MANT_DIG	CL_FLT_MANT_DIG
FLT_MAX_10_EXP	CL_FLT_MAX_10_EXP
FLT_MAX_EXP	CL_FLT_MAX_EXP
FLT_MIN_10_EXP	CL_FLT_MIN_10_EXP
FLT_MIN_EXP	CL_FLT_MIN_EXP
FLT_RADIX	CL_FLT_RADIX
FLT_MAX	CL_FLT_MAX
FLT_MIN	CL_FLT_MIN
FLT_EPSILON	CL_FLT_EPSILON

The following macros shall expand to integer constant expressions whose values are returned by **ilogb**(*x*) if *x* is zero or NaN, respectively. The value of **FP_ILOGB0** shall be either {INT_MIN} or - {INT_MAX}. The value of **FP_ILOGBNAN** shall be either {INT_MAX} or {INT_MIN}.

The following constants are also available. They are of type `float` and are accurate within the precision of the `float` type.

Constant	Description
M_E_F	Value of <i>e</i>
M_LOG2E_F	Value of $\log_2 e$
M_LOG10E_F	Value of $\log_{10} e$
M_LN2_F	Value of $\log_e 2$
M_LN10_F	Value of $\log_e 10$
M_PI_F	Value of π
M_PI_2_F	Value of $\pi / 2$
M_PI_4_F	Value of $\pi / 4$
M_1_PI_F	Value of $1 / \pi$

Provisional Specification

M_2_PI_F	Value of $2 / \pi$
M_2_SQRTPI_F	Value of $2 / \sqrt{\pi}$
M_SQRT2_F	Value of $\sqrt{2}$
M_SQRT1_2_F	Value of $1 / \sqrt{2}$

If double precision is supported by the device, the following macros and constants are available:

The macro names given in the following list must use the values specified. These constant expressions are suitable for use in #if preprocessing directives.

```
#define DBL_DIG          15
#define DBL_MANT_DIG     53
#define DBL_MAX_10_EXP   +308
#define DBL_MAX_EXP      +1024
#define DBL_MIN_10_EXP   -307
#define DBL_MIN_EXP      -1021
#define DBL_MAX           0x1.fffffffffffffp1023
#define DBL_MIN           0x1.0p-1022
#define DBL_EPSILON      0x1.0p-52
```

The following table describes the built-in macro names given above in the OpenCL C++ programming language and the corresponding macro names available to the application.

Macro in OpenCL Language	Macro for application
DBL_DIG	CL_DBL_DIG
DBL_MANT_DIG	CL_DBL_MANT_DIG
DBL_MAX_10_EXP	CL_DBL_MAX_10_EXP
DBL_MAX_EXP	CL_DBL_MAX_EXP
DBL_MIN_10_EXP	CL_DBL_MIN_10_EXP
DBL_MIN_EXP	CL_DBL_MIN_EXP
DBL_MAX	CL_DBL_MAX
DBL_MIN	CL_DBL_MIN
DBL_EPSILON	CL_DBL_EPSILON

The following constants are also available. They are of type `double` and are accurate within the precision of the `double` type.

Constant	Description
M_E	Value of e
M_LOG2E	Value of $\log_2 e$
M_LOG10E	Value of $\log_{10} e$
M_LN2	Value of $\log_e 2$

Provisional Specification

M_LN10	Value of $\log_e 10$
M_PI	Value of π
M_PI_2	Value of $\pi / 2$
M_PI_4	Value of $\pi / 4$
M_1_PI	Value of $1 / \pi$
M_2_PI	Value of $2 / \pi$
M_2_SQRTPI	Value of $2 / \sqrt{\pi}$
M_SQRT2	Value of $\sqrt{2}$
M_SQRT1_2	Value of $1 / \sqrt{2}$

If half precision arithmetic operations are supported, the following macros and constants for half precision floating-point are available:

The macro names given in the following list must use the values specified. These constant expressions are suitable for use in #if preprocessing directives.

```
#define HALF_DIG          3
#define HALF_MANT_DIG     11
#define HALF_MAX_10_EXP   +4
#define HALF_MAX_EXP       +16
#define HALF_MIN_10_EXP   -4
#define HALF_MIN_EXP      -13
#define HALF_RADIX        2
#define HALF_MAX          0x1.ffcp15h
#define HALF_MIN          0x1.0p-14h
#define HALF_EPSILON      0x1.0p-10h
```

The following table describes the built-in macro names given above in the OpenCL C++ programming language and the corresponding macro names available to the application.

Macro in OpenCL Language	Macro for application
HALF_DIG	CL_HALF_DIG
HALF_MANT_DIG	CL_HALF_MANT_DIG
HALF_MAX_10_EXP	CL_HALF_MAX_10_EXP
HALF_MAX_EXP	CL_HALF_MAX_EXP
HALF_MIN_10_EXP	CL_HALF_MIN_10_EXP
HALF_MIN_EXP	CL_HALF_MIN_EXP
HALF_RADIX	CL_HALF_RADIX
HALF_MAX	CL_HALF_MAX
HALF_MIN	CL_HALF_MIN
HALF_EPSILON	CL_HALF_EPSILON

Provisional Specification

The following constants are also available. They are of type `half` and are accurate within the precision of the `half` type.

Constant	Description
<code>M_E_H</code>	Value of e
<code>M_LOG2E_H</code>	Value of $\log_2 e$
<code>M_LOG10E_H</code>	Value of $\log_{10} e$
<code>M_LN2_H</code>	Value of $\log_e 2$
<code>M_LN10_H</code>	Value of $\log_e 10$
<code>M_PI_H</code>	Value of π
<code>M_PI_2_H</code>	Value of $\pi / 2$
<code>M_PI_4_H</code>	Value of $\pi / 4$
<code>M_1_PI_H</code>	Value of $1 / \pi$
<code>M_2_PI_H</code>	Value of $2 / \pi$
<code>M_2_SQRTPI_H</code>	Value of $2 / \sqrt{\pi}$
<code>M_SQRT2_H</code>	Value of $\sqrt{2}$
<code>M_SQRT1_2_H</code>	Value of $1 / \sqrt{2}$

2.3 Integer Functions

Table 2.3.1 describes the built-in integer functions that take scalar or vector arguments. The vector versions of the integer functions operate component-wise. The description is per-component.

We use the generic type name `gentype` to indicate that the function can take `char`, `char{2|3|4|8|16}`, `uchar`, `uchar{2|3|4|8|16}`, `short`, `short{2|3|4|8|16}`, `ushort`, `ushort{2|3|4|8|16}`, `int`, `int{2|3|4|8|16}`, `uint`, `uint{2|3|4|8|16}`, `long`, `long{2|3|4|8|16}` `ulong`, or `ulong{2|3|4|8|16}` as the type for the arguments. We use the generic type name `ugentype` to refer to unsigned versions of `gentype`. For example, if `gentype` is `char4`, `ugentype` is `uchar4`. We also use the generic type name `sgentype` to indicate that the function can take a scalar data type i.e. `char`, `uchar`, `short`, `ushort`, `int`, `uint`, `long`, or `ulong` as the type for the arguments. For built-in integer functions that take `gentype` and `sgentype` arguments, the `gentype` argument must be a vector or scalar version of the `sgentype` argument. For example, if `sgentype` is `uchar`, `gentype` must be `uchar` or `uchar{2|3|4|8|16}`. For vector versions, `sgentype` is implicitly widened to `gentype` as described in section 1.3.a.

For any specific use of a function, the actual type has to be the same for all arguments and the return type unless otherwise specified.

Defined in header `<opencl_integer>`

Function	Description
<code>ugentype abs (gentype x)</code>	Returns $ x $.
<code>ugentype abs_diff (gentype x, gentype y)</code>	Returns $ x - y $ without modulo overflow.
<code>gentype add_sat (gentype x, gentype y)</code>	Returns $x + y$ and saturates the result.
<code>gentype hadd (gentype x, gentype y)</code>	Returns $(x + y) \gg 1$. The intermediate sum does not modulo overflow.
<code>gentype rhadd (gentype x, gentype y)</code> ²¹	Returns $(x + y + 1) \gg 1$. The intermediate sum does not modulo overflow.
<code>gentype clamp (gentype x, gentype minval, gentype maxval)</code>	Returns $\min(\max(x, \text{minval}), \text{maxval})$. Results are undefined if $\text{minval} > \text{maxval}$.
<code>gentype clamp (gentype x,</code>	

²¹ Frequently vector operations need $n + 1$ bits temporarily to calculate a result. The **rhadd** instruction gives you an extra bit without needing to upsample and downsample. This can be a profound performance win.

Provisional Specification

sgentype <i>minval</i> , sgentype <i>maxval</i>)	
gentype clz (gentype <i>x</i>)	Returns the number of leading 0-bits in <i>x</i> , starting at the most significant bit position. If <i>x</i> is 0, returns the size in bits of the type of <i>x</i> or component type of <i>x</i> , if <i>x</i> is a vector.
gentype ctz (gentype <i>x</i>)	Returns the count of trailing 0-bits in <i>x</i> . If <i>x</i> is 0, returns the size in bits of the type of <i>x</i> or component type of <i>x</i> , if <i>x</i> is a vector.
gentype mad_hi (gentype <i>a</i> , gentype <i>b</i> , gentype <i>c</i>)	Returns mul_hi (<i>a</i> , <i>b</i>) + <i>c</i> .
gentype mad_sat (gentype <i>a</i> , gentype <i>b</i> , gentype <i>c</i>)	Returns $a * b + c$ and saturates the result.
gentype max (gentype <i>x</i> , gentype <i>y</i>) gentype max (gentype <i>x</i> , sgentype <i>y</i>)	Returns <i>y</i> if $x < y$, otherwise it returns <i>x</i> .
gentype min (gentype <i>x</i> , gentype <i>y</i>) gentype min (gentype <i>x</i> , sgentype <i>y</i>)	Returns <i>y</i> if $y < x$, otherwise it returns <i>x</i> .
gentype mul_hi (gentype <i>x</i> , gentype <i>y</i>)	Computes $x * y$ and returns the high half of the product of <i>x</i> and <i>y</i> .
gentype rotate (gentype <i>v</i> , gentype <i>i</i>)	For each element in <i>v</i> , the bits are shifted left by the number of bits given by the corresponding element in <i>i</i> (subject to usual shift modulo rules described in <i>section 1.3</i>). Bits shifted off the left side of the element are shifted back in from the right.
gentype sub_sat (gentype <i>x</i> , gentype <i>y</i>)	Returns $x - y$ and saturates the result.
short upsample (char <i>hi</i> , uchar <i>lo</i>) ushort upsample (uchar <i>hi</i> , uchar <i>lo</i>) shortn upsample (charn <i>hi</i> , uchar <i>lo</i>) ushortn upsample (ucharn <i>hi</i> , uchar <i>lo</i>) int upsample (short <i>hi</i> , ushort <i>lo</i>) uint upsample (ushort <i>hi</i> , ushort <i>lo</i>) intn upsample (shortn <i>hi</i> , ushortn <i>lo</i>) uintn upsample (ushortn <i>hi</i> , ushortn <i>lo</i>) long upsample (int <i>hi</i> , uint <i>lo</i>) ulong upsample (uint <i>hi</i> , uint <i>lo</i>) longn upsample (intn <i>hi</i> , uintn <i>lo</i>) ulongn upsample (uintn <i>hi</i> , uintn <i>lo</i>)	$result[i] = ((short)hi[i] << 8) lo[i]$ $result[i] = ((ushort)hi[i] << 8) lo[i]$ $result[i] = ((int)hi[i] << 16) lo[i]$ $result[i] = ((uint)hi[i] << 16) lo[i]$ $result[i] = ((long)hi[i] << 32) lo[i]$ $result[i] = ((ulong)hi[i] << 32) lo[i]$

gentype popcount (gentype x)	Returns the number of non-zero bits in x.
-------------------------------------	---

Table 2.3.1 *Scalar and Vector Integer Argument Built-in Functions*

Table 2.3.2 describes fast integer functions that can be used for optimizing performance of kernels. We use the generic type name `gentype` to indicate that the function can take `int`, `int2`, `int3`, `int4`, `int8`, `int16`, `uint`, `uint2`, `uint3`, `uint4`, `uint8` or `uint16` as the type for the arguments.

Function	Description
gentype mad24 (gentype x, gentype y, gentype z)	Multiply two 24-bit integer values x and y and add the 32-bit integer result to the 32-bit integer z. Refer to definition of mul24 to see how the 24-bit integer multiplication is performed.
gentype mul24 (gentype x, gentype y)	Multiply two 24-bit integer values x and y. x and y are 32-bit integers but only the low 24-bits are used to perform the multiplication. mul24 should only be used when values in x and y are in the range $[-2^{23}, 2^{23}-1]$ if x and y are signed integers and in the range $[0, 2^{24}-1]$ if x and y are unsigned integers. If x and y are not in this range, the multiplication result is implementation-defined.

Table 2.3.2 *Fast Integer Built-in Functions*

The macro names given in the following list must use the values specified. The values shall all be constant expressions suitable for use in `#if` preprocessing directives.

```
#define CHAR_BIT      8
#define CHAR_MAX      SCHAR_MAX
#define CHAR_MIN      SCHAR_MIN
#define INT_MAX        2147483647
#define INT_MIN        (-2147483647 - 1)
#define LONG_MAX       0x7fffffffffffffffL
#define LONG_MIN       (-0x7fffffffffffffffL - 1)
#define SCHAR_MAX      127
#define SCHAR_MIN      (-127 - 1)
#define SHRT_MAX       32767
#define SHRT_MIN       (-32767 - 1)
#define UCHAR_MAX      255
```

Provisional Specification

```
#define USHRT_MAX    65535
#define UINT_MAX     0xffffffff
#define ULONG_MAX    0xffffffffffffffffUL
```

The following table describes the built-in macro names given above in the OpenCL C++ programming language and the corresponding macro names available to the application.

Macro in OpenCL Language	Macro for application
CHAR_BIT	CL_CHAR_BIT
CHAR_MAX	CL_CHAR_MAX
CHAR_MIN	CL_CHAR_MIN
INT_MAX	CL_INT_MAX
INT_MIN	CL_INT_MIN
LONG_MAX	CL_LONG_MAX
LONG_MIN	CL_LONG_MIN
SCHAR_MAX	CL_SCHAR_MAX
SCHAR_MIN	CL_SCHAR_MIN
SHRT_MAX	CL_SHRT_MAX
SHRT_MIN	CL_SHRT_MIN
UCHAR_MAX	CL_UCHAR_MAX
USHRT_MAX	CL_USHRT_MAX
UINT_MAX	CL_UINT_MAX
ULONG_MAX	CL_ULONG_MAX

2.4 Common Functions

Table 2.4 describes the list of built-in common functions. These all operate component-wise. The description is per-component. We use the generic type name `gentype` to indicate that the function can take `half`, `half2`, `half3`, `half4`, `half8`, `half16`, `float`, `float2`, `float3`, `float4`, `float8`, `float16`, `double`, `double2`, `double3`, `double4`, `double8` or `double16` as the type for the arguments. We use the generic type name `gentypeh` to indicate that the function can take `half`, `half2`, `half3`, `half4`, `half8`, or `half16` as the type for the arguments. We use the generic type name `gentypef` to indicate that the function can take `float`, `float2`, `float3`, `float4`, `float8`, or `float16` as the type for the arguments. We use the generic type name `gentyped` to indicate that the function can take `double`, `double2`, `double3`, `double4`, `double8` or `double16` as the type for the arguments.

The built-in common functions are implemented using the round to nearest even rounding mode.

Defined in header `<opencl_common>`

Function	Description
<code>gentype clamp (gentype x, gentype minval, gentype maxval)</code> <code>gentypef clamp (gentypef x, float minval, float maxval)</code> <code>gentyped clamp (gentyped x, double minval, double maxval)</code> <code>gentypeh clamp (gentypeh x, half minval, half maxval)</code>	Returns fmin(fmax(<i>x</i>, <i>minval</i>), <i>maxval</i>) . Results are undefined if <i>minval</i> > <i>maxval</i> .
<code>gentype degrees (gentype radians)</code>	Converts <i>radians</i> to degrees, i.e. $(180 / \pi) * \textit{radians}$.
<code>gentype max (gentype x, gentype y)</code> <code>gentypef max (gentypef x, float y)</code>	Returns <i>y</i> if <i>x</i> < <i>y</i> , otherwise it returns <i>x</i> . If <i>x</i> or <i>y</i> are infinite or NaN, the return values are undefined.

Provisional Specification

gentyped max (gentyped <i>x</i> , double <i>y</i>)	
gentypeh max (gentypeh <i>x</i> , half <i>y</i>)	
gentype min (gentype <i>x</i> , gentype <i>y</i>)	Returns <i>y</i> if $y < x$, otherwise it returns <i>x</i> . If <i>x</i> or <i>y</i> are infinite or NaN, the return values are undefined.
gentypef min (gentypef <i>x</i> , float <i>y</i>)	
gentyped min (gentyped <i>x</i> , double <i>y</i>)	
gentypeh min (gentypeh <i>x</i> , half <i>y</i>)	
gentype mix (gentype <i>x</i> , gentype <i>y</i> , gentype <i>a</i>)	
gentypef mix (gentypef <i>x</i> , gentypef <i>y</i> , float <i>a</i>)	Returns the linear blend of <i>x</i> & <i>y</i> implemented as: $x + (y - x) * a$ <i>a</i> must be a value in the range 0.0 ... 1.0. If <i>a</i> is not in the range 0.0 ... 1.0, the return values are undefined.
gentyped mix (gentyped <i>x</i> , gentyped <i>y</i> , double <i>a</i>)	
gentypeh mix (gentypeh <i>x</i> , gentypeh <i>y</i> , half <i>a</i>)	
gentype radians (gentype <i>degrees</i>)	
gentype step (gentype <i>edge</i> , gentype <i>x</i>)	Returns 0.0 if $x < edge$, otherwise it returns 1.0.
gentypef step (float <i>edge</i> , gentypef <i>x</i>)	
gentyped step (double <i>edge</i> , gentyped <i>x</i>)	
gentypeh step (half <i>edge</i> , gentypeh <i>x</i>)	
gentype smoothstep (gentype <i>edge0</i> , gentype <i>edge1</i> , gentype <i>x</i>)	Returns 0.0 if $x \leq edge0$ and 1.0 if $x \geq edge1$ and performs smooth Hermite interpolation between 0 and 1 when $edge0 < x < edge1$. This is useful in cases where you would want a threshold function with a smooth transition. This is equivalent to: gentype t; t = clamp ((<i>x</i> - <i>edge0</i>) / (<i>edge1</i> - <i>edge0</i>), 0, 1); return t * t * (3 - 2 * t);
gentypef smoothstep (float <i>edge0</i> , float <i>edge1</i> , gentypef <i>x</i>)	
gentyped smoothstep (double <i>edge0</i> , double <i>edge1</i> , gentyped <i>x</i>)	
gentypeh smoothstep (half <i>edge0</i> , half <i>edge1</i> ,	

gentype x)	Results are undefined if $edge0 \geq edge1$ or if x , $edge0$ or $edge1$ is a NaN.
gentype sign (gentype x)	Returns 1.0 if $x > 0$, -0.0 if $x = -0.0$, +0.0 if $x = +0.0$, or -1.0 if $x < 0$. Returns 0.0 if x is a NaN.

Table 2.4 *Scalar and Vector Argument Built-in Common Function Table*

2.5 Geometric Functions

Table 2.5 describes the list of built-in geometric functions. These all operate component-wise. The description is per-component. `floatn` is `float`, `float2`, `float3`, or `float4`, `halfn` is `half`, `half2`, `half3`, or `half4` and `doublen` is `double`, `double2`, `double3`, or `double4`. The built-in geometric functions are implemented using the round to nearest even rounding mode.

Defined in header `<opencl_geometric>`

Function	Description
<code>float4 cross (float4 p0, float4 p1)</code> <code>float3 cross (float3 p0, float3 p1)</code> <code>double4 cross (double4 p0, double4 p1)</code> <code>double3 cross (double3 p0, double3 p1)</code> <code>half4 cross (half4 p0, half4 p1)</code> <code>half3 cross (half3 p0, half3 p1)</code>	Returns the cross product of $p0.xyz$ and $p1.xyz$. The w component of <code>float4</code> result returned will be 0.0.
<code>float dot (floatn p0, floatn p1)</code> <code>double dot (doublen p0, doublen p1)</code> <code>half dot (halfn p0, halfn p1)</code>	Compute dot product.
<code>float distance (floatn p0, floatn p1)</code> <code>double distance (doublen p0, doublen p1)</code> <code>half distance (halfn p0, halfn p1)</code>	Returns the distance between $p0$ and $p1$. This is calculated as length ($p0 - p1$).
<code>float length (floatn p)</code> <code>double length (doublen p)</code> <code>half length (halfn p)</code>	Return the length of vector p , i.e., $\sqrt{p.x^2 + p.y^2 + \dots}$
<code>floatn normalize (floatn p)</code> <code>doublen normalize (doublen p)</code> <code>halfn normalize (halfn p)</code>	Returns a vector in the same direction as p but with a length of 1.

Table 2.5 *Scalar and Vector Argument Built-in Geometric Function Table*

Provisional Specification

bool isgreater (double <i>x</i> , double <i>y</i>) booln isgreater (doublen <i>x</i> , doublen <i>y</i>) bool isgreater (half <i>x</i> , half <i>y</i>) booln isgreater (halfn <i>x</i> , halfn <i>y</i>)	
bool isgreaterequal (float <i>x</i> , float <i>y</i>) booln isgreaterequal (floatn <i>x</i> , floatn <i>y</i>) bool isgreaterequal (double <i>x</i> , double <i>y</i>) booln isgreaterequal (doublen <i>x</i> , doublen <i>y</i>) bool isgreaterequal (half <i>x</i> , half <i>y</i>) booln isgreaterequal (halfn <i>x</i> , halfn <i>y</i>)	Returns the component-wise compare of $x \geq y$.
bool isless (float <i>x</i> , float <i>y</i>) booln isless (floatn <i>x</i> , floatn <i>y</i>) bool isless (double <i>x</i> , double <i>y</i>) booln isless (doublen <i>x</i> , doublen <i>y</i>) bool isless (half <i>x</i> , half <i>y</i>) booln isless (halfn <i>x</i> , halfn <i>y</i>)	Returns the component-wise compare of $x < y$.
bool islessequal (float <i>x</i> , float <i>y</i>) booln islessequal (floatn <i>x</i> , floatn <i>y</i>) bool islessequal (double <i>x</i> , double <i>y</i>) booln islessequal (doublen <i>x</i> , doublen <i>y</i>) bool islessequal (half <i>x</i> , half <i>y</i>) booln islessequal (halfn <i>x</i> , halfn <i>y</i>)	Returns the component-wise compare of $x \leq y$.
bool islessgreater (float <i>x</i> , float <i>y</i>) booln islessgreater (floatn <i>x</i> , floatn <i>y</i>) bool islessgreater (double <i>x</i> , double <i>y</i>) booln islessgreater (doublen <i>x</i> , doublen <i>y</i>) bool islessgreater (half <i>x</i> , half <i>y</i>) booln islessgreater (halfn <i>x</i> , halfn <i>y</i>)	Returns the component-wise compare of $(x < y) \parallel (x > y)$.
bool isfinite (float) booln isfinite (floatn)	Test for finite value.

Provisional Specification

bool isfinite (double) booln isfinite (doublen) bool isfinite (half x y) booln isfinite (halfn x)	
bool isinf (float) booln isinf (floatn) bool isinf (double) booln isinf (doublen) bool isinf (half x) booln isinf (halfn x)	Test for infinity value (positive or negative) .
bool isnan (float) booln isnan (floatn) bool isnan (double) booln isnan (doublen) bool isnan (half x) booln isnan (halfn x)	Test for a NaN.
bool isnormal (float) booln isnormal (floatn) bool isnormal (double) booln isnormal (doublen) bool isnormal (half x) booln isnormal (halfn x)	Test for a normal value.
bool isordered (float x, float y) booln isordered (floatn x, floatn y) bool isordered (double x, double y) booln isordered (doublen x, doublen y) bool isordered (half x, half y) booln isordered (halfn x, halfn y)	Test if arguments are ordered. isordered() takes arguments <i>x</i> and <i>y</i> , and returns the result isequal (<i>x</i> , <i>x</i>) && isequal (<i>y</i> , <i>y</i>).
bool isunordered (float x, float y) booln isunordered (floatn x, floatn y) bool isunordered (double x, double y) booln isunordered (doublen x, doublen y)	Test if arguments are unordered. isunordered() takes arguments <i>x</i> and <i>y</i> , returning <i>true</i> if <i>x</i> or <i>y</i> is NaN, and <i>false</i> otherwise.

bool isunordered (half x , half y) booln isunordered (halfn x , halfn y)	
bool signbit (float) booln signbit (floatn) bool signbit (double) booln signbit (doublen) bool signbit (half) booln signbit (halfn)	Test for sign bit. The scalar version of the function returns a <i>true</i> if the sign bit in the float is set else returns <i>false</i> . The vector version of the function returns the following for each component in floatn: <i>true</i> (i.e all bits set) if the sign bit in the float is set else returns <i>false</i> .
bool any (booln x)	Returns <i>true</i> if any component of x is <i>true</i> ; otherwise returns <i>false</i> .
bool all (booln x)	Returns <i>true</i> if all components of x are <i>true</i> ; otherwise returns <i>false</i> .
gentype bitselect (gentype a , gentype b , gentype c)	Each bit of the result is the corresponding bit of a if the corresponding bit of c is 0. Otherwise it is the corresponding bit of b .
gentype select (gentype a , gentype b , bgentype c)	For each component of a vector type, $result[i] = \text{if } c[i] \text{ is true ? } b[i] : a[i]$. For a scalar type, $result = c ? b : a$. bgentype must have the same number of elements as gentype.

Table 2.6 *Scalar and Vector Relational Functions*

2.7 Vector Data Load and Store Functions

Table 2.7 describes the list of supported functions that allow you to read and write vector types from a pointer to memory. We use the generic type `gentype` to indicate the built-in data types `char`, `uchar`, `short`, `ushort`, `int`, `uint`, `long`, `ulong`, `half`, `float` or `double`. We use the generic type name `gentypen` to represent n -element vectors of `gentype` elements. We use the type name `halfn` to represent n -element vectors of `half` elements²³. The suffix n is also used in the function names (i.e. **vloadn**, **vstoren** etc.), where $n = 2, 3, 4, 8$ or 16 .

Defined in header `<opencl_vector_load_store>`

Function	Description
<code>gentypen vloadn</code> (<code>size_t offset</code> , <code>const gentype *p</code>)	Return <code>sizeof (gentypen)</code> bytes of data read from address $(p + (offset * n))$. The address computed as $(p + (offset * n))$ must be 8-bit aligned if <code>gentype</code> is <code>char</code> , <code>uchar</code> ; 16-bit aligned if <code>gentype</code> is <code>short</code> , <code>ushort</code> , <code>half</code> ; 32-bit aligned if <code>gentype</code> is <code>int</code> , <code>uint</code> , <code>float</code> ; 64-bit aligned if <code>gentype</code> is <code>long</code> , <code>ulong</code> .
<code>void vstoren</code> (<code>gentypen data</code> , <code>size_t offset</code> , <code>gentype *p</code>)	Write <code>sizeof (gentypen)</code> bytes given by <code>data</code> to address $(p + (offset * n))$. The address computed as $(p + (offset * n))$ must be 8-bit aligned if <code>gentype</code> is <code>char</code> , <code>uchar</code> ; 16-bit aligned if <code>gentype</code> is <code>short</code> , <code>ushort</code> , <code>half</code> ; 32-bit aligned if <code>gentype</code> is <code>int</code> , <code>uint</code> , <code>float</code> ; 64-bit aligned if <code>gentype</code> is <code>long</code> , <code>ulong</code> .
<code>float vload_half</code> (<code>size_t offset</code> , <code>const half *p</code>)	Read <code>sizeof (half)</code> bytes of data from address $(p + offset)$. The data read is interpreted as a <code>half</code> value. The <code>half</code> value is converted to a <code>float</code> value and the <code>float</code> value is returned. The read address computed as $(p + offset)$ must be 16-bit aligned.
<code>floatn vload_halfn</code> (<code>size_t offset</code> , <code>const half *p</code>)	Read <code>sizeof (halfn)</code> bytes of data from address $(p + (offset * n))$. The data read is interpreted as a <code>halfn</code> value. The <code>halfn</code> value read is converted to a <code>floatn</code> value and the <code>floatn</code> value is returned. The read address computed as $(p + (offset * n))$ must be 16-bit aligned.
<code>void vstore_half</code> (<code>float data</code> ,	The <code>float</code> value given by <code>data</code> is first converted

²³ The `halfn` type is only defined by the `cl_khr_fp16` extension described in section 9.5 of the OpenCL 1.2 Extension Specification.

<code>size_t offset, half *p)</code> void vstore_half_rte (float <i>data</i> , <code>size_t offset, half *p)</code> void vstore_half_rtz (float <i>data</i> , <code>size_t offset, half *p)</code> void vstore_half_rtp (float <i>data</i> , <code>size_t offset, half *p)</code> void vstore_half_rtn (float <i>data</i> , <code>size_t offset, half *p)</code>	<p>to a half value using the appropriate rounding mode. The half value is then written to address computed as $(p + \text{offset})$. The address computed as $(p + \text{offset})$ must be 16-bit aligned.</p> <p>vstore_half uses the default rounding mode. The default rounding mode is round to nearest even.</p>
void vstore_halfn (floatn <i>data</i> , <code>size_t offset, half *p)</code> void vstore_halfn_rte (floatn <i>data</i> , <code>size_t offset, half *p)</code> void vstore_halfn_rtz (floatn <i>data</i> , <code>size_t offset, half *p)</code> void vstore_halfn_rtp (floatn <i>data</i> , <code>size_t offset, half *p)</code> void vstore_halfn_rtn (floatn <i>data</i> , <code>size_t offset, half *p)</code>	<p>The floatn value given by <i>data</i> is converted to a halfn value using the appropriate rounding mode. The halfn value is then written to address computed as $(p + (\text{offset} * n))$. The address computed as $(p + (\text{offset} * n))$ must be 16-bit aligned.</p> <p>vstore_halfn uses the default rounding mode. The default rounding mode is round to nearest even.</p>
void vstore_half (double <i>data</i> , <code>size_t offset, half *p)</code> void vstore_half_rte (double <i>data</i> , <code>size_t offset, half *p)</code> void vstore_half_rtz (double <i>data</i> , <code>size_t offset, half *p)</code> void vstore_half_rtp (double <i>data</i> , <code>size_t offset, half *p)</code> void vstore_half_rtn (double <i>data</i> , <code>size_t offset, half *p)</code>	<p>The double value given by <i>data</i> is first converted to a half value using the appropriate rounding mode. The half value is then written to address computed as $(p + \text{offset})$. The address computed as $(p + \text{offset})$ must be 16-bit aligned.</p> <p>vstore_half use the default rounding mode. The default rounding mode is round to nearest even.</p>
void vstore_halfn (doublen <i>data</i> , <code>size_t offset, half *p)</code> void vstore_halfn_rte (doublen <i>data</i> , <code>size_t offset, half *p)</code> void vstore_halfn_rtz (doublen <i>data</i> , <code>size_t offset, half *p)</code> void vstore_halfn_rtp (doublen <i>data</i> , <code>size_t offset, half *p)</code> void vstore_halfn_rtn (doublen <i>data</i> , <code>size_t offset, half *p)</code>	<p>The doublen value given by <i>data</i> is converted to a halfn value using the appropriate rounding mode. The halfn value is then written to address computed as $(p + (\text{offset} * n))$. The address computed as $(p + (\text{offset} * n))$ must be 16-bit aligned.</p> <p>vstore_halfn uses the default rounding mode. The default rounding mode is round to nearest even.</p>
floatn vloada_halfn (size_t <i>offset</i> ,	For n = 2, 4, 8 and 16 read sizeof (halfn) bytes

<code>const half *p)</code>	<p>of data from address $(p + (\text{offset} * n))$. The data read is interpreted as a <i>halfn</i> value. The <i>halfn</i> value read is converted to a <i>floatn</i> value and the <i>floatn</i> value is returned.</p> <p>The address computed as $(p + (\text{offset} * n))$ must be aligned to <code>sizeof (halfn)</code> bytes.</p> <p>For $n = 3$, vloada_half3 reads a <i>half3</i> from address $(p + (\text{offset} * 4))$ and returns a <i>float3</i>. The address computed as $(p + (\text{offset} * 4))$ must be aligned to <code>sizeof (half) * 4</code> bytes.</p>
<code>void vstorea_halfn (floatn data, size_t offset, half *p) void vstorea_halfn_rte (floatn data, size_t offset, half *p) void vstorea_halfn_rtz (floatn data, size_t offset, half *p) void vstorea_halfn_rtp (floatn data, size_t offset, half *p) void vstorea_halfn_rtn (floatn data, size_t offset, half *p)</code>	<p>The <i>floatn</i> value given by <i>data</i> is converted to a <i>halfn</i> value using the appropriate rounding mode.</p> <p>For $n = 2, 4, 8$ and 16, the <i>halfn</i> value is written to the address computed as $(p + (\text{offset} * n))$. The address computed as $(p + (\text{offset} * n))$ must be aligned to <code>sizeof (halfn)</code> bytes.</p> <p>For $n = 3$, the <i>half3</i> value is written to the address computed as $(p + (\text{offset} * 4))$. The address computed as $(p + (\text{offset} * 4))$ must be aligned to <code>sizeof (half) * 4</code> bytes.</p> <p>vstorea_halfn uses the default rounding mode. The default rounding mode is round to nearest even.</p>
<code>void vstorea_halfn (doublen data, size_t offset, half *p) void vstorea_halfn_rte (doublen data, size_t offset, half *p) void vstorea_halfn_rtz (doublen data, size_t offset, half *p) void vstorea_halfn_rtp (doublen data, size_t offset, half *p) void vstorea_halfn_rtn (doublen data, size_t offset, half *p)</code>	<p>The <i>doublen</i> value is converted to a <i>halfn</i> value using the appropriate rounding mode.</p> <p>For $n = 2, 4, 8$ or 16, the <i>halfn</i> value is written to the address computed as $(p + (\text{offset} * n))$. The address computed as $(p + (\text{offset} * n))$ must be aligned to <code>sizeof (halfn)</code> bytes.</p> <p>For $n = 3$, the <i>half3</i> value is written to the address computed as $(p + (\text{offset} * 4))$. The address computed as $(p + (\text{offset} * 4))$ must be aligned to <code>sizeof (half) * 4</code> bytes.</p> <p>vstorea_halfn uses the default rounding mode. The default rounding mode is round to nearest</p>

	even.
--	-------

Table 2.7 *Vector Data Load and Store Functions*²⁴

The results of vector data load and store functions are undefined if the address being read from or written to is not correctly aligned as described in *table 2.7*.

²⁴ **vload3**, and **vload_half3** read *x, y, z* components from address $(p + (\text{offset} * 3))$ into a 3-component vector. **vstore3**, and **vstore_half3** write *x, y, z* components from a 3-component vector to address $(p + (\text{offset} * 3))$.

In addition **vloada_half3** reads *x, y, z* components from address $(p + (\text{offset} * 4))$ into a 3-component vector and **vstorea_half3** writes *x, y, z* components from a 3-component vector to address $(p + (\text{offset} * 4))$.

2.8 Synchronization Functions

The OpenCL C++ programming language implements the following synchronization functions.

Defined in header `<opencl_synchronization>`

Function	Description
<pre>void work_group_barrier (mem_type <i>flags</i>) void work_group_barrier (mem_type <i>flags</i>, memory_scope <i>scope</i>²⁵)</pre>	<p>All work-items in a work-group executing the kernel on a processor must execute this function before any are allowed to continue execution beyond the work_group_barrier. This function must be encountered by all work-items in a work-group executing the kernel. These rules apply to ND-ranges implemented with uniform and non-uniform work-groups.</p> <p>If work_group_barrier is inside a conditional statement, then all work-items must enter the conditional if any work-item enters the conditional statement and executes the work_group_barrier.</p> <p>If work_group_barrier is inside a loop, all work-items must execute the work_group_barrier for each iteration of the loop before any are allowed to continue execution beyond the work_group_barrier.</p> <p>The work_group_barrier function also supports a variant that specifies the memory scope. For the work_group_barrier variant that does not take a memory scope, the <i>scope</i> is <code>memory_scope_work_group</code>.</p> <p>The <i>scope</i> argument specifies whether the memory accesses of work-items in the work-group to memory address space(s) identified by <i>flags</i> become visible to all work-items in the work-group, the device or all SVM devices.</p> <p>The work_group_barrier function can also</p>

²⁵ Refer to *section 6.13.11* for description of `memory_scope`.

	<p>be used to specify which memory operations i.e. to global memory, local memory or images become visible to the appropriate memory scope identified by <i>scope</i>. The <i>flags</i> argument specifies the memory address spaces. This is a bitfield and can be set to 0 or a combination of the following values ORed together. When these flags are OR'ed together the work_group_barrier acts as a combined barrier for all address spaces specified by the flags ordering memory accesses both within and across the specified address spaces.</p> <p><i>mem_type_local</i> - The work_group_barrier function will ensure that all local memory accesses become visible to all work-items in the work-group. Note that the value of <i>scope</i> is ignored as the memory scope is always <i>memory_scope_work_group</i>.</p> <p><i>mem_type_global</i> - The work_group_barrier function ensure that all global memory accesses become visible to the appropriate scope as given by <i>scope</i>.</p> <p><i>mem_type_image</i> - The work_group_barrier function will ensure that all image memory accesses become visible to the appropriate scope as given by <i>scope</i>. The value of <i>scope</i> must be <i>memory_scope_work_group</i> or <i>memory_scope_device</i>.</p> <p>The values of <i>flags</i> and <i>scope</i> must be the same for all work-items in the work-group.</p>
<pre>void sub_group_barrier (mem_type flags) void sub_group_barrier (mem_type flags, memory_scope scope)</pre>	<p>All work-items in a sub-group executing the kernel on a processor must execute this function before any are allowed to continue execution beyond the subgroup barrier. This function must be encountered by all work-items in a sub-group executing the kernel. These rules apply to ND-ranges implemented with uniform and non-uniform work-groups.</p>

	<p>If sub_group_barrier is inside a conditional statement, then all work-items within the sub-group must enter the conditional if any work-item in the sub-group enters the conditional statement and executes the subgroup_barrier.</p> <p>If sub_group_barrier is inside a loop, all work-items within the sub-group must execute the sub_group_barrier for each iteration of the loop before any are allowed to continue execution beyond the sub_group_barrier.</p> <p>The sub_group_barrier function also queues a memory fence (reads and writes) to ensure correct ordering of memory operations to local or global memory.</p> <p>The <i>flags</i> argument specifies the memory address spaces. This is a bitfield and can be set to 0 or a combination of the following values ORed together. When these flags are OR'ed together the sub_group_barrier acts as a combined barrier for all address spaces specified by the flags ordering memory accesses both within and across the specified address spaces.</p> <p><code>mem_type_local</code> - The sub_group_barrier function will ensure that all local memory accesses become visible to all work-items in the sub-group. Note that the value of <i>scope</i> is ignored as the memory scope is always <code>memory_scope_work_group</code>.</p> <p><code>mem_type_global</code> - The sub_group_barrier function ensure that all global memory accesses become visible to the appropriate scope as given by <i>scope</i>.</p> <p><code>mem_type_image</code> - The sub_group_barrier function will ensure that all image memory accesses become visible to the appropriate scope as given by <i>scope</i>. The value of <i>scope</i></p>
--	---

Provisional Specification

	<p>must be <code>memory_scope_work_group</code> or <code>memory_scope_device</code>.</p> <p>The values of <i>flags</i> and <i>scope</i> must be the same for all work-items in the sub-group.</p>
--	---

Table 2.8 *Built-in Synchronization Functions*

2.9 Address Space Qualifier Functions

The OpenCL C++ programming language implements the following address space qualifier functions. We use the generic type name `gentype` to indicate any of the built-in data types supported by OpenCL C or a user defined type.

Function	Description
<code>mem_type</code> <code>get_mem_type</code> (<code>gentype *ptr</code>)	Returns the <code>mem_type</code> value for <i>ptr</i> .
<code>mem_type</code> <code>get_mem_type</code> (<code>const gentype *ptr</code>)	

Table 2.9 *Built-in Address Space Qualifier Functions*

2.10 Atomics

The OpenCL C++ programming language implements a subset of the C++14 atomics (refer to *section 29* of the C++14 specification) and synchronization operations. These operations play a special role in making assignments in one work-item visible to another. A synchronization operation on one or more memory locations is either an acquire operation, a release operation, or both an acquire and release operation²⁶. A synchronization operation without an associated memory location is a fence and can be either an acquire fence, a release fence or both an acquire and release fence. In addition, there are relaxed atomic operations, which are not synchronization operations, and atomic read-modify-write operations which have special characteristics.

The types include

`memory_order`

which is an enumerated type whose enumerators identify memory ordering constraints;





`memory_scope`

which is an enumerated type whose enumerators identify scope of memory ordering constraints;

`atomic_flag`

which is a 32-bit integer type representing a lock-free, primitive atomic flag; and several atomic analogs of integer types.

In the following operation definitions:

-  An A refers to one of the atomic types.
-  A C refers to its corresponding non-atomic type.
-  An M refers to the type of the other argument for arithmetic operations. For atomic integer types, M is C.
-  The functions not ending in explicit have the same semantics as the corresponding explicit function with `memory_order_seq_cst` for the `memory_order` argument.

²⁶ The C++14 consume operation is not supported.

- ✚ The functions that do not have `memory_scope` argument have the same semantics as the corresponding functions with the `memory_scope` argument set to `memory_scope_device`.

NOTE: With fine-grained system SVM, sharing happens at the granularity of individual loads and stores anywhere in host memory. Memory consistency is always guaranteed at synchronization points, but to obtain finer control over consistency, the OpenCL atomics functions may be used to ensure that the updates to individual data values made by one unit of execution are visible to other execution units. In particular, when a host thread needs fine control over the consistency of memory that is shared with one or more OpenCL devices, it must use atomic and fence operations that are compatible with the C++14 atomic operations²⁷.

The atomic functions are defined in header `<opencl_atomic>`.

2.10.1.1 The `ATOMIC_VAR_INIT` macro

The `ATOMIC_VAR_INIT` macro expands to a token sequence suitable for initializing an atomic object of a type that is initialization-compatible with value. An atomic object with automatic storage duration that is not explicitly initialized using `ATOMIC_VAR_INIT` is initially in an indeterminate state; however, the default (zero) initialization for objects with static storage duration is guaranteed to produce a valid state.

```
#define ATOMIC_VAR_INIT(C value)
```

This macro can only be used to initialize atomic objects that are declared in program scope in the global address space.

Examples:

```
atomic_int guide = ATOMIC_VAR_INIT(42);
```

Concurrent access to the variable being initialized, even via an atomic operation, constitutes a data-race.

2.10.1.2 The `atomic_init` function

²⁷ We can't require C++14 atomics since host programs can be implemented in other programming languages and versions of C or C++, but we do require that the host programs use atomics and that those atomics be compatible with those in C++14.

The `atomic_init` function non-atomically initializes the atomic object pointed to by `obj` to the value `value`.

```
void atomic_init(volatile A *obj, C value)
```

Examples:

```
local atomic_int local_guide;
if (get_local_id(0) == 0)
    atomic_init(&guide, 42);
work_group_barrier(CLK_LOCAL_MEM_FENCE);
```

2.10.1.3 Order and Consistency

The enumerated type `memory_order` specifies the detailed regular (non-atomic) memory synchronization operations and may provide for operation ordering. Its enumeration constants are as follows:

```
memory_order_relaxed
memory_order_acquire
memory_order_release
memory_order_acq_rel
memory_order_seq_cst
```

The `memory_order` can be used when performing atomic operations to global or local memory.

2.10.1.4 Memory Scope

The enumerated type `memory_scope` specifies whether the memory ordering constraints given by `memory_order` apply to work-items in a work-group or work-items of a kernel(s) executing on the device or across devices (in the case of shared virtual memory). Its enumeration constants are as follows:

```
memory_scope_work_item28
memory_scope_sub_group
memory_scope_work_group
memory_scope_device
memory_scope_all_svm_devices
```

The memory scope should only be used when performing atomic operations to

²⁸ This value for `memory_scope` can only be used with `atomic_work_item_fence` with flags set to `CLK_IMAGE_MEM_FENCE`.

global memory. Atomic operations to local memory only guarantee memory ordering in the work-group not across work-groups and therefore ignore the `memory_scope` value.

2.10.1.5 Fences

The following fence operations are supported.

The `atomic_work_item_fence` function

```
void atomic_work_item_fence(mem_type flags,
                           memory_order order,
                           memory_scope scope)
```

`flags` must be set to `mem_type_global`, `mem_type_local`, `mem_type_image` or a combination of these values ORed together; otherwise the behavior is undefined. The behavior of calling `atomic_work_item_fence` with `mem_type_global`, `mem_type_local` or `mem_type_image` ORed together is equivalent to calling `atomic_work_item_fence` individually for each of the fence values set in `flags`.

Depending on the value of `order`, this operation:

- ✚ has no effects, if `order == memory_order_relaxed`.
- ✚ is an acquire fence, if `order == memory_order_acquire`.
- ✚ is a release fence, if `order == memory_order_release`.
- ✚ is both an acquire fence and a release fence, if `order == memory_order_acq_rel`.
- ✚ is a sequentially consistent acquire and release fence, if `order == memory_order_seq_cst`.

For images declared with the `read_write` qualifier, the `atomic_work_item_fence` must be called to make sure that writes to the image by a work-item become visible to that work-item on subsequent reads to that image by that work-item.

2.10.1.6 Atomic integer and floating-point types

The list of supported atomic type names are:

```
atomic_int
atomic_uint
```

```

atomic_long29
atomic_ulong
atomic_float
atomic_double30
atomic_intptr_t31
atomic_uintptr_t
atomic_size_t
atomic_ptrdiff_t

```

Arguments to a kernel can be declared to be a pointer to the above atomic types or the `atomic_flag` type.

The representation of atomic integer, floating-point and pointer types have the same size as their corresponding regular types. The `atomic_flag` type must be implemented as a 32-bit integer.

2.10.1.7 Operations on atomic types

There are only a few kinds of operations on atomic types, though there are many instances of those kinds. This section specifies each general kind.

The `atomic_store` functions

```

void atomic_store(volatile A* object, C desired);
void atomic_store(A* object, C desired);

void atomic_store_explicit(volatile A* object,
                           C desired,
                           memory_order order);
void atomic_store_explicit(A* object,
                           C desired,
                           memory_order order);

void atomic_store_explicit(volatile A* object,
                           C desired,
                           memory_order order,
                           memory_scope scope);
void atomic_store_explicit(A* object,

```

²⁹ The `atomic_long` and `atomic_ulong` types are supported if the `cl_khr_int64_base_atomics` and `cl_khr_int64_extended_atomics` extensions are supported.

³⁰ The `atomic_double` type is only supported if double precision is supported and the `cl_khr_int64_base_atomics` and `cl_khr_int64_extended_atomics` extensions are supported.

³¹ If the device address space is 64-bits, the data types `atomic_intptr_t`, `atomic_uintptr_t`, `atomic_size_t` and `atomic_ptrdiff_t` are supported if the `cl_khr_int64_base_atomics` and `cl_khr_int64_extended_atomics` extensions are supported.

Provisional Specification

```
C desired,  
memory_order order,  
memory_scope scope);  
  
void A::store(C desired,  
             memory_order order = memory_order_seq_cst,  
             memory_scope scope = memory_scope_device)  
             volatile;  
  
void A::store(C desired,  
             memory_order order = memory_order_seq_cst,  
             memory_scope scope = memory_scope_device);
```

The order argument shall not be `memory_order_acquire`, nor `memory_order_acq_rel`. Atomically replace the value pointed to by object with the value of `desired`. Memory is affected according to the value of `order`.

The `atomic_load` functions

```
C atomic_load(const volatile A* object);  
C atomic_load(const A* object);  
  
C atomic_load_explicit(const volatile A* object,  
                      memory_order order);  
C atomic_load_explicit(const A* object,  
                      memory_order order);  
  
C atomic_load_explicit(const volatile A* object,  
                      memory_order order,  
                      memory_scope scope);  
C atomic_load_explicit(const A* object,  
                      memory_order order,  
                      memory_scope scope);  
  
C A::load(memory_order order = memory_order_seq_cst,  
         memory_scope scope = memory_scope_device)  
         const volatile;  
C A::load(memory_order order = memory_order_seq_cst,  
         memory_scope scope = memory_scope_device)  
         const;
```

The order argument shall not be `memory_order_release` nor `memory_order_acq_rel`. Memory is affected according to the value of `order`. Atomically returns the value pointed to by object.

The `atomic_exchange` functions

Provisional Specification

```
C atomic_exchange(volatile A* object, C desired);
C atomic_exchange(A* object, C desired);

C atomic_exchange_explicit(volatile A* object,
                           C desired,
                           memory_order order);
C atomic_exchange_explicit(A* object,
                           C desired,
                           memory_order order);

C atomic_exchange_explicit(volatile A* object,
                           C desired,
                           memory_order order,
                           memory_scope scope);
C atomic_exchange_explicit(A* object,
                           C desired,
                           memory_order order,
                           memory_scope scope);

C A::exchange(C desired,
              memory_order order = memory_order_seq_cst,
              memory_scope scope = memory_scope_device)
              volatile;

C A::exchange(C desired,
              memory_order order = memory_order_seq_cst,
              memory_scope scope = memory_scope_device);
```

Atomically replace the value pointed to by `object` with `desired`. Memory is affected according to the value of `order`. These operations are read-modify-write operations. `Atomic` returns the value pointed to by `object` immediately before the effects.

The `atomic_compare_exchange` functions

```
bool atomic_compare_exchange_strong(
    volatile A* object,
    C* expected,
    C desired);
bool atomic_compare_exchange_strong(
    A* object,
    C* expected,
    C desired);

bool atomic_compare_exchange_strong_explicit(
    volatile A *object,
```

Provisional Specification

```
        C* expected,
        C desired,
        memory_order success,
        memory_order failure);
bool atomic_compare_exchange_strong_explicit(
    A* object,
    C* expected,
    C desired,
    memory_order success,
    memory_order failure);

bool atomic_compare_exchange_strong_explicit(
    volatile A *object,
    C* expected,
    C desired,
    memory_order success,
    memory_order failure,
    memory_scope scope);
bool atomic_compare_exchange_strong_explicit(
    A* object,
    C* expected,
    C desired,
    memory_order success,
    memory_order failure,
    memory_scope scope);

bool atomic_compare_exchange_weak(
    volatile A* object,
    C* expected,
    C desired)
bool atomic_compare_exchange_weak(
    A* object,
    C* expected,
    C desired)

bool atomic_compare_exchange_weak_explicit(
    volatile A* object,
    C* expected,
    C desired,
    memory_order success,
    memory_order failure);
bool atomic_compare_exchange_weak_explicit(
    A* object,
    C* expected,
    C desired,
    memory_order success,
    memory_order failure);
```


Provisional Specification

```
bool atomic_compare_exchange_weak_explicit(
    volatile A* object,
    C* expected,
    C desired,
    memory_order success,
    memory_order failure,
    memory_scope scope);
bool atomic_compare_exchange_weak_explicit(
    A* object,
    C* expected,
    C desired,
    memory_order success,
    memory_order failure,
    memory_scope scope);

bool A::compare_exchange_strong(
    C& expected, C desired) volatile;
bool A::compare_exchange_strong(
    C& expected, C desired);

bool A::compare_exchange_strong(
    C& expected, C desired,
    memory_order success,
    memory_order failure,
    memory_scope scope =
        memory_scope_device) volatile;
bool A::compare_exchange_strong(
    C& expected, C desired,
    memory_order success,
    memory_order failure,
    memory_scope scope =
        memory_scope_device);

bool A::compare_exchange_weak(
    C& expected, C desired) volatile;
bool A::compare_exchange_weak(
    C& expected, C desired);

bool A::compare_exchange_weak(
    C& expected, C desired,
    memory_order success,
    memory_order failure,
    memory_scope scope =
        memory_scope_device) volatile;
bool A::compare_exchange_weak(
    C& expected, C desired,
```

Provisional Specification

```
memory_order success,  
memory_order failure,  
memory_scope scope =  
    memory_scope_device);
```

The failure argument shall not be `memory_order_release` nor `memory_order_acq_rel`. The failure argument shall be no stronger than the success argument. Atomically, compares the value pointed to by `object` for equality with that in `expected`, and if true, replaces the value pointed to by `object` with `desired`, and if false, updates the value in `expected` with the value pointed to by `object`. Further, if the comparison is true, memory is affected according to the value of `success`, and if the comparison is false, memory is affected according to the value of `failure`. These operations are atomic read-modify-write operations.

NOTE: The effect of the compare-and-exchange operations is

```
if (memcmp(object, expected, sizeof(*object)) == 0)  
    memcpy(object, &desired, sizeof(*object));  
else  
    memcpy(expected, object, sizeof(*object));
```

The weak compare-and-exchange operations may fail spuriously³². That is, even when the contents of memory referred to by `expected` and `object` are equal, it may return zero and store back to `expected` the same memory contents that were originally there.

These generic functions return the result of the comparison.

The `atomic_fetch` and `modify` functions

The following operations perform arithmetic and bitwise computations. All of these operations are applicable to an object of any atomic integer type. The key, operator, and computation correspondence is given in table below:

<i>key</i>	<i>op</i>	<i>computation</i>
add	+	addition
sub	-	subtraction
or		bitwise inclusive or
xor	^	bitwise exclusive or
and	&	bitwise and

³² This spurious failure enables implementation of compare-and-exchange on a broader class of machines, e.g. load-locked store-conditional machines.

Provisional Specification

min	min	compute min
max	max	compute max

NOTE: For `atomic_fetch` and `modify` functions with `key = add` or `sub` on atomic types `atomic_intptr_t` and `atomic_uintptr_t`, `Mis ptrdiff_t`. For `atomic_fetch` and `modify` functions with `key = or, xor, and, min` and `max` on atomic types `atomic_intptr_t` and `atomic_uintptr_t`, `Mis intptr_t` and `uintptr_t`.

```

C atomic_fetch_key(volatile A *object, M operand);
C atomic_fetch_key(A *object, M operand);

C atomic_fetch_key_explicit(
    volatile A *object,
    M operand,
    memory_order order);
C atomic_fetch_key_explicit(
    A *object,
    M operand,
    memory_order order);

C atomic_fetch_key_explicit(
    volatile A *object,
    M operand,
    memory_order order,
    memory_scope scope);
C atomic_fetch_key_explicit(
    A *object,
    M operand,
    memory_order order,
    memory_scope scope);

C A::fetch_key(M operand,
    memory_order order = memory_order_seq_cst,
    memory_scope scope =
        memory_scope_device) volatile;
C A::fetch_key(M operand,
    memory_order order = memory_order_seq_cst,
    memory_scope scope = memory_scope_device);

```

Atomically replaces the value pointed to by `object` with the result of the computation applied to the value pointed to by `object` and the given operand. Memory is affected according to the value of `order`. These operations are atomic read-modify-write operations. For signed integer types, arithmetic is defined to use two's complement representation with silent wrap-around on overflow; there are no undefined results. For address types, the result may be an undefined address, but

the operations otherwise have no undefined behavior. Returns atomically, the value pointed to by `object` immediately before the effects.

Operators

The following operators are supported for atomic types:

```
C A::operator=(C desired) volatile;  
C A::operator=(C desired);
```

Effects: `store(desired)`

Returns: `desired`

```
A::operator C() const volatile;  
A::operator C() const;
```

Effects: `load()`

Returns: The result of `load()`

```
C A::operator op=(M operand) volatile;  
C A::operator op=(M operand);
```

Effects: `fetch_key(operand)`

Returns: `fetch_key(operand) op operand`

```
C A::operator op++(int) volatile;  
C A::operator op++(int);
```

Returns: `fetch_add(int)`

```
C A::operator op--(int) volatile;  
C A::operator op--(int);
```

Returns: `fetch_sub(int)`

```
C A::operator op++() volatile;  
C A::operator op++();
```

Effects: `fetch_add(1)`

Returns: `fetch_add(1) + 1`

```
C A::operator op--() volatile;  
C A::operator op--();
```

Effects: `fetch_sub(1)`

Returns: `fetch_sub(1) - 1`

Atomic flag type and operations

The `atomic_flag` type provides the classic test-and-set functionality. It has two states, set (value is non-zero) and clear (value is 0). Operations on an object of type `atomic_flag` shall be lock free.

The macro `ATOMIC_FLAG_INIT` may be used to initialize an `atomic_flag` to the clear state. An `atomic_flag` that is not explicitly initialized with `ATOMIC_FLAG_INIT` is initially in an indeterminate state.

This macro can only be used for atomic objects that are declared in program scope in the global address space with the `atomic_flag` type.

Example:

```
atomic_flag guard = ATOMIC_FLAG_INIT;
```

The `atomic_flag_test_and_set` functions

```
bool atomic_flag_test_and_set(
    volatile atomic_flag* object);
bool atomic_flag_test_and_set(
    atomic_flag* object);

bool atomic_flag_test_and_set_explicit(
    volatile atomic_flag* object,
    memory_order order);
bool atomic_flag_test_and_set_explicit(
    atomic_flag* object,
    memory_order order);

bool atomic_flag_test_and_set_explicit(
    volatile atomic_flag* object,
    memory_order order,
    memory_scope scope);
bool atomic_flag_test_and_set_explicit(
    atomic_flag* object,
    memory_order order,
    memory_scope scope);

bool atomic_flag::test_and_set(
    memory_order order =
        memory_order_seq_cst,
```

Provisional Specification

```
memory_scope scope =  
    memory_order_device) volatile;  
bool atomic_flag::test_and_set(  
    memory_order order =  
        memory_order_seq_cst,  
    memory_scope scope =  
        memory_order_device);
```

Atomically sets the value pointed to by `object` to true. Memory is affected according to the value of `order`. These operations are atomic read-modify-write operations. Returns atomically, the value of the object immediately before the effects.

The `atomic_flag_clear` functions

```
void atomic_flag_clear(volatile atomic_flag *object);  
void atomic_flag_clear(atomic_flag *object);  
  
void atomic_flag_clear_explicit(  
    volatile atomic_flag *object,  
    memory_order order);  
void atomic_flag_clear_explicit(  
    atomic_flag *object,  
    memory_order order);  
  
void atomic_flag_clear_explicit(  
    volatile atomic_flag *object,  
    memory_order order,  
    memory_scope scope);  
void atomic_flag_clear_explicit(  
    atomic_flag *object,  
    memory_order order,  
    memory_scope scope);  
  
void atomic_flag::clear(  
    memory_order order = memory_order_seq_cst,  
    memory_scope scope = memory_scope_device)  
    volatile;  
  
void atomic_flag::clear(  
    memory_order order = memory_order_seq_cst,  
    memory_scope scope = memory_scope_device);
```

The `order` argument shall not be `memory_order_acquire` nor `memory_order_acq_rel`. Atomically sets the value pointed to by `object` to false. Memory is affected according to the value of `order`.

2.10.1.8 Restrictions

- ✚ The generic `atomic<T>` class template is only available if `T` is `int`, `uint`, `long`³³, `ulong`, `float`, `double`³⁴, `intptr_t`³⁵, `atomic_uintptr_t`, `atomic_size_t`, `atomic_ptrdiff_t`.
- ✚ The `atomic_bool`, `atomic_char`, `atomic_uchar`, `atomic_short`, `atomic_ushort`, `atomic_intmax_t` and `atomic_uintmax_t` types are not supported by OpenCL C++.
- ✚ OpenCL C++ requires that the built-in atomic functions on atomic types are lock-free.
- ✚ The atomic data types cannot be declared inside a kernel or non-kernel function unless they are declared as `static` or as declared using the `local_array` or `local` templated types.

³³ The `atomic_long` and `atomic_ulong` types are supported if the `cl_khr_int64_base_atomics` and `cl_khr_int64_extended_atomics` extensions are supported.

³⁴ The `atomic_double` type is only supported if double precision is supported and the `cl_khr_int64_base_atomics` and `cl_khr_int64_extended_atomics` extensions are supported.

³⁵ If the device address space is 64-bits, the data types `atomic_intptr_t`, `atomic_uintptr_t`, `atomic_size_t` and `atomic_ptrdiff_t` are supported if the `cl_khr_int64_base_atomics` and `cl_khr_int64_extended_atomics` extensions are supported.

2.11 printf

The OpenCL C++ programming language implements the **printf**³⁶ function. This function is defined in header `<opencl_printf>`.

Function	Description
int printf (char * restrict <i>format</i> , ...)	<p>The printf built-in function writes output to an implementation-defined stream such as stdout under control of the string pointed to by <i>format</i> that specifies how subsequent arguments are converted for output. If there are insufficient arguments for the format, the behavior is undefined. If the format is exhausted while arguments remain, the excess arguments are evaluated (as always) but are otherwise ignored. The printf function returns when the end of the format string is encountered.</p> <p>printf returns 0 if it was executed successfully and -1 otherwise.</p> <p>Restriction: Format must be constexpr i.e. a string known at compile time; otherwise it will be a compile time error.</p>

Table 2.11 Built-in printf Function

2.11.1 printf output synchronization

When the event that is associated with a particular kernel invocation is completed, the output of all `printf()` calls executed by this kernel invocation is flushed to the implementation-defined output stream. Calling **clFinish** on a command queue flushes all pending output by `printf` in previously enqueued and completed commands to the implementation-defined output stream. In the case that `printf` is executed from multiple work-items concurrently, there is no guarantee of ordering with respect to written data. For example, it is valid for the output of a work-item with a global id (0,0,1) to appear intermixed with the output of a work-item with a global id (0,0,4) and so on.

³⁶ The primary purpose of the **printf** function is to help aid in debugging OpenCL kernels.

2.11.2 printf format string

The format shall be a character sequence, beginning and ending in its initial shift state. The format is composed of zero or more directives: ordinary characters (not %), which are copied unchanged to the output stream; and conversion specifications, each of which results in fetching zero or more subsequent arguments, converting them, if applicable, according to the corresponding conversion specifier, and then writing the result to the output stream. The format is in the `constant` address space and must be resolvable at compile time i.e. cannot be dynamically created by the executing program, itself.

Each conversion specification is introduced by the character %. After the %, the following appear in sequence:

- ✚ Zero or more *flags* (in any order) that modify the meaning of the conversion specification.
- ✚ An optional minimum *field width*. If the converted value has fewer characters than the field width, it is padded with spaces (by default) on the left (or right, if the left adjustment flag, described later, has been given) to the field width. The field width takes the form of a nonnegative decimal integer.³⁷⁾
- ✚ An optional *precision* that gives the minimum number of digits to appear for the **d**, **i**, **o**, **u**, **x**, and **X** conversions, the number of digits to appear after the decimal-point character for **a**, **A**, **e**, **E**, **f**, and **F** conversions, the maximum number of significant digits for the **g** and **G** conversions, or the maximum number of bytes to be written for **s** conversions. The precision takes the form of a period (.) followed by an optional decimal integer; if only the period is specified, the precision is taken as zero. If a precision appears with any other conversion specifier, the behavior is undefined.
- ✚ An optional *vector specifier*.
- ✚ A *length modifier* that specifies the size of the argument. The *length modifier* is required with a vector specifier and together specifies the vector type. Implicit conversions between vector types are disallowed (as per *section 6.2.1*). If the *vector specifier* is not specified, the *length modifier* is optional.
- ✚ A *conversion specifier* character that specifies the type of conversion to be applied.

The flag characters and their meanings are:

- The result of the conversion is left-justified within the field. (It is right-

³⁷⁾ Note that **0** is taken as a flag, not as the beginning of a field width.

- justified if this flag is not specified.)
- +** The result of a signed conversion always begins with a plus or minus sign. (It begins with a sign only when a negative value is converted if this flag is not specified.)³⁸
 - space** If the first character of a signed conversion is not a sign, or if a signed conversion results in no characters, a space is prefixed to the result. If the *space* and **+** flags both appear, the *space* flag is ignored.
 - #** The result is converted to an “alternative form”. For **o** conversion, it increases the precision, if and only if necessary, to force the first digit of the result to be a zero (if the value and precision are both 0, a single 0 is printed). For **x** (or **X**) conversion, a nonzero result has **0x** (or **0X**) prefixed to it. For **a**, **A**, **e**, **E**, **f**, **F**, **g**, and **G** conversions, the result of converting a floating-point number always contains a decimal-point character, even if no digits follow it. (Normally, a decimal-point character appears in the result of these conversions only if a digit follows it.) For **g** and **G** conversions, trailing zeros are *not* removed from the result. For other conversions, the behavior is undefined.
 - 0** For **d**, **i**, **o**, **u**, **x**, **X**, **a**, **A**, **e**, **E**, **f**, **F**, **g**, and **G** conversions, leading zeros (following any indication of sign or base) are used to pad to the field width rather than performing space padding, except when converting an infinity or NaN. If the **0** and **-** flags both appear, the **0** flag is ignored. For **d**, **i**, **o**, **u**, **x**, and **X** conversions, if a precision is specified, the **0** flag is ignored. For other conversions, the behavior is undefined.

The vector specifier and its meaning is:

- vn** Specifies that a following **a**, **A**, **e**, **E**, **f**, **F**, **g**, **G**, **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to a vector argument, where *n* is the size of the vector and must be 2, 3, 4, 8 or 16.

The vector value is displayed in the following general form:
value1 C value2 C C valuen

where C is a separator character. The value for this separator character is a comma.

If the vector specifier is not used, the length modifiers and their meanings are:

- hh** Specifies that a following **d**, **i**, **o**, **u**, **x**, or **X** conversion specifier applies to

³⁸ The results of all floating conversions of a negative zero, and of negative values that round to zero, include a minus sign.

a **char** or **uchar** argument (the argument will have been promoted according to the integer promotions, but its value shall be converted to **char** or **uchar** before printing).

- h** Specifies that a following **d, i, o, u, x,** or **X** conversion specifier applies to a **short** or **ushort** argument (the argument will have been promoted according to the integer promotions, but its value shall be converted to **short** or **unsigned short** before printing).
- l (ell)** Specifies that a following **d, i, o, u, x,** or **X** conversion specifier applies to a **long** or **ulong** argument. The **l** modifier is supported by the full profile. For the embedded profile, the **l** modifier is supported only if 64-bit integers are supported by the device.

If the vector specifier is used, the length modifiers and their meanings are:

- hh** Specifies that a following **d, i, o, u, x,** or **X** conversion specifier applies to a **charn** or **ucharn** argument (the argument will not be promoted).
- h** Specifies that a following **d, i, o, u, x,** or **X** conversion specifier applies to a **shortn** or **ushortn** argument (the argument will not be promoted); that a following **a, A, e, E, f, F, g,** or **G** conversion specifier applies to a **halfn** argument.
- hl** This modifier can only be used with the vector specifier. Specifies that a following **d, i, o, u, x,** or **X** conversion specifier applies to a **intn** or **uintn** argument; that a following **a, A, e, E, f, F, g,** or **G** conversion specifier applies to a **floatn** argument.
- l (ell)** Specifies that a following **d, i, o, u, x,** or **X** conversion specifier applies to a **longn** or **ulongn** argument; that a following **a, A, e, E, f, F, g,** or **G** conversion specifier applies to a **doublen** argument. The **l** modifier is supported by the full profile. For the embedded profile, the **l** modifier is supported only if 64-bit integers or double-precision floating-point are supported by the device.

If a vector specifier appears without a length modifier, the behavior is undefined. The vector data type described by the vector specifier and length modifier must match the data type of the argument; otherwise the behavior is undefined.

If a length modifier appears with any conversion specifier other than as specified above, the behavior is undefined.

The conversion specifiers and their meanings are:

- d,i** The **int**, **charn**, **shortn**, **intn** or **longn** argument is converted to signed

decimal in the style *[-]dddd*. The precision specifies the minimum number of digits to appear; if the value being converted can be represented in fewer digits, it is expanded with leading zeros. The default precision is 1. The result of converting a zero value with a precision of zero is no characters.

o,u,

x,X

The **unsigned int**, **ucharn**, **ushortn**, **uintn** or **ulongn** argument is converted to unsigned octal (**o**), unsigned decimal (**u**), or unsigned hexadecimal notation (**x** or **X**) in the style *dddd*; the letters **abcdef** are used for **x** conversion and the letters **ABCDEF** for **X** conversion. The precision specifies the minimum number of digits to appear; if the value being converted can be represented in fewer digits, it is expanded with leading zeros. The default precision is 1. The result of converting a zero value with a precision of zero is no characters.

f,F

A **double**, **halfn**, **floatn** or **doublen** argument representing a floating-point number is converted to decimal notation in the style *[-]ddd.ddd*, where the number of digits after the decimal-point character is equal to the precision specification. If the precision is missing, it is taken as 6; if the precision is zero and the **#** flag is not specified, no decimal-point character appears. If a decimal-point character appears, at least one digit appears before it. The value is rounded to the appropriate number of digits. A **double**, **halfn**, **floatn** or **doublen** argument representing an infinity is converted in one of the styles *[-]inf* or *[-]infinity* — which style is implementation-defined. A **double**, **halfn**, **floatn** or **doublen** argument representing a NaN is converted in one of the styles *[-]nan* or *[-]nan(n-char-sequence)* — which style, and the meaning of any *n-char-sequence*, is implementation-defined. The **F** conversion specifier produces **INF**, **INFINITY**, or **NAN** instead of **inf**, **infinity**, or **nan**, respectively.³⁹⁾

e,E

A **double**, **halfn**, **floatn** or **doublen** argument representing a floating-point number is converted in the style *[-]d.ddd e±dd*, where there is one digit (which is nonzero if the argument is nonzero) before the decimal-point character and the number of digits after it is equal to the precision; if the precision is missing, it is taken as 6; if the precision is zero and the **#** flag is not specified, no decimal-point character appears. The value is rounded to the appropriate number of digits. The **E** conversion specifier produces a number with **E** instead of **e** introducing the exponent. The exponent always contains at least two digits, and only as many more digits as necessary to represent the exponent. If the value is zero, the exponent is zero. A **double**, **halfn**, **floatn** or **doublen** argument representing an infinity or NaN is converted in the style of an **f** or **F** conversion specifier.

³⁹⁾ When applied to infinite and NaN values, the **-**, **+**, and *space* flag characters have their usual meaning; the **#** and **0** flag characters have no effect.

- g,G** A **double**, **halfn**, **floatn** or **doublen** argument representing a floating-point number is converted in style **f** or **e** (or in style **F** or **E** in the case of a **G** conversion specifier), depending on the value converted and the precision. Let P equal the precision if nonzero, 6 if the precision is omitted, or 1 if the precision is zero. Then, if a conversion with style **E** would have an exponent of X : — if $P > X \geq -4$, the conversion is with style **f** (or **F**) and precision $P - (X + 1)$. — otherwise, the conversion is with style **e** (or **E**) and precision $P - 1$. Finally, unless the **#** flag is used, any trailing zeros are removed from the fractional portion of the result and the decimal-point character is removed if there is no fractional portion remaining. A **double**, **halfn**, **floatn** or **doublen e** argument representing an infinity or NaN is converted in the style of an **f** or **F** conversion specifier.
- a,A** A **double**, **halfn**, **floatn** or **doublen** argument representing a floating-point number is converted in the style `[-]0xh.hhhh p±d`, where there is one hexadecimal digit (which is nonzero if the argument is a normalized floating-point number and is otherwise unspecified) before the decimal-point character⁴⁰) and the number of hexadecimal digits after it is equal to the precision; if the precision is missing, then the precision is sufficient for an exact representation of the value; if the precision is zero and the **#** flag is not specified, no decimal point character appears. The letters **abcdef** are used for **a** conversion and the letters **ABCDEF** for **A** conversion. The **A** conversion specifier produces a number with **X** and **P** instead of **x** and **p**. The exponent always contains at least one digit, and only as many more digits as necessary to represent the decimal exponent of 2. If the value is zero, the exponent is zero. A **double**, **halfn**, **floatn** or **doublen** argument representing an infinity or NaN is converted in the style of an **f** or **F** conversion specifier.

NOTE: The conversion specifiers **e,E,g,G,a,A** convert a float or half argument that is a scalar type to a double only if the double data type is supported. If the double data type is not supported, the argument will be a **float** instead of a **double** and the half type will be converted to a float.

- c** The **int** argument is converted to an **unsigned char**, and the resulting character is written.
- s** The argument shall be a literal string.⁴¹ Characters from the literal string array are written up to (but not including) the terminating null character. If the precision is specified, no more than that many bytes are written. If the precision is not specified or is greater than the size of the array, the array shall contain a null character.

⁴⁰ Binary implementations can choose the hexadecimal digit to the left of the decimal-point character so that subsequent digits align to nibble (4-bit) boundaries.

⁴¹ No special provisions are made for multibyte characters. The behavior of **printf** with the **s** conversion specifier is undefined if the argument value is not a pointer to a literal string.

- p** The argument shall be a pointer to **void**. The pointer can refer to a memory region in the `global`, `constant`, `local`, `private` or `generic` address space. The value of the pointer is converted to a sequence of printing characters in an implementation-defined manner.
- %** A **%** character is written. No argument is converted. The complete conversion specification shall be **%%**.

If a conversion specification is invalid, the behavior is undefined. If any argument is not the correct type for the corresponding conversion specification, the behavior is undefined.

In no case does a nonexistent or small field width cause truncation of a field; if the result of a conversion is wider than the field width, the field is expanded to contain the conversion result.

For **a** and **A** conversions, the value is correctly rounded to a hexadecimal floating number with the given precision.

A few examples of `printf` are given below:

```
float4 f = float4(1.0f, 2.0f, 3.0f, 4.0f);
uchar4 uc = uchar4(0xFA, 0xFB, 0xFC, 0xFD);

printf("f4 = %2.2v4hlf\n", f);
printf("uc = %#v4hhx\n", uc);
```

The above two `printf` calls print the following:

```
f4 = 1.00,2.00,3.00,4.00
uc = 0xfa,0xfb,0xfc,0xfd
```

A few examples of valid use cases of `printf` for the conversion specifier **s** are given below. The argument value must be a pointer to a literal string.

```
kernel void my_kernel( ... )
{
    printf("%s\n", "this is a test string\n");
}
```

A few examples of invalid use cases of `printf` for the conversion specifier **s** are given below:

```
kernel void my_kernel(char *s, ... )
{
```

```

        printf("%s\n", s);
    }

```

A few examples of invalid use cases of `printf` where data types given by the vector specifier and length modifier do not match the argument type are given below:

```

kernel void my_kernel(char *s, ... )
{
    uint2 ui = (uint2)(0x12345678, 0x87654321);
    printf("unsigned short value = (%#v2hx)\n", ui)
    printf("unsigned char value = (%#v2hhx)\n", ui)
}

```

2.11.3 Differences between OpenCL C++ and C++14 `printf`

- ✚ The **l** modifier followed by a **c** conversion specifier or **s** conversion specifier is not supported by OpenCL C.
- ✚ The **ll**, **j**, **z**, **t**, and **L** length modifiers are not supported by OpenCL C++ but are reserved.
- ✚ The **n** conversion specifier is not supported by OpenCL C++ but is reserved.
- ✚ OpenCL C++ adds the optional **vn** vector specifier to support printing of vector types.
- ✚ The conversion specifiers **f**, **F**, **e**, **E**, **g**, **G**, **a**, **A** convert a float argument to a double only if the double data type is supported. Refer to the description of `CL_DEVICE_DOUBLE_FP_CONFIG` in *table 4.3*. If the double data type is not supported, the argument will be a float instead of a double.
- ✚ For the embedded profile, the **l** length modifier is supported only if 64-bit integers are supported.
- ✚ In OpenCL C++, **printf** returns 0 if it was executed successfully and -1 otherwise vs. C++14 where **printf** returns the number of characters printed or a negative value if an output or encoding error occurred.
- ✚ In OpenCL C++, the conversion specifier **s** can only be used for arguments that are literal strings.

2.12 Images

This section describes the image and sampler types and functions that can be used to read from and/or write to an image.

2.12.1 Image Types

The following image types are supported:

```
image1d<T, image_access::a=image_access::read>
image1d_buffer<T, image_access::a=image_access::read>
image1d_array<T, image_access::a=image_access::read>
image2d<T, image_access::a=image_access::read>
image2d_array<T, image_access::a=image_access::read>
image3d<T, image_access::a=image_access::read>
image2d_depth<T, image_access::a=image_access::read>
image2d_array_depth<T, image_access::a=image_access::read>
```

T specifies the type of the value returned when reading from an image or the type of the color value specified when writing to an image. For non-depth image types, supported values of T are `float4`, `half4`, `int4` or `uint4`. For depth image types, T must be `float` or `half`.




The image access qualifier `a` specifies how the image will be accessed and is the following enumerated type:

```
enum class image_access { sample,
                          read,
                          write,
                          read_write };
```

2.12.2 Samplers

The image read functions take a sampler argument. The sampler can be passed as an argument to the kernel using **`clSetKernelArg`**, or can be declared in the outermost scope of kernel functions, or it can be a constant variable of type `sampler` declared in the program source.

Sampler variables in a program are declared to be of type `sampler`. A variable of `sampler` type declared in the program source must be initialized with a 32-bit unsigned integer constant, which is interpreted as a bit-field specifying the following properties:

-  Addressing Mode
-  Filter Mode
-  Normalized Coordinates

These properties control how elements of an image object are read by the **image_sample** function.

Samplers can also be declared as global constants in the program source using the following syntax.

```
constexpr sampler    <sampler name> = <value>
```

The sampler fields are described in *table 2.12.1*.

Sampler State	Description
<normalized coords>	<p>Specifies whether the <i>x</i>, <i>y</i> and <i>z</i> coordinates are passed in as normalized or unnormalized values. This must be a literal value and can be one of the following predefined enums:</p> <p>CLK_NORMALIZED_COORDS_TRUE or CLK_NORMALIZED_COORDS_FALSE.</p> <p>The samplers used with an image in multiple calls to sample_image declared in a kernel must use the same value for <normalized coords>.</p>
<addressing mode>	<p>Specifies the image addressing-mode i.e. how out-of-range image coordinates are handled. This must be a literal value and can be one of the following predefined enums:</p> <p>CLK_ADDRESS_MIRRORED_REPEAT - Flip the image coordinate at every integer junction. This addressing mode can only be used with normalized coordinates. If normalized coordinates are not used, this addressing mode may generate image coordinates that are undefined.</p> <p>CLK_ADDRESS_REPEAT – out-of-range image coordinates are wrapped to the valid range. This addressing mode can only be used with normalized coordinates. If normalized coordinates are not used, this addressing mode may generate image coordinates that are undefined.</p>

	<p>CLK_ADDRESS_CLAMP_TO_EDGE – out-of-range image coordinates are clamped to the extent.</p> <p>CLK_ADDRESS_CLAMP⁴² – out-of-range image coordinates will return a border color.</p> <p>CLK_ADDRESS_NONE – for this addressing mode the programmer guarantees that the image coordinates used to sample elements of the image refer to a location inside the image; otherwise the results are undefined.</p> <p>For 1D and 2D image arrays, the addressing mode applies only to the x and (x, y) coordinates. The addressing mode for the coordinate which specifies the array index is always CLK_ADDRESS_CLAMP_TO_EDGE.</p>
<filter mode>	<p>Specifies the filter mode to use. This must be a literal value and can be one of the following predefined enums: CLK_FILTER_NEAREST or CLK_FILTER_LINEAR.</p> <p>Refer to <i>section 4.2</i> in the OpenCL API specification for a description of these filter modes.</p>
<mipfilter mode>	<p>Specifies the mipmap filter mode to use. This must be a literal value and can be one of the following predefined enums: CLK_MIP_FILTER_NONE, CLK_MIP_FILTER_NEAREST or CLK_MIP_FILTER_LINEAR.</p> <p>Refer to <i>section 4.2</i> in the OpenCL API specification for a description of these filter modes.</p>

Table 2.12.1 *Sampler Descriptor*

Examples:

```
constexpr sampler samplerA =
    CLK_NORMALIZED_COORDS_TRUE |
    CLK_ADDRESS_REPEAT          |
    CLK_FILTER_NEAREST          |
    CLK_MIP_FILTER_NONE;
```

⁴² This is similar to the GL_ADDRESS_CLAMP_TO_BORDER addressing mode.

samplerA specifies a sampler that uses normalized coordinates, the repeat addressing mode and a nearest filter.

The maximum number of samplers that can be declared in a kernel can be queried using the CL_DEVICE_MAX_SAMPLERS token in **clGetDeviceInfo**.

2.12.2.1 Determining the border color or value

If <addressing mode> in sampler is CLK_ADDRESS_CLAMP, then out-of-range image coordinates return the border color. The border color selected depends on the image channel order and can be one of the following values:

- ✚ If the image channel order is `image_channel_order::a`,
`image_channel_order::intensity`, `image_channel_order::rx`,
`image_channel_order::ra`, `image_channel_order::rgx`,
`image_channel_order::rgbx`, `image_channel_order::srgbx`,
`image_channel_order::argb`, `image_channel_order::bgra`,
`image_channel_order::abgr`, `image_channel_order::rgba`,
`image_channel_order::srgba` or
`image_channel_order::sbgra`, the border color is (0.0f, 0.0f,
0.0f, 0.0f).
- ✚ If the image channel order is `image_channel_order::r`,
`image_channel_order::rg`, `image_channel_order::rgb`, or
`image_channel_order::luminance`, the border color is (0.0f,
0.0f, 0.0f, 1.0f).
- ✚ If the image channel order is `image_channel_order::depth`, the border value is 0.0f.

2.12.2.2 sRGB Images

The built-in image read functions perform sRGB to linear RGB conversions if the image is an sRGB image. Writes to sRGB images from a kernel is an optional extension. The **cl_khr_srgb_image_writes** extension will be reported in the CL_DEVICE_EXTENSIONS string if a device supports writing to sRGB images using **image_write**. **clGetSupportedImageFormats** will return the supported sRGB images if CL_MEM_READ_WRITE or CL_MEM_WRITE_ONLY is specified in *flags* argument and the device supports writing to an sRGB image. If the **cl_khr_srgb_image_writes** extension is supported, the built-in image write functions perform the linear to sRGB conversion.

Only the R, G and B components are converted from linear to sRGB and vice-versa. The alpha component is returned as is.

The image and sampler types and the image functions are defined in header `<opencl_image>`.

2.12.3 Image Sample Functions

These functions are available if the image is declared with the `image_access:sample` qualifier.

T image1d<T>::sample (sampler *s*, float *coord*) const

T image1d<T>::sample (sampler *s*, float *coord*, float *lod*) const

T image1d_array<T>::sample (sampler *s*, float2 *coord*) const

T image1d_array<T>::sample (sampler *s*, float2 *coord*, float *lod*) const

T image2d<T>::sample (sampler *s*, float2 *coord*) const

T image2d<T>::sample (sampler *s*, float2 *coord*, float *lod*) const

T image2d_array<T>::sample (sampler *s*, float3 *coord*) const

T image2d_array<T>::sample (sampler *s*, float3 *coord*, float *lod*) const

T image3d<T>::sample (sampler *s*, float3 *coord*) const

T image3d<T>::sample (sampler *s*, float3 *coord*, float *lod*) const

T image2d_depth<T>::sample (sampler *s*, float2 *coord*) const

T image2d_depth<T>::sample (sampler *s*, float2 *coord*, float *lod*) const


T image2d_array_depth<T>::sample (sampler *s*, float3 *coord*) const

T image2d_array_depth<T>::sample (sampler *s*, float3 *coord*, float *lod*) const

Description:

Use the coordinate *coord* and level of detail *lod* to sample from the image object specified by *image*. The image must be declared with the `read` access qualifier.

The **image_sample** functions that take an image object where **T** is a floating-point type behaves as follows:

-  returns floating-point values in the range [0.0 ... 1.0] for image objects created with `image_channel_data_type` set to one of the pre-defined packed formats or `image_channel_type::unorm_int8` or `image_channel_type::unorm_int16`.

- ✚ returns floating-point values in the range [-1.0 ... 1.0] for image objects created with `image_channel_type::snorm_int8` or `image_channel_type::snorm_int16`.
- ✚ returns floating-point values for image objects created with `image_channel_type::float16` or `image_channel_type::float32`.

Values returned by **image_sample** where **T** is a floating-point type for image objects with *image_channel_data_type* values not specified in the description above are undefined.

The **image_sample** functions that take an image object where **T** is a signed integer type can only be used with image objects created with:

```
image_channel_type::sint8,
image_channel_type::sint16 and
image_channel_type::sint32.
```

If the *image_channel_data_type* is not one of the above values, the values returned by **image_sample** are undefined.

The **image_sample** functions that take an image object where **T** is an unsigned integer type can only be used with image objects created with:

```
image_channel_type::uint8,
image_channel_type::uint16 and
image_channel_type::uint32.
```

If the *image_channel_data_type* is not one of the above values, the values returned by **image_sample** are undefined.

2.12.4 Image Read Functions

These functions are available if the image is declared with the `image_access:sample`, `image_access:read` or `image_access:read_write` qualifier.

T image1d<T>::read (int *coord*, int *lod*=0) const

T image1d_array<T>::read (int2 *coord*, int *lod*=0) const

T image2d<T>::read (int2 *coord*, int *lod*=0) const

T image2d_array<T>::read (int3 *coord*, int *lod*=0) const

T image3d<T>::read (int3 *coord*, int *lod*=0) const

T image2d_depth<T>::read (int2 *coord*, int *lod*=0) const

T image2d_array_depth<T>::read (int3 *coord*, int *lod*=0) const

Description:

Use the coordinate *coord* to do an element lookup in the mip-level specified by *lod* of the image object specified by *image*. The **image_read** functions behave exactly as the corresponding **image_sample** functions described in *section 2.12.3* that take integer coordinates and a sampler with filter mode set to CLK_FILTER_NEAREST, normalized coordinates set to CLK_NORMALIZED_COORDS_FALSE and addressing mode to CLK_ADDRESS_NONE. The *coord* components are considered to be unnormalized coordinates and must be in the range 0 ... image width - 1, 0 ... image height - 1 and 0 ... image depth - 1 (for 3D image objects) or 0 ... image array size - 1 (for image array objects), where image width, height and depth are the width, height and depth of the mip-level specified by *lod*; otherwise the behavior of **image_read** is undefined.

The image must be declared with the `read` or `read_write` access qualifier.

The **image_read** functions that take an image object where **T** is a floating-point type behaves as follows:

- ✚ returns floating-point values in the range [0.0 ... 1.0] for image objects created with *image_channel_data_type* set to one of the pre-defined packed formats or `image_channel_type::unorm_int8` or `image_channel_type::unorm_int16`.
- ✚ returns floating-point values in the range [-1.0 ... 1.0] for image objects created with `image_channel_type::snorm_int8` or `image_channel_type::snorm_int16`.
- ✚ returns floating-point values for image objects created with `image_channel_type::float16` or `image_channel_type::float32`.

Values returned by **image_read** where **T** is a floating-point type for image objects with *image_channel_data_type* values not specified in the description above are undefined.

The **image_read** functions that take an image object where **T** is a signed integer type can only be used with image objects created with:

```
image_channel_type::sint8,
image_channel_type::sint16 and
image_channel_type::sint32.
```

If the *image_channel_data_type* is not one of the above values, the values returned by **image_read** are undefined.

The **image_read** functions that take an image object where **T** is an unsigned integer type can only be used with image objects created with *image_channel_data_type* set to one of the following values:

```
image_channel_type::uint8,
image_channel_type::uint16 and
image_channel_type::uint32.
```

If the *image_channel_data_type* is not one of the above values, the values returned by **image_read** are undefined.

2.12.5 Image Write Functions

These functions are available if the image is declared with the *image_access:write* or *image_access:read_write* qualifier.

```
void image1d<T>::write (int coord, T color, int lod=0)
```

```
void image1d_array<T>::write (int2 coord, T color, int lod=0)
```

```
void image2d<T>::write (int2 coord, T color, int lod=0)
```

```
void image2d_array<T>::write (int3 coord, T color, int lod=0)
```

```
void image3d<T>::write (int3 coord, T color, int lod=0)
```

```
void image2d_depth<T>::write (int2 coord, T color, int lod=0)
```

```
void image2d_array_depth<T>::write (int3 coord, T color, int lod=0)
```

Description:

Write *color* value to location specified by *coord* in the mip-level specified by *lod* of the image object specified by *image*. Appropriate data format conversion to the

specified image format is done before writing the color value. The *coord* components are considered to be unnormalized coordinates and must be in the range 0 ... image width - 1, 0 ... image height - 1 and 0 ... image depth - 1 (for 3D image objects) or 0 ... image array size - 1 (for image array objects), where image width, image height and image depth are the width, height and depth of the mip-level specified by *lod*; otherwise the behavior of **image_write** is undefined.

The image must be declared with the `write` or `read_write` access qualifier.

The **image_write** functions that take an image object where **T** is a floating-point type can only be used with image objects created with *image_channel_data_type* set to one of the pre-defined packed formats or set to

```
image_channel_type::snorm_int8,
image_channel_type::unorm_int8,
image_channel_type::snorm_int16,
image_channel_type::unorm_int16, image_channel_type::float16
or image_channel_type::float32. Appropriate data format conversion will
be done to convert channel data from a floating-point value to actual data format in
which the channels are stored.
```

The **image_write** functions that take an image object where **T** is a signed integer type can only be used with image objects created with:

```
image_channel_type::sint8,
image_channel_type::sint16 and
image_channel_type::sint32.
```

The **image_write** functions that take an image object where **T** is an unsigned integer type can only be used with image objects created with:

```
image_channel_type::uint8,
image_channel_type::uint16 and
image_channel_type::uint32.
```

The behavior of **image_write** for image objects created with *image_channel_data_type* values not specified in the description above is undefined.

2.12.6 Image Query Functions

```
int image1d<T>::width (int lod=0) const
```

```
int image1d_array<T>::width (int lod=0) const
```

```
int image2d<T>::width (int lod=0) const
```


int **image2d_array**<T>::**width** (int *lod*=0) const

int **image3d**<T>::**width** (int *lod*=0) const

int **image2d_depth**<T>::**width** (int *lod*=0) const

int **image2d_array_depth**<T>::**width** (int *lod*=0) const

Description: Returns the image width of the mip-level specified by *lod* in pixels.

int **image2d**<T>::**height** (int *lod*=0) const

int **image2d_array**<T>::**height** (int *lod*=0) const

int **image3d**<T>::**height** (int *lod*=0) const

int **image2d_depth**<T>::**height** (int *lod*=0) const

int **image2d_array_depth**<T>::**height** (int *lod*=0) const

Description: Returns the image height of the mip-level specified by *lod* in pixels.

int **image3d**<T>::**depth** (int *lod*=0) const

Description: Returns the image depth of the mip-level specified by *lod* in pixels.

int **image1d_array**<T>::**array_size** () const

int **image2d_array**<T>::**array_size** () const

int **image2d_array_depth**<T>::**array_size** () const

Description: Returns the image array size in pixels.

int **image1d**<T>::**miplevels** () const

int **image1d_array**<T>::**miplevels** () const

int **image2d**<T>::**miplevels** () const

int **image2d_array**<T>::**miplevels** () const

int **image3d**<T>::**miplevels** () const

int **image2d_depth**<T>::**miplevels** () const

`int image2d_array_depth<T>::miplevels () const`

Description: Returns the number of miplevels.

`image_channel_type image1d<T>::channel_data_type () const`

`image_channel_type image1d_array<T>::channel_data_type () const`

`image_channel_type image2d<T>::channel_data_type () const`

`image_channel_type image2d_array<T>::channel_data_type () const`

`image_channel_type image3d<T>::channel_data_type () const`

`image_channel_type image2d_depth<T>::channel_data_type () const`

`image_channel_type image2d_array_depth<T>::channel_data_type () const`

Description: Returns the image channel data type. Valid values are:

```
enum class image_channel_type {
    snorm_int8,
    snorm_int16,
    unorm_int8,
    unorm_int16,
    unorm_short565,
    unorm_short555,
    unorm_short101010,
    sint8,
    sint16,
    sint32,
    uint8,
    uint16,
    uint32,
    float16,
    float32,
};
```

`image_channel_order image1d<T>::channel_order () const`

`image_channel_order image1d_array<T>::channel_order () const`

`image_channel_order image2d<T>::channel_order () const`

`image_channel_order image2d_array<T>::channel_order () const`

image_channel_order **image3d**<T>::channel_order () const

image_channel_order **image2d_depth**<T>::channel_order () const

image_channel_order **image2d_array_depth**<T>::channel_order () const

Description: Returns the image channel data order. Valid values are:

```
enum class image_channel_order {
    a,
    r,
    rx,
    rg,
    rgx,
    ra,
    rgb,
    rgbx,
    rgba,
    argb,
    bgra,
    intensity,
    luminance,
    abgr,
    depth,
    srgb,
    srgbx,
    srgba,
    sbgra
};
```

2.12.7 Reading and writing to the same image in a kernel

To read and write to the same image in a kernel, the image must be declared with the `read_write` access qualifier. Only sampler-less reads and write functions can be called on an image declared with the `read_write` access qualifier. Calling the **image_sample** functions on an image declared with the `read_write` access qualifier will result in a compilation error.

The `atomic_work_item_fence(mem_type_image)` built-in function can be used to make sure that writes are visible to later sampler-less reads by the same work-item. Only a scope of `memory_order_acq_rel` is valid for `atomic_work_item_fence` when passed the `mem_type_image` flag. If multiple work-items are writing to and reading from multiple locations in an image, the `work_group_barrier(mem_type_image)` should be used.

Consider the following example:

```
#include <opencl_stdlib>
using namespace cl;

kernel void
foo(image2d<float4, image_access::read_write> img, ... )
{
    int2 coord;
    coord.x = (int)get_global_id(0);
    coord.y = (int)get_global_id(1);

    float4 clr = img.read(coord);
    ...
    img.write(coord, clr);

    // required to ensure that following read from image at
    // location coord returns the latest color value.
    atomic_work_item_fence(mem_type_image,
                           memory_order_acq_rel,
                           memory_scope_work_item);

    float4 clr_new = img.read(coord);
    ...
}
```

2.12.8 Mapping image channels to color values returned by `image_sample`, `image_read` and color values passed to `image_write` to image channels

The following table describes the mapping of the number of channels of an image element to the appropriate components in the float4, int4 or uint4 vector data type for the color values returned by **image_sample**, **image_read** or supplied to **image_write**. The unmapped components will be set to 0.0 for red, green and blue channels and will be set to 1.0 for the alpha channel.

Image Channel Order	float4, int4 or uint4 components of channel data
r, rx	(r, 0.0, 0.0, 1.0)
a	(0.0, 0.0, 0.0, a)
rg, rgx	(r, g, 0.0, 1.0)
ra	(r, 0.0, 0.0, a)
rgb, rgbx,	(r, g, b, 1.0)

Provisional Specification

srgb, srgbx	
rgba, bgra, argb, abgr, srgba, sbgra	(r, g, b, a)
intensity	(I, I, I, I)
luminance	(L, L, L, 1.0)

For `image_channel_order::depth` images, a scalar value is returned by **image_sample**, **image_read** or supplied to **image_write**.

NOTE: A kernel that uses a sampler with the `CL_ADDRESS_CLAMP` addressing mode with multiple images may result in additional samplers being used internally by an implementation. If the same sampler is used with multiple images called via **image_sample**, then it is possible that an implementation may need to allocate an additional sampler to handle the different border color values that may be needed depending on the image formats being used. The implementation allocated samplers will count against the maximum sampler values supported by the device and given by `CL_DEVICE_MAX_SAMPLERS`. Enqueuing a kernel that requires more samplers than the implementation can support will result in a `CL_OUT_OF_RESOURCES` error being returned.

2.13 Work-group Functions

The OpenCL C++ programming language implements the following built-in functions that operate on a work-group level. These built-in functions must be encountered by all work-items in a work-group executing the kernel. We use the generic type name `gentype` to indicate the built-in data types `half`, `int`, `uint`, `long`, `ulong`, `float` or `double`⁴³ as the type for the arguments.

Defined in header: `<opencl_workgroup>`

```
enum class work_group_op { add, min, max };
```

Function	Description
<code>bool work_group_all (bool <i>predicate</i>)</code>	Evaluates <i>predicate</i> for all work-items in the work-group and returns true if <i>predicate</i> evaluates to true for all work-items in the work-group.
<code>bool work_group_any (bool <i>predicate</i>)</code>	Evaluates <i>predicate</i> for all work-items in the work-group and returns true if <i>predicate</i> evaluates to true for any work-items in the work-group.
<code>gentype work_group_broadcast (</code> <code>gentype a,</code> <code>size_t local_id)</code> <code>gentype work_group_broadcast (</code> <code>gentype a,</code> <code>size_t local_id_x,</code> <code>size_t local_id_y)</code> <code>gentype work_group_broadcast (</code> <code>gentype a,</code> <code>size_t local_id_x,</code> <code>size_t local_id_y,</code> <code>size_t local_id_z)</code>	Broadcast the value of <i>x</i> for work-item identified by <i>local_id</i> to all work-items in the work-group. <i>local_id</i> must be the same value for all work-items in the work-group.
<code>gentype</code> <code>work_group_reduce<work_group_op::op> (</code> <code>gentype x)</code>	Return result of reduction operation specified by <op> for all values of <i>x</i> specified by work-items in a work-group.
<code>gentype</code> <code>work_group_scan_</code>	Do an exclusive scan operation specified by <op> of all values specified by work-

⁴³ Only if double precision is supported.

exclusive<work_group_op::op> (gentype <i>x</i>)	<p>items in the work-group. The scan results are returned for each work-item.</p> <p>The scan order is defined by increasing 1D linear global ID within the work-group.</p>
gentype work_group_scan_inclusive<work_group_op::op> (gentype <i>x</i>)	<p>Do an inclusive scan operation specified by <op> of all values specified by work-items in the work-group. The scan results are returned for each work-item.</p> <p>The scan order is defined by increasing 1D linear global ID within the work-group.</p>
bool sub_group_all (bool <i>predicate</i>)	Evaluates predicate for all work-items in the sub-group and returns a non-zero value if predicate evaluates to non-zero for all work-items in the sub-group.
bool sub_group_any (bool <i>predicate</i>)	Evaluates predicate for all work-items in the sub-group and returns a non-zero value if predicate evaluates to non-zero for any work-items in the sub-group.
gentype sub_group_broadcast (gentype <i>x</i> , size_t <i>sub_group_local_id</i>)	Broadcast the value of <i>x</i> for work-item identified by <i>sub_group_local_id</i> (value returned by get_sub_group_local_id) to all work-items in the sub-group. <i>sub_group_local_id</i> must be the same value for all work-items in the sub-group.
gentype sub_group_reduce<work_group_op::op> (gentype <i>x</i>)	Return result of reduction operation specified by <op> for all values of <i>x</i> specified by work-items in a sub-group.
gentype sub_group_scan_exclusive< work_group_op::op> (gentype <i>x</i>)	<p>Do an exclusive scan operation specified by <op> of all values specified by work-items in a sub-group. The scan results are returned for each work-item.</p> <p>The scan order is defined by increasing 1D linear global ID within the sub-group.</p>
gentype sub_group_scan_inclusive< work_group_op::op> (gentype <i>x</i>)	<p>Do an inclusive scan operation specified by <op> of all values specified by work-items in a sub-group. The scan results are returned for each work-item.</p> <p>The scan order is defined by increasing</p>

	1D linear global ID within the sub-group.
--	---

Table 2.13 *Built-in Work-group Functions*

The inclusive scan operation takes a binary operator *op* with an identity *I* and *n* (where *n* is the size of the work-group) elements [*a*₀, *a*₁, ... *a*_{*n*-1}] and returns [*a*₀, (*a*₀ *op* *a*₁), ... (*a*₀ *op* *a*₁ *op* ... *op* *a*_{*n*-1})]. If <*op*> = add, the identity *I* is 0. If <*op*> = min, the identity *I* is INT_MAX, UINT_MAX, LONG_MAX, ULONG_MAX, for int, uint, long, ulong types and is +INF for floating-point types. Similarly if <*op*> = max, the identity *I* is INT_MIN, 0, LONG_MIN, 0 and -INF.

Consider the following example:

```
void foo(int *p)
{
    ...
    int prefix_sum_val =
        work_group_scan_inclusive<work_group_op::add>(
            p[get_local_id(0)]);
}
```

For the example above, let's assume that the work-group size is 8 and *p* points to the following elements [3 1 7 0 4 1 6 3]. Work-item 0 calls **work_group_scan_inclusive_add** with 3 and returns 3. Work-item 1 calls **work_group_scan_inclusive_add** with 1 and returns 4. The full set of values returned by **work_group_scan_inclusive_add** for work-items 0 ... 7 are [3 4 11 11 15 16 22 25].

The exclusive scan operation takes a binary associative operator *op* with an identity *I* and *n* (where *n* is the size of the work-group) elements [*a*₀, *a*₁, ... *a*_{*n*-1}] and returns [*I*, *a*₀, (*a*₀ *op* *a*₁), ... (*a*₀ *op* *a*₁ *op* ... *op* *a*_{*n*-2})]. For the example above, the exclusive scan add operation on the ordered set [3 1 7 0 4 1 6 3] would return [0 3 4 11 11 15 16 22].

NOTE: The order of floating-point operations is not guaranteed for the **work_group_reduce<op>**, **work_group_scan_inclusive<op>** and **work_group_scan_exclusive<op>** built-in functions that operate on half, float and double data types. The order of these floating-point operations is also non-deterministic for a given work-group.

2.14 Pipe Functions

A pipe is specified as the templated type:

```
pipe<typename T, pipe_access::a=pipe_access::read>
```

`pipe_access` is an enumerated class described as:

```
enum class pipe_access { read, write };
```

The data type `T` specifies the size of each packet in the pipe. `T` may be any supported OpenCL C++ scalar and vector integer or floating-point data type, or a user-defined type built from these scalar and vector data types.

The `pipe_access` qualifier describes how the pipe is accessed. The supported `pipe_access` qualifier values are:

`read` – a kernel and its callees and enqueued child kernels read from a pipe

`write` – a kernel and its callees and enqueued child kernels can write to the pipe.

A pipe object declared in program scope is specified as the template type:

```
program_pipe<typename T>
```

Such variables declare an allocation of a pipe object and are sized by passing the storage capacity in packets to the pipe constructor:

```
program_pipe<T>::program_pipe(size_t max_packets,  
                               const pipe_properties *properties=null)
```

`max_packets` specifies the pipe capacity by specifying the maximum number of packets the pipe can hold. This must be a compile time constant.

`properties` specifies a list of properties for the pipe and their corresponding values. Each property name is immediately followed by the corresponding desired value. The list is terminated with 0. In OpenCL 2.1, `properties` must be null.

The `program_pipe` type does not support the pipe builtin functions. Instead a kernel may construct a read-only or write-only pipe with kernel storage duration from a `program_pipe`.

2.14.1.1 Restrictions

- ✚ Pipes can only be passed as arguments to a function (including kernel functions).
- ✚ The `pipe` type cannot be used with variables declared inside a kernel, a structure or union field, a pointer type, an array, global variables declared in program scope or the return type of a function.
- ✚ The `program_pipe` type cannot be used with variables declared inside a kernel, passed as an argument to a kernel or non-kernel function, a structure or union field, a pointer type, an array or the return type of a function.
- ✚ Variables of type `program_pipe` can only be declared in program scope.
- ✚ Variables of type `pipe` created from `program_pipe` can only be declared inside a kernel function at kernel scope.
- ✚ A kernel cannot read from and write to the same pipe object.
- ✚ A kernel cannot read from and write to the same program pipe object.

Examples:

```
pipe<int4> pipeA; // a pipe with int4 packets
pipe<user_type> pipeB; // a pipe with user_type
                      // packets
```

In the following example

```
kernel void
foo (pipe<float4> pipeA,
    pipe<int4, pipe_access::write> pipeB)
{
    ...
}
```

`pipeA` is a read-only pipe object, and `pipeB` is a write-only pipe object.

`null_reserve_id` is a predefined reservation ID of type `reserve_id` that refers to an invalid reservation ID.

The following examples describes program pipe objects:

Example 1:

Provisional Specification

```
program_pipe<int> myGlobalPipe0(100);

kernel void producer()
{
    pipe<int, pipe_access::read> p(myGlobalPipe0);
    p.read_pipe(...);
}

kernel void consumer()
{
    pipe<int, pipe_access::write> p(myGlobalPipe0);
    p.write_pipe(...);
}
```

Example 2:

```
program_pipe<float2> myGlobalPipe1(200);

kernel void pipeline()
{
    pipe<int, pipe_access::read> rp(myGlobalPipe1);
    pipe<int, pipe_access::write> wp(myGlobalPipe1);

    event evt;
    q.enqueue_kernel(flags, ndrange, 0, null, &evt,
        [=]{
            // produce
            wp.write_pipe(...);
        });

    q.enqueue_kernel(flags, ndrange, 1, &evt, null,
        [=]{
            // consume
            rp.write_pipe(...);
        });
}
```

The pipe types and member functions are defined in header `<opencl_pipe>`.

2.14.1.2 Built-in Pipe Read Functions

The following functions are available for a pipe declared with the `read` access qualifier.

Function	Description
int pipe<T>::read (T& <i>ptr</i>)	Read packet from pipe <i>p</i> into <i>ptr</i> . Returns 0 if read is successful and a negative value if the pipe is empty.
reserve_id pipe<T>::reserve_read (uint <i>num_packets</i>)	Reserve <i>num_packets</i> entries for reading from pipe <i>p</i> . Returns a valid reservation ID if the reservation is successful.
void pipe<T>::commit_read (reserve_id <i>reserve_id</i>)	Indicates that all reads to <i>num_packets</i> associated with reservation <i>reserve_id</i> are completed.

Table 2.14.1 *Built-in Pipe Read Functions*

2.14.1.3 Built-in Pipe Write Functions

The following functions are available for a pipe declared with the `write` access qualifer.

Function	Description
int pipe<T>::write (const T& <i>ptr</i>)	Write packet specified by <i>ptr</i> to pipe <i>p</i> . Returns 0 if write is successful and a negative value if the pipe is full.
int pipe<T>::write (reserve_id <i>reserve_id</i> , uint <i>index</i> , const gentype& <i>ptr</i>)	Write packet specified by <i>ptr</i> to the reserved area of the pipe referred to by <i>reserve_id</i> and <i>index</i> . The reserved pipe entries are referred to by indices that go from 0 ... <i>num_packets</i> – 1. Returns 0 if write is successful and a negative value otherwise.
reserve_id pipe<T>::reserve_write (uint <i>num_packets</i>)	Reserve <i>num_packets</i> entries for writing to pipe <i>p</i> . Returns a valid reservation ID if the reservation is successful.
void pipe<T>::commit_write (reserve_id <i>reserve_id</i>)	Indicates that all writes to <i>num_packets</i> associated with reservation <i>reserve_id</i> are completed.

Table 2.14.2 *Built-in Pipe Write Functions*

2.14.1.4 Built-in Work-group Pipe Read Functions

The following functions are available for a pipe declared with the `read` access qualifier. These functions must be encountered by all work-items in a work-group executing the kernel with the same argument values; otherwise the behavior is undefined.

Function	Description
<code>reserve_id</code> <code>pipe<T>::work_group_reserve_read (</code> <code>uint num_packets)</code>	Reserve <i>num_packets</i> entries for reading from to pipe <i>p</i> . Returns a valid reservation ID if the reservation is successful. The reserved pipe entries are referred to by indices that go from 0 ... <i>num_packets</i> – 1.
<code>void</code> <code>pipe<T>::work_group_commit_read (</code> <code>reserve_id reserve_id)</code>	Indicates that all reads to <i>num_packets</i> associated with reservation <i>reserve_id</i> are completed.

Table 2.14.3 *Built-in Pipe Work-group Read Functions*

2.14.1.5 Built-in Work-group Pipe Write Functions

The following member functions are available for a pipe declared with the `write` access qualifier. These member functions must be encountered by all work-items in a work-group executing the kernel with the same argument values; otherwise the behavior is undefined.

Function	Description
<code>reserve_id</code> <code>pipe<T>::work_group_reserve_write (</code> <code>uint num_packets)</code>	Reserve <i>num_packets</i> entries for writing to pipe <i>p</i> . Returns a valid reservation ID if the reservation is successful. The reserved pipe entries are referred to by indices that go from 0 ... <i>num_packets</i> – 1.
<code>void</code> <code>pipe<T>::work_group_commit_write (</code> <code>reserve_id reserve_id)</code>	Indicates that all writes to <i>num_packets</i> associated with reservation <i>reserve_id</i> are completed.

Table 2.14.4 *Built-in Pipe Work-group Write Functions*

2.14.1.6 Built-in Sub-group Pipe Read Functions

The following member functions are available for a pipe declared with the `read` access qualifier. These functions must be encountered with the same argument values by all work-items in a sub-group executing the kernel; otherwise the behavior is undefined.

Function	Description
<code>reserve_id</code> <code>pipe<T>::sub_group_reserve_read (</code> <code>uint num_packets)</code> 	Reserve <i>num_packets</i> entries for reading from <i>pipe</i> . Returns a valid reservation ID if the reservation is successful and a null reservation otherwise. The reserved pipe entries are referred to by indices that go from 0 ... <i>num_packets</i> – 1.
<code>void</code> <code>pipe<T>::sub_group_commit_read (</code> <code>reserve_id_t reserve_id)</code> 	Indicates that all reads associated with reservation <i>reserve_id</i> are completed.

Table 2.14.4 *Built-in Pipe Sub-group Read Functions*

2.14.1.1 Built-in Sub-group Pipe Write Functions

The following member functions are available for a pipe declared with the `write` access qualifier. These functions must be encountered with the same argument values by all work-items in a sub-group executing the kernel; otherwise the behavior is undefined.

Function	Description
<code>reserve_id</code> <code>pipe<T>::sub_group_reserve_write (</code> <code>uint num_packets)</code> 	Reserve <i>num_packets</i> entries for writing to <i>pipe</i> . Returns a valid reservation ID if the reservation is successful and a null reservation otherwise. The reserved pipe entries are referred to by indices that go from 0 ... <i>num_packets</i> – 1.
<code>void</code> <code>pipe<T>::sub_group_commit_write (</code> <code>reserve_id_t reserve_id)</code> 	Indicates that all writes associated with reservation <i>reserve_id</i> are completed.

Table 2.14.6 *Built-in Pipe Sub-group Write Functions*

NOTE: The **pipe<T>::read** and **pipe<T>::write** pipe functions that take a reservation ID as an argument can be used to read from or write to a packet index. These functions can be used to read from or write to a packet index one or multiple times. If a packet index that is reserved for writing is not written to using the **pipe<T>::write** function, the contents of that packet in the pipe are undefined. **pipe<T>::commit_read** and **pipe<T>::work_group_commit_read** remove the entries reserved for reading from the pipe. **pipe<T>::commit_write** and **pipe<T>::work_group_commit_write** ensures that the entries reserved for writing are all added in-order as one contiguous set of packets to the pipe.

There can only be `CL_DEVICE_PIPE_MAX_ACTIVE_RESERVATIONS` (refer to *table 4.3*) reservations active (i.e. reservation IDs that have been reserved but not committed) per work-item or work-group for a pipe in a kernel executing on a device.

Work-item based reservations made by a work-item are ordered in the pipe as they are ordered in the program. Reservations made by different work-items that belong to the same work-group can be ordered using the work-group barrier function. The order of work-item based reservations that belong to different work-groups is implementation defined.

Work-group based reservations made by a work-group are ordered in the pipe as they are ordered in the program. The order of work-group based reservations by different work-groups is implementation defined.

2.14.1.2 Built-in Pipe Query Functions

The following functions are available:

Function	Description
uint pipe<T>::num_packets (void)	Returns the number of available entries in the pipe. The number of available entries in a pipe is a dynamic value. The value returned should be considered immediately stale.
uint pipe<T>::max_packets (void)	Returns the maximum number of packets specified when <i>pipe</i> was created.

Table 2.14.5 *Built-in Pipe Query Functions*

2.14.1.3 Built-in Reservation ID Functions

The following functions are available for the `reserve_id` type.

Function	Description
bool reserve_id::is_valid (void)	Return true if <i>reserve_id</i> is a valid reservation ID and false otherwise.

Table 2.14.6 *Built-in Reservation ID Functions*

2.14.1.4 Restrictions

The following behavior is undefined:

- ✚ A kernel fails to call **pipe<T>::reserve** before calling **pipe<T>::read** or **pipe<T>::write** that take a reservation ID.
- ✚ A kernel calls **pipe<T>::read**, **pipe<T>::write**, **pipe<T>::commit_read** or **pipe<T>::commit_write** with an invalid reservation ID.
- ✚ A kernel calls **pipe<T>::read** or **pipe<T>::write** with an valid reservation ID but with an *index* that is not a value from 0 ... *num_packets* – 1 specified to the corresponding call to *reserve_pipe*.
- ✚ A kernel calls **pipe<T>::read** or **pipe<T>::write** with a reservation ID that has already been committed (i.e. a **pipe<T>::commit_read** or **pipe<T>::commit_write** with this reservation ID has already been called).
- ✚ A kernel fails to call **pipe<T>::commit_read** for any reservation ID obtained by a prior call to **pipe<T>::reserve_read**.
- ✚ A kernel fails to call **pipe<T>::commit_write** for any reservation ID obtained by a prior call to **pipe<T>::reserve_write**.
- ✚ The contents of the reserved data packets in the pipe are undefined if the kernel does not call **pipe<T>::write** for all entries that were reserved by the corresponding call to **pipe<T>::reserve**.
- ✚ Calls to **pipe<T>::read** that takes a reservation ID and **pipe<T>::commit_read** or **pipe<T>::write_pipe** that takes a reservation ID and **pipe<T>::commit_write** for a given reservation ID must be called by the same kernel that made the reservation using **pipe<T>::reserve_read** or **pipe<T>::reserve_write**. The reservation ID cannot be passed to another kernel including child kernels.

2.15 Enqueuing kernels

OpenCL 2.0 allows a kernel to independently enqueue to the same device, without host interaction. A kernel may enqueue a kernel function or code represented as a kernel lambda function, and control execution order with event dependencies including user events and markers.

The following table describes the list of built-in functions that can be used to enqueue a kernel(s).

The keyword `nullevent` refers to an invalid device event. The keyword `nullqueue` refers to an invalid device queue.

The queue and event types and functions described in this section are defined in header `<opencl_queue>`.

A kernel lambda function is described as:

```
[ capture-list ] ( params ) kernel { body }
```

2.15.1.1 Built-in Functions – Enqueuing a kernel

The following functions can be used to enqueue kernels to a device queue.

```
template<typename... Args>
int queue::enqueue_kernel (kernel_enqueue_flags flags,
                           const ndrange& ndrange,
                           kernel_function kernel_func,
                           Args ...args)
```

```
template<typename... Args>
int queue::enqueue_kernel (kernel_enqueue_flags flags,
                           const ndrange& ndrange,
                           uint num_events_in_wait_list,
                           const event* event_wait_list,
                           event* event_ret,
                           kernel_function kernel_func,
                           Args ...args)
```

```
template<typename... Args>
int queue::enqueue_kernel (kernel_enqueue_flags flags,
                           const ndrange& ndrange,
                           kernel_function kernel_func,
                           Args ...args)
```

```
template<typename... Args>
int queue::enqueue_kernel (kernel_enqueue_flags flags,
                           const ndrange& ndrange,
                           uint num_events_in_wait_list,
                           const event* event_wait_list,
                           event* event_ret,
                           kernel_function kernel_func,
                           Args ...args)
```

Description

The **enqueue_kernel** function allows a work-item to enqueue a function to a device queue. Work-items can enqueue multiple functions to a device queue(s).

The **enqueue_kernel** function returns CLK_SUCCESS if the block is enqueued successfully and returns CLK_ENQUEUE_FAILURE otherwise. If the `-g` compile option is specified in compiler options passed to **clCompileProgram** or **clBuildProgram** when compiling or building the parent program, the following errors may be returned instead of CLK_ENQUEUE_FAILURE to indicate why **enqueue_kernel** failed to enqueue the block:

- ✚ CLK_INVALID_QUEUE if *queue* is not a valid device queue.
- ✚ CLK_INVALID_NDRANGE if *ndrange* is not a valid ND-range descriptor or if the program was compiled with `-cl-uniform-work-group-size` and the *local_work_size* is specified in *ndrange* but the *global_work_size* specified in *ndrange* is not a multiple of the *local_work_size*.
- ✚ CLK_INVALID_EVENT_WAIT_LIST if *event_wait_list* is NULL and *num_events_in_wait_list* > 0, or if *event_wait_list* is not NULL and *num_events_in_wait_list* is 0, or if event objects in *event_wait_list* are not valid events.
- ✚ CLK_DEVICE_QUEUE_FULL if *queue* is full.
- ✚ CLK_INVALID_ARG_SIZE if size of local memory arguments is 0.
- ✚ CLK_EVENT_ALLOCATION_FAILURE if *event_ret* is not NULL and an event could not be allocated.
- ✚ CLK_OUT_OF_RESOURCES if there is a failure to queue the kernel in *queue* because of insufficient resources needed to execute the kernel.

If an event is returned, **enqueue_kernel** performs an implicit retain on the returned event.

Below are some examples of how to enqueue a block.

```
#include <opencl_stdlib>
using namespace cl;

kernel void
my_func_A(int *a, int *b, int *c)
{
    ...
}

kernel void
my_func_B(int *a, int *b, int *c)
{
    ndrange range;
    // build ndrange information
    ...
    queue q = get_default_queue();

    // example - enqueue a kernel as a block
    q.enqueue_kernel(ndrange,
                     my_func_A,
                     a, b, c);
    ...
}

kernel void
my_func_C(int *a, int *b, int *c)
{
    ndrange range;
    // build ndrange information
    ...

    // note that a, b and c are variables in scope of
    // the block
    auto x = [=] kernel {my_func_A(a, b, c);};
    queue q = get_default_queue();

    // enqueue the block variable
    q.enqueue_kernel(CLK_ENQUEUE_FLAGS_WAIT_KERNEL,
                     ndrange,
                     x);
    ...
}
```

NOTE: Lambdas passed to `enqueue_kernel` cannot use global variables or stack variables local to the enclosing lexical scope that are a pointer type in the local address space.

Example:

```
kernel void
foo(int *a, local_ptr<int> lptr, ...)
{
    queue q = get_default_queue();
    q.enqueue_kernel(CLK_ENQUEUE_FLAGS_WAIT_KERNEL,
                    ndrange,
                    [=] kernel
                    {
                        size_t id = get_global_id(0);
                        local_ptr<int> p = lptr; // undefined behavior
                    });
}
```

2.15.1.2 Arguments that are a pointer type to local address space

A block passed to `enqueue_kernel` can have arguments declared to be a pointer to local memory. The `enqueue_kernel` built-in function variants allow blocks to be enqueued with a variable number of arguments. Each argument must be declared to be a pointer of a data type to local memory. These `enqueue_kernel` built-in function variants also have a corresponding number of arguments each of type `uint` that follow the block argument. These arguments specify the size of each local memory pointer argument of the enqueued block.

Some examples follow:

```
#include <opencl_stdlib>
using namespace cl;

kernel void
my_func_A_local_arg1(int *a, local_ptr<int> lptr)
{
    ...
}

kernel void
my_func_A_local_arg2(int *a,
                    local_ptr<int> lptr1,
                    local_ptr<float4> lptr2)
{
    ...
}
```

Provisional Specification

```
    ...
}

kernel void
my_func_B(int *a, ...)
{
    ...
    ndrange range = ndrange(...);

    uint local_mem_size = compute_local_mem_size();

    queue q = get_default_queue();

    q.enqueue_kernel(CLK_ENQUEUE_FLAGS_WAIT_KERNEL,
                     ndrange,
                     [=](local int *p) kernel {
                         my_func_A_local_arg1(a, p, ...);},
                     local_mem_size);
}

kernel void
my_func_C(int *a, ...)
{
    ...
    ndrange range = ndrange(...);

    auto x = [] (local_ptr<int>, local_ptr<float4>) kernel
    {
        my_func_A_local_arg2(a, lptr1, lptr2);
    }

    // calculate local memory size for lptr
    // argument in local address space for my_blk_A
    uint local_mem_size = compute_local_mem_size();

    queue q = get_default_queue();
    q.enqueue_kernel<int, float4>(
        CLK_ENQUEUE_FLAGS_WAIT_KERNEL
        range,
        x,
        local_mem_size,
        local_mem_size*4);
}
```

2.15.1.3 A Complete Example

The example below shows how to implement an iterative algorithm where the host enqueues the first instance of the nd-range kernel (`dp_func_A`). The kernel `dp_func_A` will launch a kernel (`evaluate_dp_work_A`) that will determine if new nd-range work needs to be performed. If new nd-range work does need to be performed, then `evaluate_dp_work_A` will enqueue a new instance of `dp_func_A`. This process is repeated until all the work is completed.

```
#include <opencl_stdlib>
using namespace cl;

kernel void
dp_func_A(queue q, ...)
{
    ...

    // queue a single instance of evaluate_dp_work_A to
    // device queue q. queued kernel begins execution after
    // kernel dp_func_A finishes

    if (get_global_id(0) == 0)
    {
        q.enqueue_kernel(CLK_ENQUEUE_FLAGS_WAIT_KERNEL,
                        ndrange(1),
                        [=] kernel {
                            evaluate_dp_work_A(q, ...);});
    }
}

kernel void
evaluate_dp_work_A(queue_t q,...)
{
    // check if more work needs to be performed
    bool more_work = check_new_work(...);
    if (more_work)
    {
        size_t global_work_size = compute_global_size(...);

        auto x = [=] kernel {dp_func_A(q, ...)};

        // get local WG-size for kernel dp_func_A
        size_t local_work_size =
            get_kernel_work_group_size(x);

        // build nd-range descriptor
```

Provisional Specification

```
ndrange range = ndrange(global_work_size,
                        local_work_size);

// enqueue dp_func_A
q.enqueue_kernel(CLK_ENQUEUE_FLAGS_WAIT_KERNEL,
                range,
                x);
}
...
}
```

2.15.1.4 Determining when a child kernel begins execution

The `kernel_enqueue_flags`⁴⁴ argument to `enqueue_kernel` functions can be used to specify when the child kernel begins execution. Supported values are described in the table below:

kernel_enqueue_flags enum	Description
<code>CLK_ENQUEUE_FLAGS_NO_WAIT</code>	Indicates that the enqueued kernels do not need to wait for the parent kernel to finish execution before they begin execution.
<code>CLK_ENQUEUE_FLAGS_WAIT_KERNEL</code>	Indicates that all work-items of the parent kernel must finish executing and all immediate ⁴⁵ side effects committed before the enqueued child kernel may begin execution.
<code>CLK_ENQUEUE_FLAGS_WAIT_WORK_GROUP</code> ⁴⁶	Indicates that the enqueued kernels wait only for the workgroup that enqueued the kernels to finish before they begin execution.

Table 2.15.1 *Kernel Enqueue Flags*

NOTE: The `kernel_enqueue_flags` flags are useful when a kernel enqueued from the host and executing on a device enqueues kernels on the device. The kernel enqueued from the host may not have an event associated with it. The

⁴⁴ Implementations are not required to honor this flag. Implementations may not schedule kernel launch earlier than the point specified by this flag, however.

⁴⁵ Immediate meaning not side effects resulting from child kernels. The side effects would include stores to global memory and pipe reads and writes.

⁴⁶ This acts as a memory synchronization point between work-items in a work-group and child kernels enqueued by work-items in the work-group.

`kernel_enqueue_flags` flags allow the developer to indicate when the child kernels can begin execution.

2.15.1.5 Determining when a parent kernel has finished execution

A parent kernel's execution status is considered to be complete when it and all its child kernels have finished execution. The execution status of a parent kernel will be `CL_COMPLETE` if this kernel and all its child kernels finish execution successfully. The execution status of the kernel will be an error code (given by a negative integer value) if it or any of its child kernels encounter an error, or are abnormally terminated.

For example, assume that the host enqueues a kernel `k` for execution on a device. Kernel `k` when executing on the device enqueues kernels `A` and `B` to a device queue(s). The `enqueue_kernel` call to enqueue kernel `B` specifies the event associated with kernel `A` in the `event_wait_list` argument i.e. wait for kernel `A` to finish execution before kernel `B` can begin execution. Let's assume kernel `A` enqueues kernels `X`, `Y` and `Z`. kernel `A` is considered to have finished execution i.e. its execution status is `CL_COMPLETE` only after `A` and the kernels `A` enqueued (and any kernels these enqueued kernels enqueue and so on) have finished execution.

2.15.1.6 Built-in Functions – Kernel Query Functions

Built-in Function	Description
<code>uint get_kernel_work_group_size (</code> <code>kernel_function kernel_func);</code>	This provides a mechanism to query the maximum work-group size that can be used to execute a function on a specific device given by <i>device</i> . <i>kernel_func</i> specifies the lambda to be enqueued.
<code>uint get_kernel_preferred_</code> <code>work_group_size_multiple (</code> <code>kernel_function kernel_func);</code>	Returns the preferred multiple of work-group size for launch. This is a performance hint. Specifying a work-group size that is not a multiple of the value returned by this query as the value of the local work size argument to <code>enqueue_kernel</code> will not fail to enqueue the lambda for execution unless the work-group size specified is larger than the device maximum.
<code>uint get_kernel_sub_group_</code> <code>count_for_ndrange (</code>	Returns the number of subgroups in each workgroup of the dispatch (except for the

const ndrange& ndrange, kernel_function <i>kernel_func</i>);	last in cases where the global size does not divide cleanly into work-groups) given the combination of the passed ndrange and the function to be enqueued.
uint get_kernel_max_sub_group_size_for_ndrange (const ndrange& ndrange, kernel_function <i>kernel_func</i>);	Returns the maximum sub-group size for a function that may be enqueued as a child kernel and the ndrange that will be used to enqueue it.

Table 2.15.2 *Built-in Kernel Query Functions*

2.15.1.7 Built-in Functions – Queuing other commands





The following functions can be used to enqueue a marker to a device queue.

```
int queue::enqueue_marker (uint num_events_in_wait_list,
                           const event* event_wait_list,
                           event& event_ret)
```

Description

The **enqueue_marker** function queues a marker command to *queue*. The marker command waits for a list of events specified by *event_wait_list* to complete before the marker completes.

enqueue_marker performs an implicit retain on the returned event. The **enqueue_marker** built-in function returns CLK_SUCCESS if the marked command is enqueued successfully and returns CLK_ENQUEUE_FAILURE otherwise. If the `-g` compile option is specified in compiler options passed to **clCompileProgram** or **clBuildProgram**, the following errors may be returned instead of CLK_ENQUEUE_FAILURE to indicate why **enqueue_marker** failed to enqueue the marker command:

-  CLK_INVALID_QUEUE if *queue* is not a valid device queue.
-  CLK_INVALID_EVENT_WAIT_LIST if *event_wait_list* is NULL, or if *event_wait_list* is not NULL and *num_events_in_wait_list* is 0, or if event objects in *event_wait_list* are not valid events.
-  CLK_DEVICE_QUEUE_FULL if *queue* is full.
-  CLK_EVENT_ALLOCATION_FAILURE if an event could not be allocated.

- 🚦 **CLK_OUT_OF_RESOURCES** if there is a failure to queue the marker in *queue* because of insufficient resources needed to execute the marker.

2.15.1.8 Built-in Functions – Event Functions

The following functions can be used with events.

int event::release(void)

Decrements the event reference count. The event object is deleted once the event reference count is zero, the specific command identified by this event has completed (or terminated) and there are no commands in any device command queue that require a wait for this event to complete.

event must be an event returned by `enqueue_kernel`, `enqueue_marker` or a user event.

int event::retain(void)

Increments the event reference count. *event* must be an event returned by `enqueue_kernel` or `enqueue_marker` or a user event.

event create_user_event(void)

Create a user event. Returns the user event. The execution status of the user event created is set to `CL_SUBMITTED`.

bool is_valid_event(event e);

Returns true if *event* is a valid event. Otherwise returns false.

void event::set_status()

Sets the execution status of a user event. *event* must be a user-event. *status* can be either `CL_COMPLETE` or a negative integer value indicating an error.

void event::capture_profiling_info (profiling_info name, void *value)

Captures the profiling information for functions that are enqueued as commands. The specific function being referred to is: `enqueue_kernel`. These enqueued commands are identified by unique event objects. The profiling information will be available in *value* once the command identified by *event* has completed.

event must be an event returned by `enqueue_kernel`.

name identifies which profiling information is to be queried and can be:

CLK_PROFILING_COMMAND_EXEC_TIME

value is a pointer to two 64-bit values.

The first 64-bit value describes the elapsed time CL_PROFILING_COMMAND_END – CL_PROFILING_COMMAND_START for the command identified by *event* in nanoseconds.

The second 64-bit value describes the elapsed time CL_PROFILING_COMMAND_COMPLETE – CL_PROFILING_COMMAND_START for the command identified by *event* in nanoseconds.

NOTE: The behavior of **event_capture_profiling_info** and **event::capture_profiling_info** when called multiple times for the same *event* is undefined.

Events can be used to identify commands enqueued to a command-queue from the host. These events created by the OpenCL runtime can only be used on the host i.e. as events passed in *event_wait_list* argument to various clEnqueue APIs or runtime APIs that take events as arguments such as clRetainEvent, clReleaseEvent, clGetEventProfilingInfo.

Similarly, events can be used to identify commands enqueued to a device queue (from a kernel). These event objects cannot be passed to the host or used by OpenCL runtime APIs such as the clEnqueue APIs or runtime APIs that take event arguments.

clRetainEvent and clReleaseEvent will return CL_INVALID_OPERATION if *event* specified is an event that refers to any kernel enqueued to a device queue using enqueue_kernel or enqueue_marker or is a user event created by create_user_event.

Similarly, clSetUserEventStatus can only be used to set the execution status of events created using clCreateUserEvent. User events created on the device can be set using set_user_event_status built-in function.

The example below shows how events can be used with kernels enqueued to multiple device queues.

```
#include <opencl_stdlib>
using namespace cl;
```

Provisional Specification

```
extern void barA_kernel(...);
extern void barB_kernel(...);

kernel void
foo(queue_t q0, queue q1, ...)
{
    ...
    clk_event_t evt0;

    // enqueue kernel to queue q0
    q0.enqueue_kernel(CLK_ENQUEUE_FLAGS_NO_WAIT,
                      range_A,
                      0, NULL, &evt0,
                      [=] kernel {barA_kernel(...);} );

    // enqueue kernel to queue q1
    q1.enqueue_kernel(CLK_ENQUEUE_FLAGS_NO_WAIT,
                      range_B,
                      1, &evt0, NULL,
                      [=] kernel {barB_kernel(...);} );

    // release event evt0. This will get released
    // after barA_kernel enqueued in queue q0 has finished
    // execution and barB_kernel enqueued in queue q1 and
    // waits for evt0 is submitted for execution i.e. wait
    // for evt0 is satisfied.
    evt0.release();
}
```

The example below shows how the marker command can be used with kernels enqueued to a device queue.

```
kernel void
foo(queue_t q, ...)
{
    ...
    event marker_event;
    event events[2];

    q.enqueue_kernel(CLK_ENQUEUE_FLAGS_NO_WAIT,
                     range,
                     0, NULL, &events[0],
                     [=] kernel {barA_kernel(...);} );

    q.enqueue_kernel(CLK_ENQUEUE_FLAGS_NO_WAIT,
                     range,
                     0, NULL, &events[1],
```

Provisional Specification

```
[=] kernel {barB_kernel(...);} );

// barA_kernel and barB_kernel can be executed
// out of order. we need to wait for both these
// kernels to finish execution before barC_kernel
// starts execution so we enqueue a marker command and
// then enqueue barC_kernel that waits on the event
// associated with the marker.

q.enqueue_marker(2, events, &marker_event);

q.enqueue_kernel(CLK_ENQUEUE_FLAGS_NO_WAIT,
                 1, &marker_event, NULL,
                 [=] kernel {barC_kernel(...);} );

events[0].release();
events[1].release();
marker_event.release();
}
```

2.15.1.9 Built-in Functions – Helper Functions

The following helper functions are available.

queue **get_default_queue**(void)

Returns the default device queue. If a default device queue has not been created, `nullqueue` is returned.

The following constructors are available to build a 1D, 2D or 3D ndrange descriptor:

```
ndrange::ndrange(size_t global_work_size)
ndrange::ndrange(size_t global_work_size, size_t local_work_size)
ndrange::ndrange(size_t global_work_offset,
                  size_t global_work_size, size_t local_work_size)

ndrange::ndrange(size_t global_work_size[2])
ndrange::ndrange(size_t global_work_size[2], size_t local_work_size[2])
ndrange::ndrange(size_t global_work_offset[2],
                  size_t global_work_size[2], size_t local_work_size[2])

ndrange::ndrange(size_t global_work_size[3])
ndrange::ndrange(size_t global_work_size[3], size_t local_work_size[3])
ndrange::ndrange(size_t global_work_offset[3],
                  size_t global_work_size[3], size_t local_work_size[3])
```

2.16 Metaprogramming and type traits

The feature set defined by C++14 section 20.11 is also supported in OpenCL C++. This section describes components used by OpenCL C++ programs, particularly in templates, to support the widest possible range of types, optimise template code usage, detect type related user errors, and perform type inference and transformation at compile time. It includes type classification traits, type property inspection traits, and type transformations. The type classification traits describe a complete taxonomy of all possible OpenCL C++ types, and state where in that taxonomy a given type belongs. The type property inspection traits allow important characteristics of types or of combinations of types to be inspected. The type transformations allow certain properties of types to be manipulated.

Header: `<opencl_type_traits>`

2.17 Diagnostics

The header `<opencl_assert>` provides a macro for documenting OpenCL C++ kernel assertions and a mechanism for disabling the assertion checks.

```
#ifndef NDEBUG
#define assert(condition) ((void)0)
#else
#define assert(condition) /*implementation defined*/
#endif
```

The definition of the macro `assert` depends on another macro, `NDEBUG`, which is not defined by the standard library but by the implementation. If `NDEBUG` is defined as a macro name at the point in the source code where `<opencl_assert>` is included, then `assert` does nothing.

If `NDEBUG` is not defined, then `assert` checks if its argument (which must have scalar type) compares equal to zero. If it does, `assert` outputs implementation-specific diagnostic information on the standard error output and aborts the kernel. The diagnostic information is required to include the text of expression, as well as the values of the standard macros `__FILE__`, `__LINE__`, and the standard variable `__func__`.





3 OpenCL Numerical Compliance

This section describes features of the C++ 14 and IEEE 754 standards that must be supported by all OpenCL compliant devices.

This section describes the functionality that must be supported by all OpenCL devices for single precision floating-point numbers. Currently, only single precision and half precision floating-point is a requirement. Double precision floating-point is an optional feature.

3.1 Rounding Modes

Floating-point calculations may be carried out internally with extra precision and then rounded to fit into the destination type. IEEE 754 defines four possible rounding modes:

-  Round to nearest even
-  Round toward $+\infty$
-  Round toward $-\infty$
-  Round toward zero

Round to nearest even is currently the only rounding mode required⁴⁷ by the OpenCL specification for single precision and double precision operations and is therefore the default rounding mode. In addition, only static selection of rounding mode is supported. Dynamically reconfiguring the rounding modes as specified by the IEEE 754 spec is unsupported.

3.2 INF, NaN and Denormalized Numbers

INF and NaNs must be supported. Support for signaling NaNs is not required.

Support for denormalized numbers with single precision floating-point is optional. Denormalized single precision floating-point numbers passed as input or produced as the output of single precision floating-point operations such as add, sub, mul, divide, and the functions defined in *sections 2.2* (math functions), *2.4* (common functions) and *2.5* (geometric functions) may be flushed to zero.

⁴⁷ Except for the embedded profile whether either round to zero or round to nearest rounding mode may be supported for single precision floating-point.

3.3 Floating-Point Exceptions

Floating-point exceptions are disabled in OpenCL. The result of a floating-point exception must match the IEEE 754 spec for the exceptions not enabled case. Whether and when the implementation sets floating-point flags or raises floating-point exceptions is implementation-defined. This standard provides no method for querying, clearing or setting floating-point flags or trapping raised exceptions. Due to non-performance, non-portability of trap mechanisms and the impracticality of servicing precise exceptions in a vector context (especially on heterogeneous hardware), such features are discouraged.

Implementations that nevertheless support such operations through an extension to the standard shall initialize with all exception flags cleared and the exception masks set so that exceptions raised by arithmetic operations do not trigger a trap to be taken. If the underlying work is reused by the implementation, the implementation is however not responsible for reclearing the flags or resetting exception masks to default values before entering the kernel. That is to say that kernels that do not inspect flags or enable traps are licensed to expect that their arithmetic will not trigger a trap. Those kernels that do examine flags or enable traps are responsible for clearing flag state and disabling all traps before returning control to the implementation. Whether or when the underlying work-item (and accompanying global floating-point state if any) is reused is implementation-defined.

The expressions **math_errorhandling** and **MATH_ERREXCEPT** are reserved for use by this standard, but not defined. Implementations that extend this specification with support for floating-point exceptions shall define **math_errorhandling** and **MATH_ERREXCEPT** per ISO / IEC 9899 : TC2.

3.4 Relative Error as ULPs

In this section we discuss the maximum relative error defined as `ulp` (units in the last place). Addition, subtraction, multiplication, fused multiply-add and conversion between integer and a single precision floating-point format are IEEE 754 compliant and are therefore correctly rounded. Conversion between floating-point formats and explicit conversions specified in *section 1.2.3* must be correctly rounded.

The ULP is defined as follows:

If x is a real number that lies between two finite consecutive floating-point numbers a and b , without being equal to one of them, then $\text{ulp}(x) = |b - a|$, otherwise $\text{ulp}(x)$ is the distance between the two non-

Provisional Specification

equal finite floating-point numbers nearest x .
Moreover, $\text{ulp}(\text{NaN})$ is NaN.

Attribution: This definition was taken with consent from Jean-Michel Muller with slight clarification for behavior at zero. Refer to <ftp://ftp.inria.fr/INRIA/publication/publi-pdf/RR/RR-5504.pdf>.

Table 3.4.1⁴⁸ describes the minimum accuracy of single precision floating-point arithmetic operations given as ULP values. The reference value used to compute the ULP value of an arithmetic operation is the infinitely precise result.

Function	Min Accuracy - ULP values ⁴⁹
$x + y$	Correctly rounded
$x - y$	Correctly rounded
$x * y$	Correctly rounded
$1.0 / x$	≤ 2.5 ulp
x / y	≤ 2.5 ulp
acos	≤ 4 ulp
acospi	≤ 5 ulp
asin	≤ 4 ulp
asinpi	≤ 5 ulp
atan	≤ 5 ulp
atan2	≤ 6 ulp
atanpi	≤ 5 ulp
atan2pi	≤ 6 ulp
acosh	≤ 4 ulp
asinh	≤ 4 ulp
atanh	≤ 5 ulp
cbrt	≤ 2 ulp
ceil	Correctly rounded
copysign	0 ulp
cos	≤ 4 ulp
cosh	≤ 4 ulp
cospi	≤ 4 ulp
erfc	≤ 16 ulp
erf	≤ 16 ulp
exp	≤ 3 ulp
exp2	≤ 3 ulp
exp10	≤ 3 ulp

⁴⁸ The ULP values for built-in math functions **lgamma** and **lgamma_r** is currently undefined.

⁴⁹ 0 ulp is used for math functions that do not require rounding.

Provisional Specification

expm1	<= 3 ulp
fabs	0 ulp
fdim	Correctly rounded
floor	Correctly rounded
fma	Correctly rounded
fmax	0 ulp
fmin	0 ulp
fmod	0 ulp
fract	Correctly rounded
frexp	0 ulp
hypot	<= 4 ulp
ilogb	0 ulp
ldexp	Correctly rounded
log	<= 3 ulp
log2	<= 3 ulp
log10	<= 3 ulp
log1p	<= 2 ulp
logb	0 ulp
mad	Implemented either as a correctly rounded fma or as a multiply followed by an add both of which are correctly rounded
maxmag	0 ulp
minmag	0 ulp
modf	0 ulp
nan	0 ulp
nextafter	0 ulp
pow(x, y)	<= 16 ulp
pown(x, y)	<= 16 ulp
powr(x, y)	<= 16 ulp
remainder	0 ulp
remquo	0 ulp
rint	Correctly rounded
rootn	<= 16 ulp
round	Correctly rounded
rsqrt	<= 2 ulp
sin	<= 4 ulp
sincos	<= 4 ulp for sine and cosine values
sinh	<= 4 ulp
sinpi	<= 4 ulp
sqrt	<= 3 ulp
tan	<= 5 ulp
tanh	<= 5 ulp
tanpi	<= 6 ulp

tgamma	≤ 16 ulp
trunc	Correctly rounded
native_cos	Implementation-defined
native_divide	Implementation-defined
native_exp	Implementation-defined
native_exp2	Implementation-defined
native_exp10	Implementation-defined
native_log	Implementation-defined
native_log2	Implementation-defined
native_log10	Implementation-defined
native_powr	Implementation-defined
native_recip	Implementation-defined
native_rsqr	Implementation-defined
native_sin	Implementation-defined
native_sqrt	Implementation-defined
native_tan	Implementation-defined

Table 3.4.1 *ULP values for single precision built-in math functions*

Table 3.4.2 describes the minimum accuracy of commonly used single precision floating-point arithmetic operations given as ULP values if the `-cl-fast-relaxed-math` compiler option is specified when compiling or building an OpenCL program. The minimum accuracy of math functions not defined in table 3.4.2 when the `-cl-fast-relaxed-math` compiler option is specified is as defined in table 3.4.1. The reference value used to compute the ULP value of an arithmetic operation is the infinitely precise result.

Function	Min Accuracy - ULP values ⁵⁰
1.0 / x	≤ 2.5 ulp for x in the domain of 2^{-126} to 2^{126}
x / y	≤ 2.5 ulp for x in the domain of 2^{-62} to 2^{62} and y in the domain of 2^{-62} to 2^{62} .
acos(x)	≤ 4096 ulp
acospi(x)	Implemented as acos(x) * M_PI_F. For non-derived implementations, the error is ≤ 8192 ulp.
asin(x)	≤ 4096 ulp
asinpi(x)	Implemented as asin(x) * M_PI_F. For non-derived implementations, the error is ≤ 8192 ulp.
atan(x)	≤ 4096 ulp

⁵⁰ 0 ulp is used for math functions that do not require rounding.

Provisional Specification

atan2(y, x)	Implemented as atan(y/x) for $x > 0$, atan(y/x) + M_PI_F for $x < 0$ and $y > 0$ and atan(y/x) - M_PI_F for $x < 0$ and $y < 0$.
atanpi(x)	Implemented as atan(x) * M_PI_F . For non-derived implementations, the error is ≤ 8192 ulp.
atan2pi(y, x)	Implemented as atan2(y, x) * M_PI_F . For non-derived implementations, the error is ≤ 8192 ulp.
acosh(x)	Implemented as log(x + sqrt(x*x - 1)) .
asinh(x)	Implemented as log(x + sqrt(x*x + 1)) .
cbrt(x)	Implemented as rootn(x, 3) . For non-derived implementations, the error is ≤ 8192 ulp.
cos(x)	For x in the domain $[-\pi, \pi]$, the maximum absolute error is $\leq 2^{-11}$ and larger otherwise.
cosh(x)	Implemented as 0.5 * exp(x) + exp(-x) . For non-derived implementations, the error is ≤ 8192 ulp.
cospi(x)	For x in the domain $[-1, 1]$, the maximum absolute error is $\leq 2^{-11}$ and larger otherwise.
exp(x)	$\leq 3 + \text{floor}(\text{fabs}(2 * x))$ ulp
exp2(x)	$\leq 3 + \text{floor}(\text{fabs}(2 * x))$ ulp
exp10(x)	Derived implementations implement this as exp2(x * log2(10)) . For non-derived implementations, the error is ≤ 8192 ulp.
expm1(x)	Derived implementations implement this as exp(x) - 1 . For non-derived implementations, the error is ≤ 8192 ulp.
log(x)	For x in the domain $[0.5, 2]$ the maximum absolute error is $\leq 2^{-21}$; otherwise the maximum error is ≤ 3 ulp for the full profile and ≤ 4 ulp for the embedded profile
log2(x)	For x in the domain $[0.5, 2]$ the maximum absolute error is $\leq 2^{-21}$; otherwise the maximum error is ≤ 3 ulp for the full profile and ≤ 4 ulp for the embedded profile
log10(x)	For x in the domain $[0.5, 2]$ the maximum absolute error is $\leq 2^{-21}$; otherwise the maximum error is ≤ 3 ulp for the full profile and ≤ 4 ulp for the embedded profile
log1p(x)	Derived implementations implement this as log(x + 1) . For non-derived implementations, the error is ≤ 8192 ulp.
pow(x, y)	Undefined for $x = 0$ and $y = 0$ or for $x < 0$ and non-integer y . For $x \geq 0$ or $x < 0$ and even y ,

	derived implementations implement this as $\exp2(y * \log2(x))$. For $x < 0$ and odd y , derived implementations implement this as $-\exp2(y * \log2(\text{fabs}(x)))$. For non-derived implementations, the error is ≤ 8192 ulp.
$\text{pown}(x, y)$	Defined only for integer values of y . Undefined for $x = 0$ and $y = 0$. For $x \geq 0$ or $x < 0$ and even y , derived implementations implement this as $\exp2(y * \log2(x))$. For $x < 0$ and odd y , derived implementations implement this as $-\exp2(y * \log2(\text{fabs}(x)))$. For non-derived implementations, the error is ≤ 8192 ulp.
$\text{powr}(x, y)$	Defined only for $x \geq 0$. Undefined for $x = 0$ and $y = 0$. Derived implementations implement this as $\exp2(y * \log2(x))$. For non-derived implementations, the error is ≤ 8192 ulp.
$\text{rootn}(x, y)$	Defined for $x > 0$ and for $x < 0$ when y is odd. Undefined for $x = 0$ and $y = 0$. Derived implementations implement this as $\exp2(\log2(x) / y)$ for $x > 0$. Derived implementations implement this as $-\exp2(\log2(-x) / y)$ for $x < 0$. For non-derived implementations, the error is ≤ 8192 ulp.
$\sin(x)$	For x in the domain $[-\pi, \pi]$, the maximum absolute error is $\leq 2^{-11}$ and larger otherwise.
$\text{sincos}(x)$	ulp values as defined for $\sin(x)$ and $\cos(x)$
$\sinh(x)$	Implemented as $0.5 * \exp(x) - \exp(-x)$. For non-derived implementations, the error is ≤ 8192 ulp.
$\sinpi(x)$	For x in the domain $[-1, 1]$, the maximum absolute error is $\leq 2^{-11}$ and larger otherwise.
$\tan(x)$	Derived implementations implement this as $\sin(x) * (1.0f / \cos(x))$. For non-derived implementations, the error is ≤ 8192 ulp.
$\tanpi(x)$	Derived implementations implement this as $\tan(x * \text{M_PI_F})$. For non-derived implementations, the error is ≤ 8192 ulp for x in the domain $[-1, 1]$.
$x * y + z$	Implemented either as a correctly rounded fma or as a multiply and an add both of which are correctly rounded.

Table 3.4.2 *ULP values for single precision built-in math functions with fast relaxed math*

Table 3.4.3 describes the minimum accuracy of double precision floating-point arithmetic operations given as ULP values. The reference value used to compute the ULP value of an arithmetic operation is the infinitely precise result.

Function	Min Accuracy - ULP values ⁵¹
$x + y$	Correctly rounded
$x - y$	Correctly rounded
$x * y$	Correctly rounded
$1.0 / x$	Correctly rounded
x / y	Correctly rounded
acos	≤ 4 ulp
acospi	≤ 5 ulp
asin	≤ 4 ulp
asinpi	≤ 5 ulp
atan	≤ 5 ulp
atan2	≤ 6 ulp
atanpi	≤ 5 ulp
atan2pi	≤ 6 ulp
acosh	≤ 4 ulp
asinh	≤ 4 ulp
atanh	≤ 5 ulp
cbrt	≤ 2 ulp
ceil	Correctly rounded
copysign	0 ulp
cos	≤ 4 ulp
cosh	≤ 4 ulp
cospi	≤ 4 ulp
erfc	≤ 16 ulp
erf	≤ 16 ulp
exp	≤ 3 ulp
exp2	≤ 3 ulp
exp10	≤ 3 ulp
expm1	≤ 3 ulp
fabs	0 ulp
fdim	Correctly rounded
floor	Correctly rounded
fma	Correctly rounded
fmax	0 ulp
fmin	0 ulp
fmod	0 ulp
fract	Correctly rounded

⁵¹ 0 ulp is used for math functions that do not require rounding.

frexp	0 ulp
hypot	≤ 4 ulp
ilogb	0 ulp
ldexp	Correctly rounded
log	≤ 3 ulp
log2	≤ 3 ulp
log10	≤ 3 ulp
log1p	≤ 2 ulp
logb	0 ulp
mad	Any value allowed (infinite ulp)
maxmag	0 ulp
minmag	0 ulp
modf	0 ulp
nan	0 ulp
nextafter	0 ulp
pow(x, y)	≤ 16 ulp
pown(x, y)	≤ 16 ulp
powr(x, y)	≤ 16 ulp
remainder	0 ulp
remquo	0 ulp
rint	Correctly rounded
rootn	≤ 16 ulp
round	Correctly rounded
rsqrt	≤ 2 ulp
sin	≤ 4 ulp
sincos	≤ 4 ulp for sine and cosine values
sinh	≤ 4 ulp
sinpi	≤ 4 ulp
sqrt	Correctly rounded
tan	≤ 5 ulp
tanh	≤ 5 ulp
tanpi	≤ 6 ulp
tgamma	≤ 16 ulp
trunc	Correctly rounded

Table 3.4.3 *ULP values for double precision built-in math functions*

Table 3.4.4 describes the minimum accuracy of half precision floating-point arithmetic operations given as ULP values. The reference value used to compute the ULP value of an arithmetic operation is the infinitely precise result.

Function	Min Accuracy - ULP values ⁵²
----------	---

⁵² 0 ulp is used for math functions that do not require rounding.

Provisional Specification

$x + y$	Correctly rounded
$x - y$	Correctly rounded
$x * y$	Correctly rounded
$1.0 / x$	Correctly rounded
x / y	Correctly rounded
acos	≤ 2 ulp
acospi	≤ 2 ulp
asin	≤ 2 ulp
asinpi	≤ 2 ulp
atan	≤ 2 ulp
atan2	≤ 2 ulp
atanpi	≤ 2 ulp
atan2pi	≤ 2 ulp
acosh	≤ 2 ulp
asinh	≤ 2 ulp
atanh	≤ 2 ulp
cbrt	≤ 2 ulp
ceil	Correctly rounded
copysign	0 ulp
cos	≤ 2 ulp
cosh	≤ 2 ulp
cospi	≤ 2 ulp
erfc	≤ 4 ulp
erf	≤ 4 ulp
exp	≤ 2 ulp
exp2	≤ 2 ulp
exp10	≤ 2 ulp
expm1	≤ 2 ulp
fabs	0 ulp
fdim	Correctly rounded
floor	Correctly rounded
fma	Correctly rounded
fmax	0 ulp
fmin	0 ulp
fmod	0 ulp
fract	Correctly rounded
frexp	0 ulp
hypot	≤ 2 ulp
ilogb	0 ulp
ldexp	Correctly rounded
log	≤ 2 ulp
log2	≤ 2 ulp
log10	≤ 2 ulp

log1p	≤ 2 ulp
logb	0 ulp
mad	Any value allowed (infinite ulp)
maxmag	0 ulp
minmag	0 ulp
modf	0 ulp
nan	0 ulp
nextafter	0 ulp
pow(x, y)	≤ 4 ulp
pown(x, y)	≤ 4 ulp
powr(x, y)	≤ 4 ulp
remainder	0 ulp
remquo	0 ulp
rint	Correctly rounded
rootn	≤ 4 ulp
round	Correctly rounded
rsqrt	≤ 1 ulp
sin	≤ 2 ulp
sincos	≤ 2 ulp for sine and cosine values
sinh	≤ 2 ulp
sinpi	≤ 2 ulp
sqrt	Correctly rounded
tan	≤ 2 ulp
tanh	≤ 2 ulp
tanpi	≤ 2 ulp
tgamma	≤ 4 ulp
trunc	Correctly rounded

Table 3.4.4 *ULP values for half precision built-in math functions*

3.5 Edge Case Behavior

The edge case behavior of the math functions (*section 1.2*) shall conform to sections F.9 and G.6 of ISO/IEC 9899:TC 2, except where noted below in *section 3.5.1*.

3.5.1 Additional Requirements Beyond ISO/IEC 9899:TC2

Functions that return a NaN with more than one NaN operand shall return one of the NaN operands. Functions that return a NaN operand may silence the NaN if it is a signaling NaN. A non-signaling NaN shall be converted to a non-signaling NaN. A signaling NaN shall be converted to a NaN, and should be converted to a non-signaling NaN. How the rest of the NaN payload bits or the sign of NaN is converted is undefined.

The usual allowances for rounding error (*section 3.4*) or flushing behavior (*section 3.5.3*) shall not apply for those values for which *section F.9* of ISO/IEC 9899:TC2, or *sections 3.5.1* and *3.5.3* below (and similar sections for other floating-point precisions) prescribe a result (e.g. **ceil** ($-1 < x < 0$) returns -0). Those values shall produce exactly the prescribed answers, and no other. Where the \pm symbol is used, the sign shall be preserved. For example, $\sin(\pm 0) = \pm 0$ shall be interpreted to mean $\sin(+0)$ is $+0$ and $\sin(-0)$ is -0 .

acospi (1) = $+0$.

acospi (x) returns a NaN for $|x| > 1$.

asinpi (± 0) = ± 0 .

asinpi (x) returns a NaN for $|x| > 1$.

atanpi (± 0) = ± 0 .

atanpi ($\pm\infty$) = ± 0.5 .

atan2pi ($\pm 0, -0$) = ± 1 .

atan2pi ($\pm 0, +0$) = ± 0 .

atan2pi ($\pm 0, x$) returns ± 1 for $x < 0$.

atan2pi ($\pm 0, x$) returns ± 0 for $x > 0$.

atan2pi ($y, \pm 0$) returns -0.5 for $y < 0$.

atan2pi ($y, \pm 0$) returns 0.5 for $y > 0$.

atan2pi ($\pm y, -\infty$) returns ± 1 for finite $y > 0$.

atan2pi ($\pm y, +\infty$) returns ± 0 for finite $y > 0$.

atan2pi ($\pm\infty, x$) returns ± 0.5 for finite x .

atan2pi ($\pm\infty, -\infty$) returns ± 0.75 .

atan2pi ($\pm\infty, +\infty$) returns ± 0.25 .

ceil ($-1 < x < 0$) returns -0 .

cospi (± 0) returns 1

cospi ($n + 0.5$) is $+0$ for any integer n where $n + 0.5$ is representable.

cospi ($\pm\infty$) returns a NaN.

exp10 (± 0) returns 1.

exp10 ($-\infty$) returns $+0$.

exp10 ($+\infty$) returns $+\infty$.

distance (x, y) calculates the distance from x to y without overflow or extraordinary precision loss due to underflow.

fdim (any, NaN) returns NaN.

fdim (NaN, any) returns NaN.

fmod (± 0 , NaN) returns NaN.

frexp ($\pm\infty$, *exp*) returns $\pm\infty$ and stores 0 in *exp*.

frexp (NaN, *exp*) returns the NaN and stores 0 in *exp*.

fract (*x*, *iptr*) shall not return a value greater than or equal to 1.0, and shall not return a value less than 0.

fract (+0, *iptr*) returns +0 and +0 in *iptr*.

fract (-0, *iptr*) returns -0 and -0 in *iptr*.

fract (+inf, *iptr*) returns +0 and +inf in *iptr*.

fract (-inf, *iptr*) returns -0 and -inf in *iptr*.

fract (NaN, *iptr*) returns the NaN and NaN in *iptr*.

length calculates the length of a vector without overflow or extraordinary precision loss due to underflow.

lgamma_r (*x*, *signp*) returns 0 in *signp* if *x* is zero or a negative integer.

nextafter (-0, $y > 0$) returns smallest positive denormal value.

nextafter (+0, $y < 0$) returns smallest negative denormal value.

normalize shall reduce the vector to unit length, pointing in the same direction without overflow or extraordinary precision loss due to underflow.

normalize (*v*) returns *v* if all elements of *v* are zero.

normalize (*v*) returns a vector full of NaNs if any element is a NaN.

normalize (*v*) for which any element in *v* is infinite shall proceed as if the elements in *v* were replaced as follows:

```
for( i = 0; i < sizeof(v) / sizeof(v[0]); i++ )
    v[i] = isinf(v[i]) ? copysign(1.0, v[i]) : 0.0 * v[i];
```

pow (± 0 , $-\infty$) returns $+\infty$

pown (*x*, 0) is 1 for any *x*, even zero, NaN or infinity.

pown (± 0 , *n*) is $\pm\infty$ for odd $n < 0$.

pown (± 0 , *n*) is $+\infty$ for even $n < 0$.

pown (± 0 , *n*) is +0 for even $n > 0$.

pown (± 0 , *n*) is ± 0 for odd $n > 0$.

powr (*x*, ± 0) is 1 for finite $x > 0$.

powr (± 0 , *y*) is $+\infty$ for finite $y < 0$.

powr (± 0 , $-\infty$) is $+\infty$.

powr (± 0 , *y*) is +0 for $y > 0$.

powr (+1, *y*) is 1 for finite *y*.

powr (x, y) returns NaN for $x < 0$.
powr ($\pm 0, \pm 0$) returns NaN.
powr ($+\infty, \pm 0$) returns NaN.
powr ($+1, \pm\infty$) returns NaN.
powr (x, NaN) returns the NaN for $x \geq 0$.
powr (NaN, y) returns the NaN.

rint ($-0.5 \leq x < 0$) returns -0.

remquo ($x, y, \&quo$) returns a NaN and 0 in *quo* if x is $\pm\infty$, or if y is 0 and the other argument is non-NaN or if either argument is a NaN.

rootn ($\pm 0, n$) is $\pm\infty$ for odd $n < 0$.
rootn ($\pm 0, n$) is $+\infty$ for even $n < 0$.
rootn ($\pm 0, n$) is +0 for even $n > 0$.
rootn ($\pm 0, n$) is ± 0 for odd $n > 0$.
rootn (x, n) returns a NaN for $x < 0$ and n is even.
rootn ($x, 0$) returns a NaN.

round ($-0.5 < x < 0$) returns -0.

sinpi (± 0) returns ± 0 .
sinpi ($+n$) returns +0 for positive integers n .
sinpi ($-n$) returns -0 for negative integers n .
sinpi ($\pm\infty$) returns a NaN.

tanpi (± 0) returns ± 0 .
tanpi ($\pm\infty$) returns a NaN.
tanpi (n) is **copysign**(0.0, n) for even integers n .
tanpi (n) is **copysign**(0.0, $-n$) for odd integers n .
tanpi ($n + 0.5$) for even integer n is $+\infty$ where $n + 0.5$ is representable.
tanpi ($n + 0.5$) for odd integer n is $-\infty$ where $n + 0.5$ is representable.

trunc ($-1 < x < 0$) returns -0.

3.5.2 Changes to ISO/IEC 9899: TC2 Behavior

modf behaves as though implemented by:

```
gentype modf ( gentype value, gentype *iptr )
{
    *iptr = trunc( value );
    return copysign( isinf( value ) ? 0.0 : value - *iptr, value );
}
```

rint always rounds according to round to nearest even rounding mode even if the caller is in some other rounding mode.

3.5.3 Edge Case Behavior in Flush To Zero Mode

If denormals are flushed to zero, then a function may return one of four results:

1. Any conforming result for non-flush-to-zero mode
2. If the result given by 1. is a sub-normal before rounding, it may be flushed to zero
3. Any non-flushed conforming result for the function if one or more of its sub-normal operands are flushed to zero.
4. If the result of 3. is a sub-normal before rounding, the result may be flushed to zero.

In each of the above cases, if an operand or result is flushed to zero, the sign of the zero is undefined.

If subnormals are flushed to zero, a device may choose to conform to the following edge cases for **nextafter** instead of those listed in *section 3.5.1*:

nextafter (+smallest normal, $y < +\text{smallest normal}$) = +0.
nextafter (-smallest normal, $y > -\text{smallest normal}$) = -0.
nextafter (-0, $y > 0$) returns smallest positive normal value.
nextafter (+0, $y < 0$) returns smallest negative normal value.

For clarity, subnormals or denormals are defined to be the set of representable numbers in the range $0 < x < \text{TYPE_MIN}$ and $-\text{TYPE_MIN} < x < -0$. They do not include ± 0 . A non-zero number is said to be sub-normal before rounding if after normalization, its radix-2 exponent is less than $(\text{TYPE_MIN_EXP} - 1)$.⁵³

⁵³ Here TYPE_MIN and TYPE_MIN_EXP should be substituted by constants appropriate to the floating-point type under consideration, such as FLT_MIN and FLT_MIN_EXP for float.

4 Image Addressing and Filtering

Let w_t , h_t and d_t be the width, height (or image array size for a 1D image array) and depth (or image array size for a 2D image array) of the image in pixels. Let `coord.xy` also referred to as (s, t) or `coord.xyz` also referred to as (s, t, r) be the coordinates specified to **image::read**. The sampler specified in **image::read** is used to determine how to sample the image and return an appropriate color.

4.1 Image Coordinates

This affects the interpretation of image coordinates. If image coordinates specified to **image::read** are normalized (as specified in the sampler), the s , t , and r coordinate values are multiplied by w_t , h_t , and d_t respectively to generate the unnormalized coordinate values. For image arrays, the image array coordinate (i.e. t if it is a 1D image array or r if it is a 2D image array) specified to **image::read** must always be the un-normalized image coordinate value.

Let (u, v, w) represent the unnormalized image coordinate values.

4.2 Addressing and Filter Modes

We first describe how the addressing and filter modes are applied to generate the appropriate sample locations to read from the image if the addressing mode is not `CLK_ADDRESS_REPEAT` nor `CLK_ADDRESS_MIRRORED_REPEAT`.

After generating the image coordinate (u, v, w) we apply the appropriate addressing and filter mode to generate the appropriate sample locations to read from the image.

If values in (u, v, w) are INF or NaN, the behavior of **image::read** is undefined.

Filter Mode = `CLK_FILTER_NEAREST`

When filter mode is `CLK_FILTER_NEAREST`, the image element in the image that is nearest (in Manhattan distance) to that specified by (u, v, w) is obtained. This means the image element at location (i, j, k) becomes the image element value, where

$$\begin{aligned} i &= \text{address_mode}((\text{int})\text{floor}(u)) \\ j &= \text{address_mode}((\text{int})\text{floor}(v)) \end{aligned}$$

Provisional Specification

```
k      =      address_mode((int) floor(w))
```

For a 3D image, the image element at location (i, j, k) becomes the color value. For a 2D image, the image element at location (i, j) becomes the color value.

Table 4.2.1 describes the `address_mode` function.

Addressing Mode	Result of <code>address_mode(coord)</code>
<code>CLK_ADDRESS_CLAMP_TO_EDGE</code>	<code>clamp(coord, 0, size - 1)</code>
<code>CLK_ADDRESS_CLAMP</code>	<code>clamp(coord, -1, size)</code>
<code>CLK_ADDRESS_NONE</code>	<code>coord</code>

Table 4.2.1 Addressing modes to generate texel location.

The `size` term in table 4.2.1 is w_t for u , h_t for v and d_t for w .

The `clamp` function used in *table 4.2.1* is defined as:

```
clamp(a, b, c) = return (a < b) ? b
                  : ((a > c) ? c : a)
```

If the selected texel location (i, j, k) refers to a location outside the image, the border color is used as the color value for this texel.

Filter Mode = `CLK_FILTER_LINEAR`

When filter mode is `CLK_FILTER_LINEAR`, a 2×2 square of image elements for a 2D image or a $2 \times 2 \times 2$ cube of image elements for a 3D image is selected. This 2×2 square or $2 \times 2 \times 2$ cube is obtained as follows.

Let

```
i0 = address_mode((int) floor(u - 0.5))
j0 = address_mode((int) floor(v - 0.5))
k0 = address_mode((int) floor(w - 0.5))
i1 = address_mode((int) floor(u - 0.5) + 1)
j1 = address_mode((int) floor(v - 0.5) + 1)
k1 = address_mode((int) floor(w - 0.5) + 1)
a  = frac(u - 0.5)
b  = frac(v - 0.5)
c  = frac(w - 0.5)
```

where `frac(x)` denotes the fractional part of x and is computed as $x - \text{floor}(x)$.

Provisional Specification

For a 3D image, the image element value is found as

$$\begin{aligned} T = & (1 - a) * (1 - b) * (1 - c) * T_{i0j0k0} \\ & + a * (1 - b) * (1 - c) * T_{i1j0k0} \\ & + (1 - a) * b * (1 - c) * T_{i0j1k0} \\ & + a * b * (1 - c) * T_{i1j1k0} \\ & + (1 - a) * (1 - b) * c * T_{i0j0k1} \\ & + a * (1 - b) * c * T_{i1j0k1} \\ & + (1 - a) * b * c * T_{i0j1k1} \\ & + a * b * c * T_{i1j1k1} \end{aligned}$$

where T_{ijk} is the image element at location (i, j, k) in the 3D image.

For a 2D image, the image element value is found as

$$\begin{aligned} T = & (1 - a) * (1 - b) * T_{i0j0} \\ & + a * (1 - b) * T_{i1j0} \\ & + (1 - a) * b * T_{i0j1} \\ & + a * b * T_{i1j1} \end{aligned}$$

where T_{ij} is the image element at location (i, j) in the 2D image.

If any of the selected T_{ijk} or T_{ij} in the above equations refers to a location outside the image, the border color is used as the color value for T_{ijk} or T_{ij} .

We now discuss how the addressing and filter modes are applied to generate the appropriate sample locations to read from the image if the addressing mode is CLK_ADDRESS_REPEAT.

If values in (s, t, r) are INF or NaN, the behavior of the built-in image read functions is undefined.

Filter Mode = CLK_FILTER_NEAREST

When filter mode is CLK_FILTER_NEAREST, the image element at location (i, j, k) becomes the image element value, with i, j and k computed as

$$\begin{aligned} u &= (s - \text{floor}(s)) * w_t \\ i &= (\text{int})\text{floor}(u) \\ \text{if } (i > w_t - 1) \\ &\quad i = i - w_t \\ v &= (t - \text{floor}(t)) * h_t \end{aligned}$$

Provisional Specification

```
j = (int)floor(v)
if (j > ht - 1)
    j = j - ht

w = (r - floor(r)) * dt
k = (int)floor(w)
if (k > dt - 1)
    k = k - dt
```

For a 3D image, the image element at location (i, j, k) becomes the color value. For a 2D image, the image element at location (i, j) becomes the color value.

Filter Mode = CLK_FILTER_LINEAR

When filter mode is CLK_FILTER_LINEAR, a 2 × 2 square of image elements for a 2D image or a 2 × 2 × 2 cube of image elements for a 3D image is selected. This 2 × 2 square or 2 × 2 × 2 cube is obtained as follows.

Let

```
u = (s - floor(s)) * wt
i0 = (int)floor(u - 0.5)
i1 = i0 + 1
if (i0 < 0)
    i0 = wt + i0
if (i1 > wt - 1)
    i1 = i1 - wt

v = (t - floor(t)) * ht
j0 = (int)floor(v - 0.5)
j1 = j0 + 1
if (j0 < 0)
    j0 = ht + j0
if (j1 > ht - 1)
    j1 = j1 - ht

w = (r - floor(r)) * dt
k0 = (int)floor(w - 0.5)
k1 = k0 + 1
if (k0 < 0)
    k0 = dt + k0
if (k1 > dt - 1)
    k1 = k1 - dt

a = frac(u - 0.5)
b = frac(v - 0.5)
```

Provisional Specification

$$c = \text{frac}(w - 0.5)$$

where $\text{frac}(x)$ denotes the fractional part of x and is computed as $x - \text{floor}(x)$.

For a 3D image, the image element value is found as

$$\begin{aligned} T = & (1 - a) * (1 - b) * (1 - c) * T_{i0j0k0} \\ & + a * (1 - b) * (1 - c) * T_{i1j0k0} \\ & + (1 - a) * b * (1 - c) * T_{i0j1k0} \\ & + a * b * (1 - c) * T_{i1j1k0} \\ & + (1 - a) * (1 - b) * c * T_{i0j0k1} \\ & + a * (1 - b) * c * T_{i1j0k1} \\ & + (1 - a) * b * c * T_{i0j1k1} \\ & + a * b * c * T_{i1j1k1} \end{aligned}$$

where T_{ijk} is the image element at location (i, j, k) in the 3D image.

For a 2D image, the image element value is found as

$$\begin{aligned} T = & (1 - a) * (1 - b) * T_{i0j0} \\ & + a * (1 - b) * T_{i1j0} \\ & + (1 - a) * b * T_{i0j1} \\ & + a * b * T_{i1j1} \end{aligned}$$

where T_{ij} is the image element at location (i, j) in the 2D image.

We now discuss how the addressing and filter modes are applied to generate the appropriate sample locations to read from the image if the addressing mode is `CLK_ADDRESS_MIRRORED_REPEAT`. The `CLK_ADDRESS_MIRRORED_REPEAT` addressing mode causes the image to be read as if it is tiled at every integer seam with the interpretation of the image data flipped at each integer crossing. For example, the (s, t, r) coordinates between 2 and 3 are addressed into the image as coordinates from 1 down to 0. If values in (s, t, r) are INF or NaN, the behavior of the built-in image read functions is undefined.

Filter Mode = `CLK_FILTER_NEAREST`

When filter mode is `CLK_FILTER_NEAREST`, the image element at location (i, j, k) becomes the image element value, with i, j and k computed as

$$\begin{aligned} s' &= 2.0f * \text{rint}(0.5f * s) \\ s' &= \text{fabs}(s - s') \end{aligned}$$

```

u = s' * wt
i = (int)floor(u)
i = min(i, wt - 1)

t' = 2.0f * rint(0.5f * t)
t' = fabs(t - t')
v = t' * ht
j = (int)floor(v)
j = min(j, ht - 1)

r' = 2.0f * rint(0.5f * r)
r' = fabs(r - r')
w = r' * dt
k = (int)floor(w)
k = min(k, dt - 1)

```

For a 3D image, the image element at location (i, j, k) becomes the color value. For a 2D image, the image element at location (i, j) becomes the color value.

Filter Mode = CLK_FILTER_LINEAR

When filter mode is CLK_FILTER_LINEAR, a 2 × 2 square of image elements for a 2D image or a 2 × 2 × 2 cube of image elements for a 3D image is selected. This 2 × 2 square or 2 × 2 × 2 cube is obtained as follows.

Let

```

s' = 2.0f * rint(0.5f * s)
s' = fabs(s - s')
u = s' * wt
i0 = (int)floor(u - 0.5f)
i1 = i0 + 1
i0 = max(i0, 0)
i1 = min(i1, wt - 1)

t' = 2.0f * rint(0.5f * t)
t' = fabs(t - t')
v = t' * ht
j0 = (int)floor(v - 0.5f)
j1 = j0 + 1
j0 = max(j0, 0)
j1 = min(j1, ht - 1)

r' = 2.0f * rint(0.5f * r)
r' = fabs(r - r')
w = r' * dt

```

Provisional Specification

```
k0 = (int) floor(w - 0.5f)
k1 = k0 + 1
k0 = max(k0, 0)
k1 = min(k1, dt - 1)

a = frac(u - 0.5)
b = frac(v - 0.5)
c = frac(w - 0.5)
```

where $\text{frac}(x)$ denotes the fractional part of x and is computed as $x - \text{floor}(x)$.

For a 3D image, the image element value is found as

$$\begin{aligned} T = & (1 - a) * (1 - b) * (1 - c) * T_{i0j0k0} \\ & + a * (1 - b) * (1 - c) * T_{i1j0k0} \\ & + (1 - a) * b * (1 - c) * T_{i0j1k0} \\ & + a * b * (1 - c) * T_{i1j1k0} \\ & + (1 - a) * (1 - b) * c * T_{i0j0k1} \\ & + a * (1 - b) * c * T_{i1j0k1} \\ & + (1 - a) * b * c * T_{i0j1k1} \\ & + a * b * c * T_{i1j1k1} \end{aligned}$$

where T_{ijk} is the image element at location (i, j, k) in the 3D image.

For a 2D image, the image element value is found as

$$\begin{aligned} T = & (1 - a) * (1 - b) * T_{i0j0} \\ & + a * (1 - b) * T_{i1j0} \\ & + (1 - a) * b * T_{i0j1} \\ & + a * b * T_{i1j1} \end{aligned}$$

where T_{ij} is the image element at location (i, j) in the 2D image.

For a 1D image, the image element value is found as

$$\begin{aligned} T = & (1 - a) * T_{i0} \\ & + a * T_{i1} \end{aligned}$$

where T_i is the image element at location (i) in the 1D image.

NOTE

If the sampler is specified as using unnormalized coordinates (floating-point or integer coordinates), filter mode set to CLK_FILTER_NEAREST and addressing mode set to one of the following modes - CLK_ADDRESS_NONE, CLK_ADDRESS_CLAMP_TO_EDGE or CLK_ADDRESS_CLAMP, the location of the image element in the image given by (i, j, k) will be computed without any loss of precision.

*For all other sampler combinations of normalized or unnormalized coordinates, filter and addressing modes, the relative error or precision of the addressing mode calculations and the image filter operation are not defined by this revision of the OpenCL specification. To ensure a minimum precision of image addressing and filter calculations across any OpenCL device, for these sampler combinations, developers should unnormalize the image coordinate in the kernel and implement the linear filter in the kernel with appropriate calls to **read_image{f|i|ui}** with a sampler that uses unnormalized coordinates, filter mode set to CLK_FILTER_NEAREST, addressing mode set to CLK_ADDRESS_NONE, CLK_ADDRESS_CLAMP_TO_EDGE or CLK_ADDRESS_CLAMP and finally performing the interpolation of color values read from the image to generate the filtered color value.*

4.3 Conversion Rules

In this section we discuss conversion rules that are applied when reading and writing images in a kernel.

4.3.1 Conversion rules for normalized integer channel data types

In this section we discuss converting normalized integer channel data types to floating-point values and vice-versa.

4.3.1.1 Converting normalized integer channel data types to floating-point values

For images created with image channel data type of CL_UNORM_INT8 and CL_UNORM_INT16, **image::read** will convert the channel values from an 8-bit or 16-bit unsigned integer to normalized floating-point values in the range $[0.0f \dots 1.0]$.

For images created with image channel data type of CL_SNORM_INT8 and CL_SNORM_INT16, **image::read** will convert the channel values from an 8-bit or 16-bit signed integer to normalized floating-point values in the range $[-1.0 \dots 1.0]$.

These conversions are performed as follows:

CL_UNORM_INT8 (8-bit unsigned integer) \rightarrow float

$$\text{normalized float value} = (\text{float})c / 255.0f$$

Provisional Specification

CL_UNORM_INT_101010 (10-bit unsigned integer) → float

normalized float value = (float)c / 1023.0f

CL_UNORM_INT16 (16-bit unsigned integer) → float

normalized float value = (float)c / 65535.0f

CL_SNORM_INT8 (8-bit signed integer) → float

normalized float value = max(-1.0f, (float)c / 127.0f)

CL_SNORM_INT16 (16-bit signed integer) → float

normalized float value = max(-1.0f, (float)c / 32767.0f)

The precision of the above conversions is <= 1.5 ulp except for the following cases.

For CL_UNORM_INT8

0 must convert to 0.0f and
255 must convert to 1.0f

For CL_UNORM_INT_101010

0 must convert to 0.0f and
1023 must convert to 1.0f

For CL_UNORM_INT16

0 must convert to 0.0f and
65535 must convert to 1.0f

For CL_SNORM_INT8

-128 and -127 must convert to -1.0f,
0 must convert to 0.0f and
127 must convert to 1.0f

For CL_SNORM_INT16

-32768 and -32767 must convert to -1.0f,
0 must convert to 0.0f and
32767 must convert to 1.0f

4.3.1.2 Converting floating-point values to normalized integer channel data types

For images created with image channel data type of CL_UNORM_INT8 and CL_UNORM_INT16, **image::write** will convert the floating-point color value to an 8-bit or 16-bit unsigned integer.

For images created with image channel data type of CL_SNORM_INT8 and CL_SNORM_INT16, **image::write** will convert the floating-point color value to an 8-bit or 16-bit signed integer.

The preferred method for how conversions from floating-point values to normalized integer values are performed is as follows:

float → CL_UNORM_INT8 (8-bit unsigned integer)

```
convert_cast<uchar,
              saturate::on,
              roundingmode::rte>(f * 255.0f)
```

float → CL_UNORM_INT_101010 (10-bit unsigned integer)

```
min(convert_cast<ushort,
               saturate::on,
               roundingmode::rte>(f * 1023.0f),
    0x3fff)
```

float → CL_UNORM_INT16 (16-bit unsigned integer)

```
convert_cast<ushort,
             saturate::on,
             roundingmode::rte>(f * 65535.0f)
```

float → CL_SNORM_INT8 (8-bit signed integer)

```
convert_cast<char,
            saturate::on,
            roundingmode::rte>(f * 127.0f)
```

float → CL_SNORM_INT16 (16-bit signed integer)

```
convert_cast<short,
            saturate::on,
            roundingmode::rte>(f * 32767.0f)
```

Please refer to *section 1.2.3.3* for out-of-range behavior and saturated conversions rules.

OpenCL implementations may choose to approximate the rounding mode used in the conversions described above. If a rounding mode other than round to nearest even (`_rte`) is used, the absolute error of the implementation dependant rounding mode vs. the result produced by the round to nearest even rounding mode must be ≤ 0.6 .

float \rightarrow CL_UNORM_INT8 (8-bit unsigned integer)

```

Let fpreferred =
    convert_cast<uchar,
        saturate::on,
        rounding_mode::rte>(f * 255.0f)
Let fapprox =
    convert_uchar_sat_impl-rounding-mode(
        f * 255.0f)

fabs(fpreferred - fapprox) must be  $\leq 0.6$ 

```

float \rightarrow CL_UNORM_INT_101010 (10-bit unsigned integer)

```

Let fpreferred =
    convert_cast<ushort,
        saturate::on,
        roundingmode::rte>(f * 1023.0f)
Let fapprox =
    convert_ushort_sat_impl-rounding-mode(
        f * 1023.0f)

fabs(fpreferred - fapprox) must be  $\leq 0.6$ 

```

float \rightarrow CL_UNORM_INT16 (16-bit unsigned integer)

```

Let fpreferred =
    convert_cast<ushort,
        saturate::on,
        roundingmode::rte>(f * 65535.0f)
Let fapprox =
    convert_ushort_sat_impl-rounding-mode(
        f * 65535.0f)

fabs(fpreferred - fapprox) must be  $\leq 0.6$ 

```

float \rightarrow CL_SNORM_INT8 (8-bit signed integer)

```

Let fpreferred =

```

Provisional Specification

```
convert_cast<char,  
            saturate::on,  
            roundingmode::rte>(f * 127.0f)  
Let fapprox =  
    convert_char_sat_impl_rounding_mode(  
        f * 127.0f)
```

$\text{fabs}(f_{\text{preferred}} - f_{\text{approx}})$ must be ≤ 0.6

float \rightarrow CL_SNORM_INT16 (16-bit signed integer)

```
Let fpreferred =  
    convert_cast<short,  
                saturate::on,  
                roundingmode::rte>(f * 32767.0f)  
Let fapprox =  
    convert_short_sat_impl-rounding-mode(  
        f * 32767.0f)
```




$\text{fabs}(f_{\text{preferred}} - f_{\text{approx}})$ must be ≤ 0.6

4.3.2 Conversion rules for half precision floating-point channel data type

For images created with a channel data type of CL_HALF_FLOAT, the conversions from half to float are lossless (as described in *section 1.1.1.1*). Conversions from float to half round the mantissa using the round to nearest even or round to zero rounding mode. Denormalized numbers for the half data type which may be generated when converting a float to a half may be flushed to zero. A float NaN must be converted to an appropriate NaN in the half type. A float INF must be converted to an appropriate INF in the half type.

4.3.3 Conversion rules for floating-point channel data type

The following rules apply for reading and writing images created with channel data type of CL_FLOAT.

-  NaNs may be converted to a NaN value(s) supported by the device.
-  Denorms can be flushed to zero.
-  All other values must be preserved.

4.3.4 Conversion rules for signed and unsigned 8-bit, 16-bit and 32-bit integer channel data types

Calls to **image::read** with channel data type values of CL_SIGNED_INT8, CL_SIGNED_INT16 and CL_SIGNED_INT32 return the unmodified integer values stored in the image at specified location.

Calls to **image::read** with channel data type values of CL_UNSIGNED_INT8, CL_UNSIGNED_INT16 and CL_UNSIGNED_INT32 return the unmodified integer values stored in the image at specified location.

Calls to **image::write** will perform one of the following conversions:

32 bit signed integer → 8-bit signed integer

```
convert_cast<char, saturate::on>(i)
```

32 bit signed integer → 16-bit signed integer

```
convert_cast<short, saturate::on>(i)
```

32 bit signed integer → 32-bit signed integer

```
no conversion is performed
```

Calls to **image::write** will perform one of the following conversions:

32 bit unsigned integer → 8-bit unsigned integer

```
convert_cast<uchar, saturate::on>(i)
```

32 bit unsigned integer → 16-bit unsigned integer

```
convert_cast<ushort, saturate::on>(i)
```

32 bit unsigned integer → 32-bit unsigned integer

```
no conversion is performed
```

The conversions described in this section must be correctly saturated.

4.3.5 Conversion rules for sRGBA and sBGRA images

Standard RGB data, which roughly displays colors in a linear ramp of luminosity levels such that an average observer, under average viewing conditions, can view them as perceptually equal steps on an average display. All 0's maps to 0.0f, and all 1's maps to 1.0f. The sequence of unsigned integer encodings between all 0's and all 1's represent a nonlinear progression in the floating-point interpretation of the numbers between 0.0f to 1.0f. For more detail, see the SRGB color standard, IEC 61996-2-1, at IEC (International Electrotechnical Commission).

Conversion from sRGB space is automatically done by **image::read** built-in functions if the image channel order is one of the sRGB values described above. When reading from an sRGB image, the conversion from sRGB to linear RGB is performed before the filter specified in the sampler specified to read_imagef is applied. If the format has an alpha channel, the alpha data is stored in linear color space. Conversion to sRGB space is automatically done by **image::write** built-in functions if the image channel order is one of the sRGB values described above and the device supports writing to sRGB images.

If the format has an alpha channel, the alpha data is stored in linear color space.

The following is the conversion rule for converting a normalized 8-bit unsigned integer sRGB color value to a floating-point linear RGB color value using **image::read**.

```
Convert the normalized 8-bit unsigned integer R, G and
B channel values to a floating-point value (call it c)
as per rules described in section 4.3.1.1.
```

```
if (c <= 0.04045),
    result = c / 12.92;
else
    result = powr((c + 0.055) / 1.055, 2.4);
```

The resulting floating point value, if converted back to an sRGB value without rounding to a 8-bit unsigned integer value, must be within 0.5 ulp of the original sRGB value.

The following are the conversion rules for converting a linear RGB floating-point color value (call it c) to a normalized 8-bit unsigned integer sRGB value using write_imagef.

```
if (isnan(c)) c = 0.0;
if (c > 1.0)
    c = 1.0;
else if (c < 0.0)
```

Provisional Specification

```
c = 0.0;
else if (c < 0.0031308)
    c = 12.92 * c;
else
    c = 1.055 * powr(c, 1.0/2.4) - 0.055;

convert to integer scale i.e. c = c * 255.0
convert to integer:
    c = c + 0.5
    drop the decimal fraction, and the remaining
    floating-point(integral) value is converted
    directly to an integer.
```

The precision of the above conversion should be such that $\text{fabs}(\text{reference result} - \text{integer result}) \leq 0.6$.

4.4 Selecting an Image from an Image Array

Let (u, v, w) represent the unnormalized image coordinate values for reading from and/or writing to a 2D image in a 2D image array.

The 2D image layer selected is computed as:

```
layer = clamp(rint(w), 0, dt - 1)
```

Let (u, v) represent the unnormalized image coordinate values for reading from and/or writing to a 1D image in a 1D image array.

The 1D image layer selected is computed as:

```
layer = clamp(rint(v), 0, ht - 1)
```